
Introduction to Process

Operating Systems

School of Data & Computer Science
Sun Yat-sen University

Lecture Notes: os_sysu@163.com
Instructor: Guoyang Cai
email: isscgy@mail.sysu.edu.cn



■ Contents

- Basic Concepts
- Process Table and Process Control Block
- Process States and Transitions
- Operations on Process
 - Process Creation
 - Process Termination
- Unix and Linux Examples
- Process Scheduling
- Process Switching

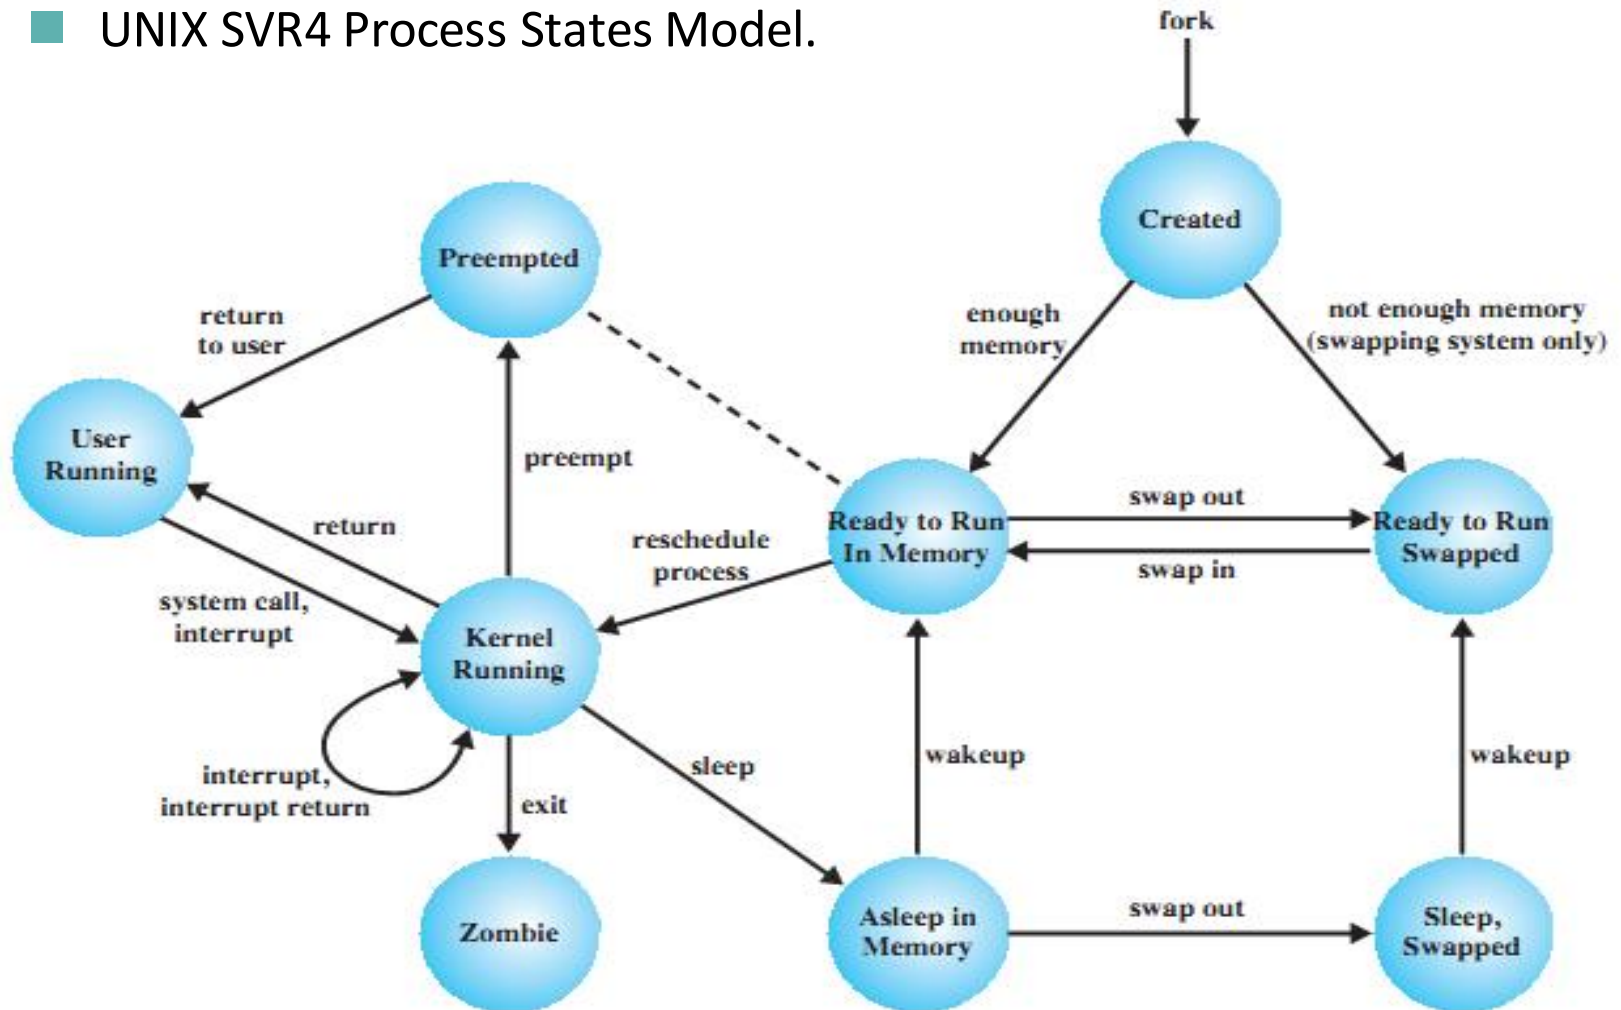


■ UNIX SVR4 Process States

- **User Running** - Executing in user mode.
- **Kernel Running** - Executing in kernel mode.
- **Ready to Run, in Memory** - Ready to run as soon as the kernel schedules it.
- **Asleep in Memory** - Unable to execute until an event occurs; process is in main memory (a blocked state).
- **Ready to Run, Swapped** - Process is ready to run, but the swapper must swap the process into main memory before the kernel can schedule it to execute.
- **Sleeping, Swapped** - The process is awaiting an event and has been swapped to secondary storage (a blocked state).
- **Preempted** - Process is returning from kernel to user mode, but the kernel preempts it and does a process switch to schedule another process.
- **Created** - Process is newly created and not yet ready to run.
- **Zombie** - Process no longer exists, but it leaves a record for its parent process to collect.

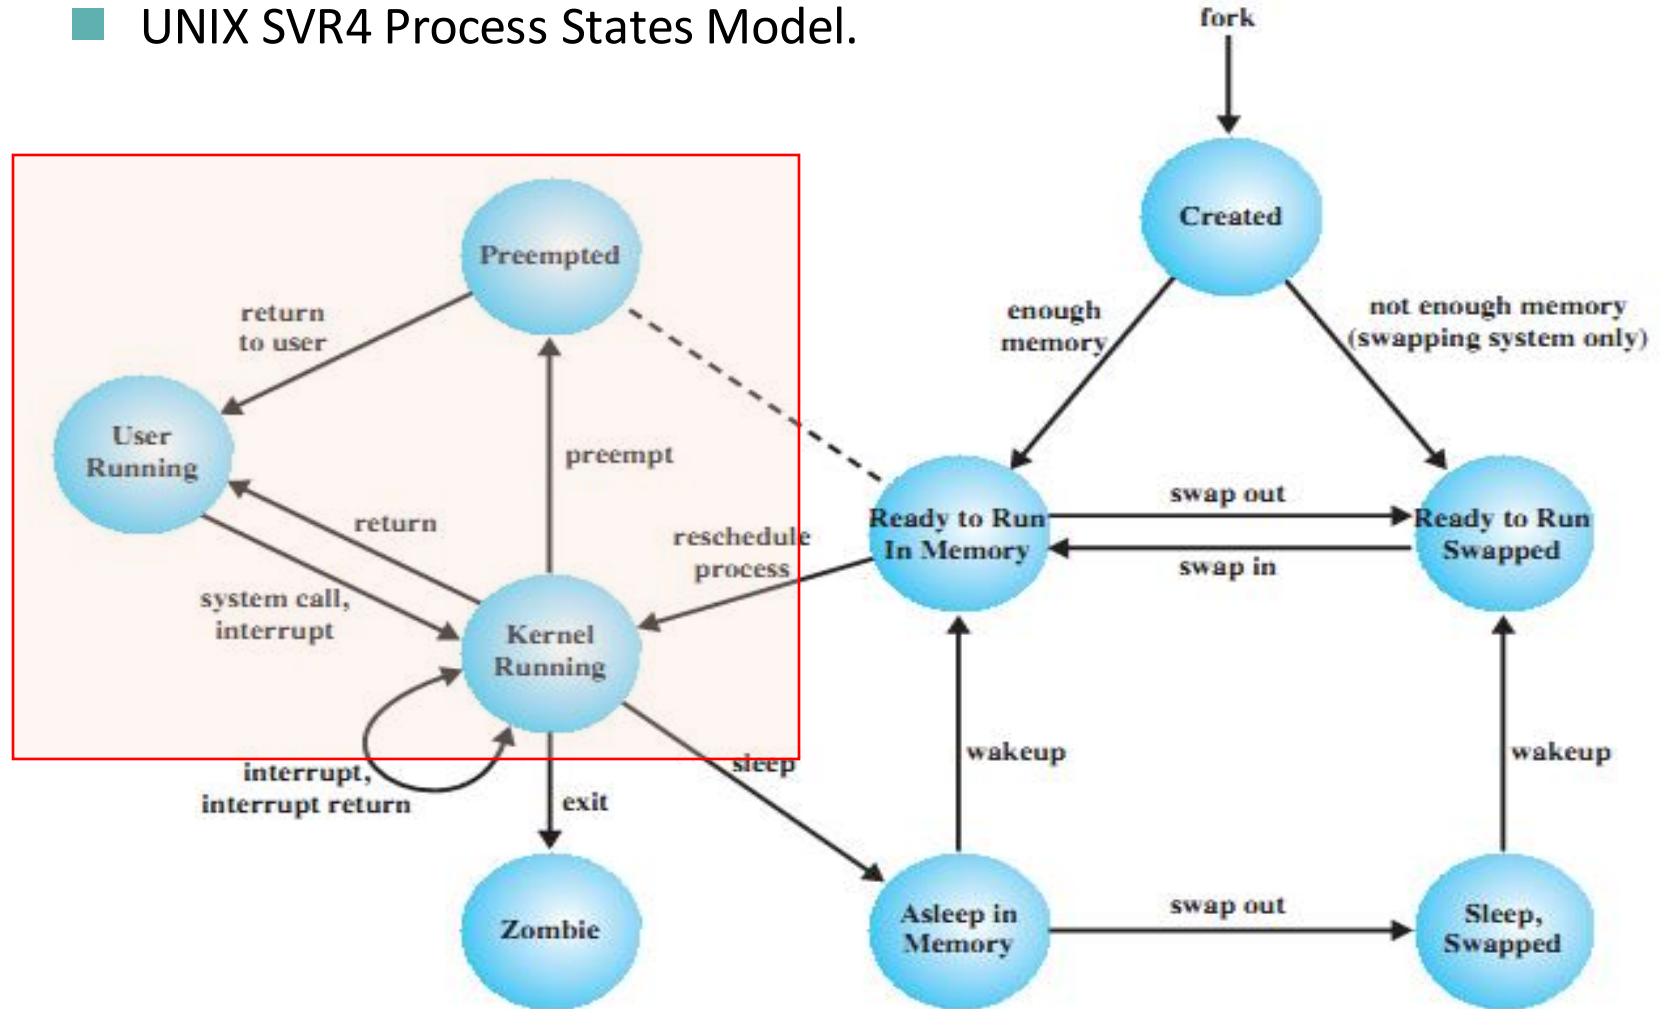
UNIX SVR4 Process States

UNIX SVR4 Process States Model.



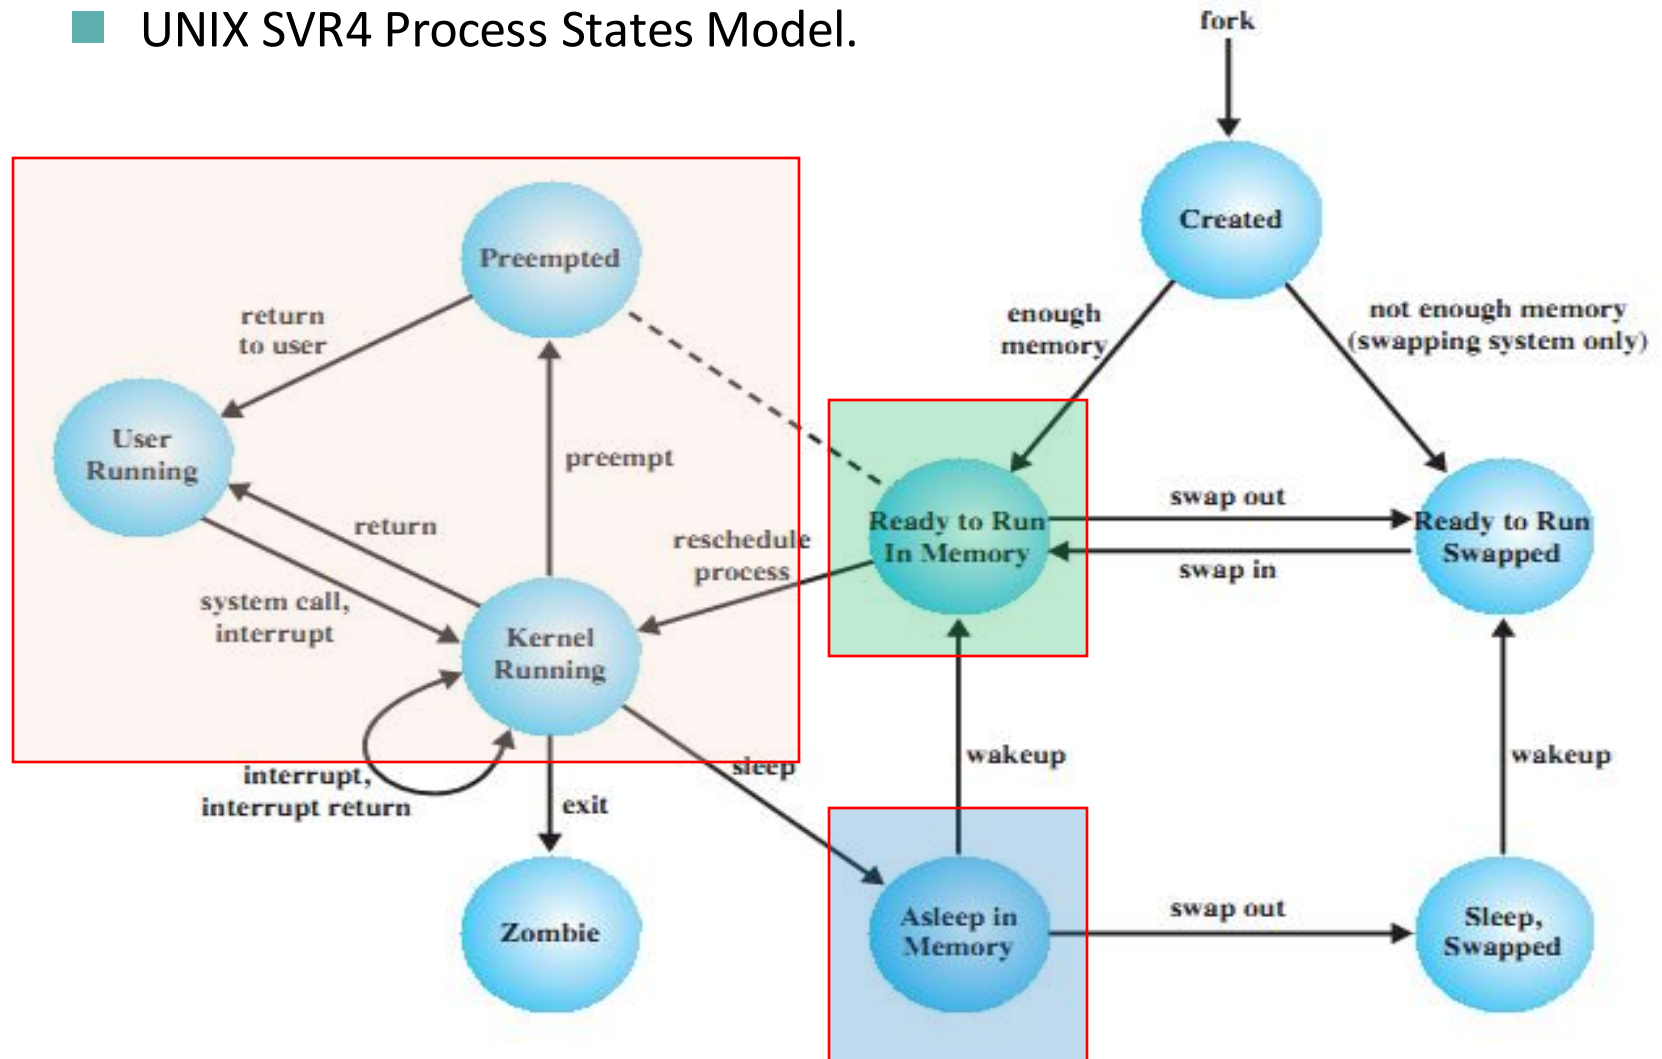
UNIX SVR4 Process States

UNIX SVR4 Process States Model.



UNIX SVR4 Process States

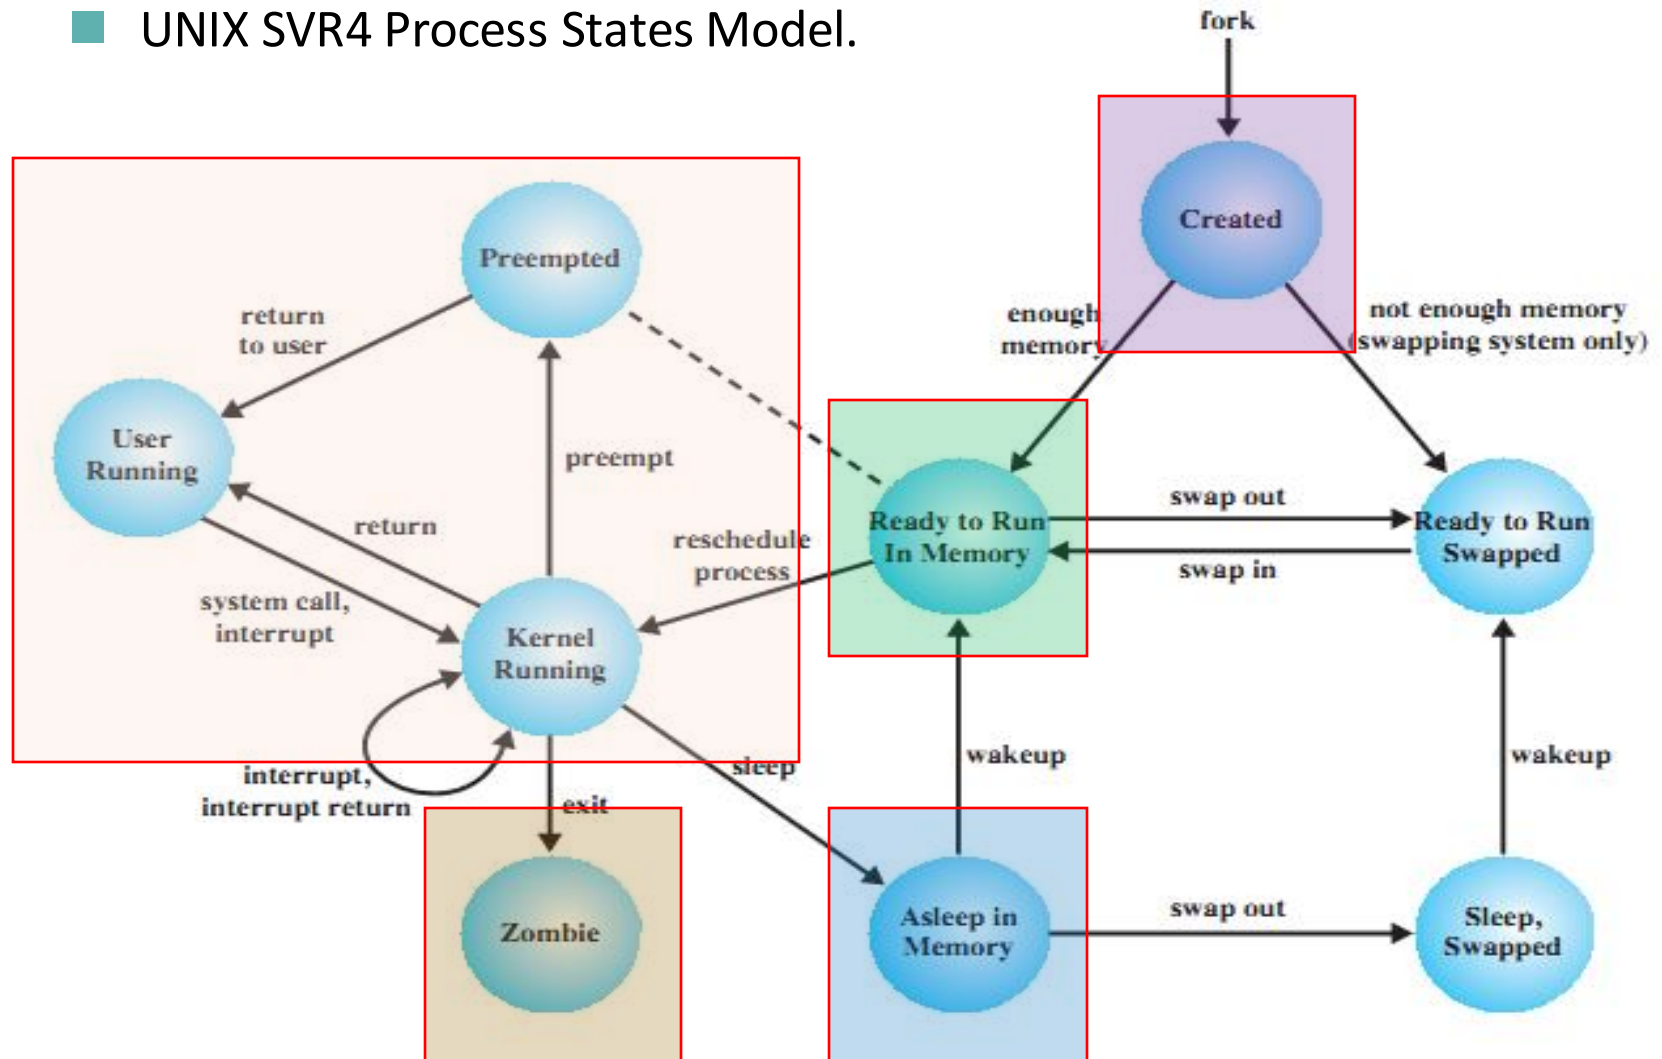
UNIX SVR4 Process States Model.



Three-state Transition

UNIX SVR4 Process States

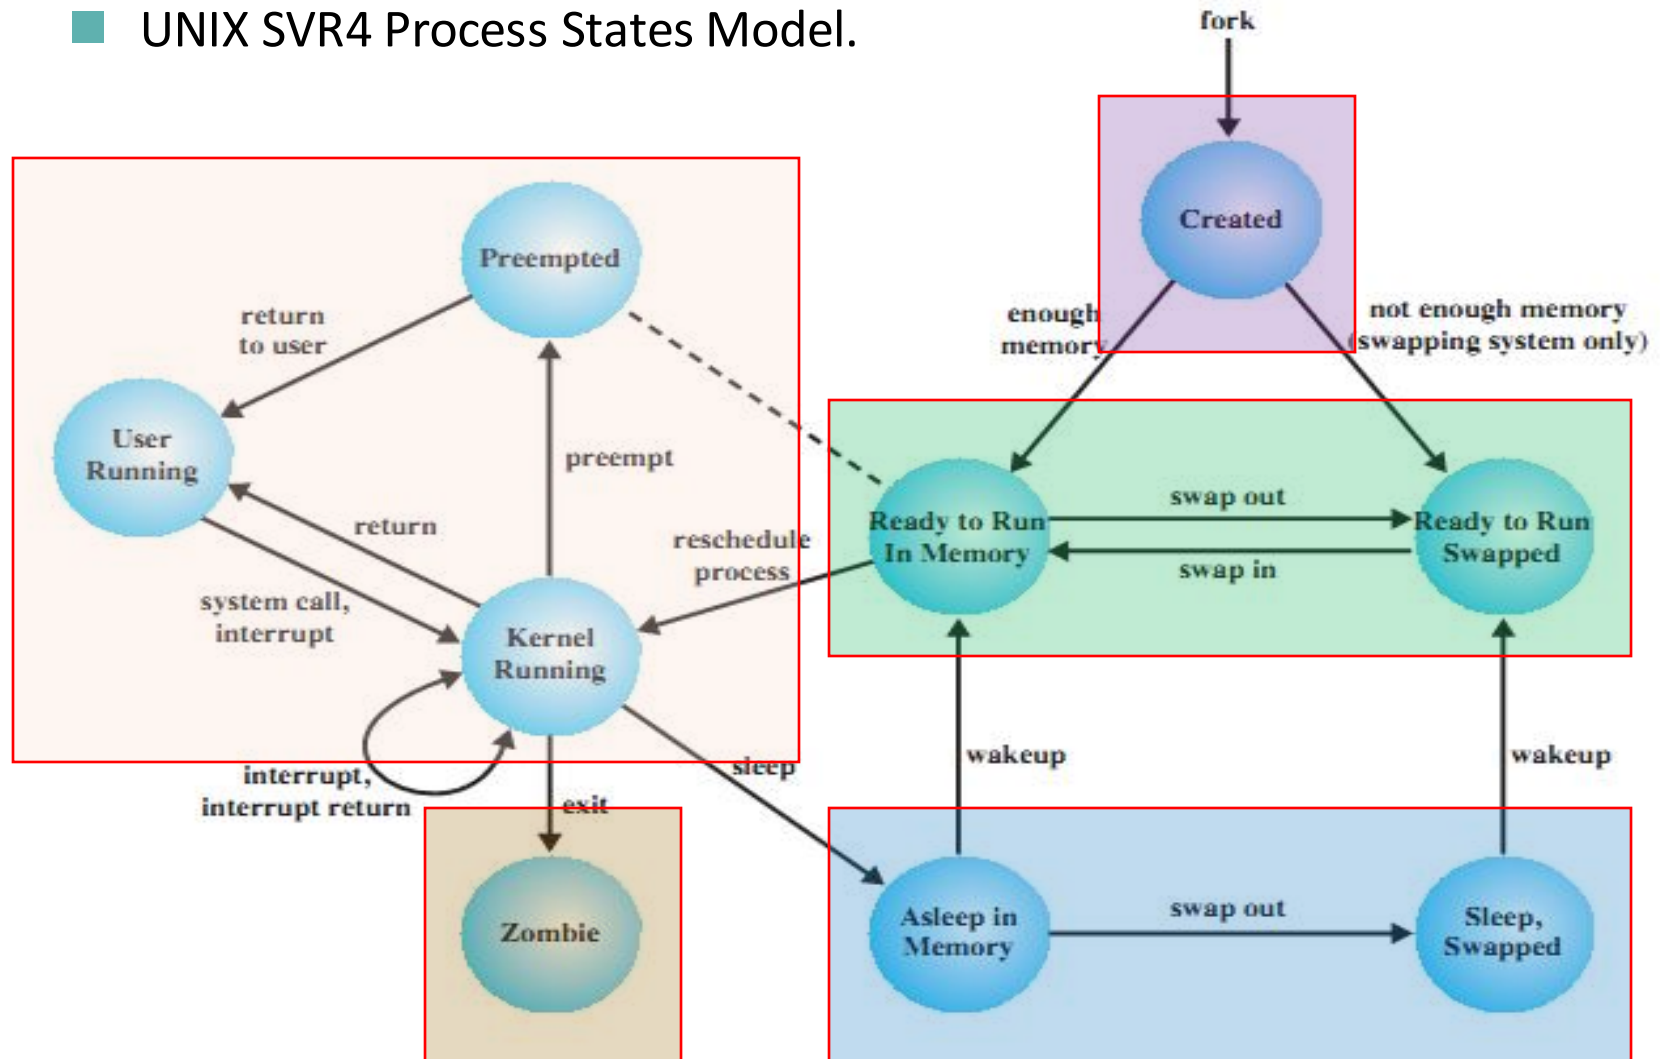
UNIX SVR4 Process States Model.



Five-state Transition

UNIX SVR4 Process States

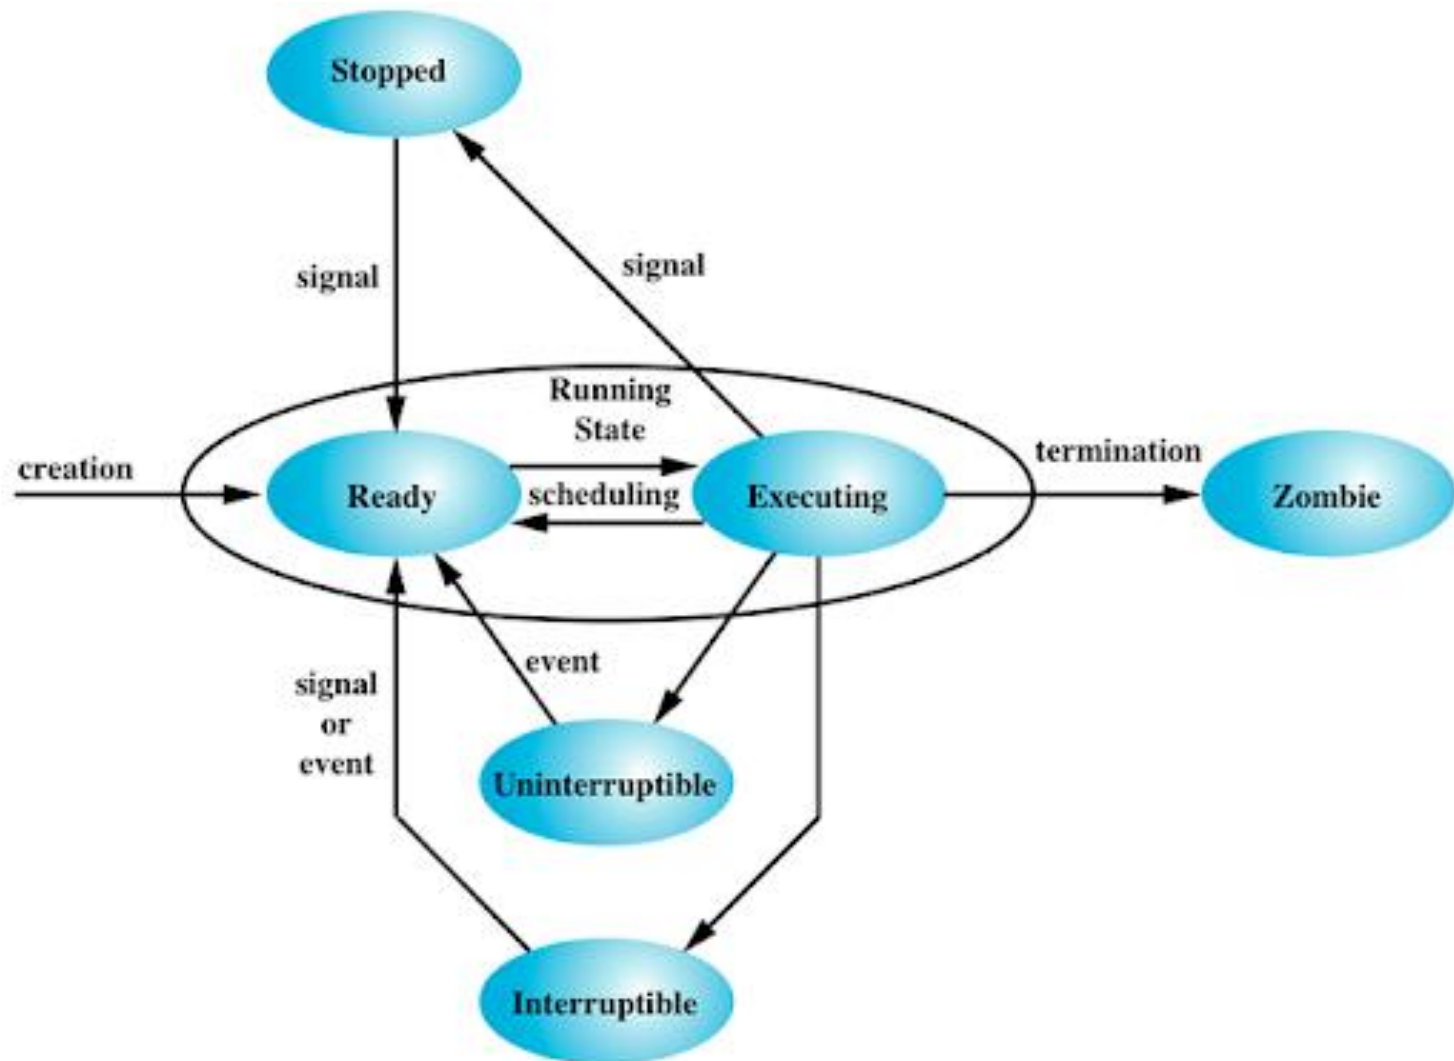
UNIX SVR4 Process States Model.



Five-state Transition with Swapping

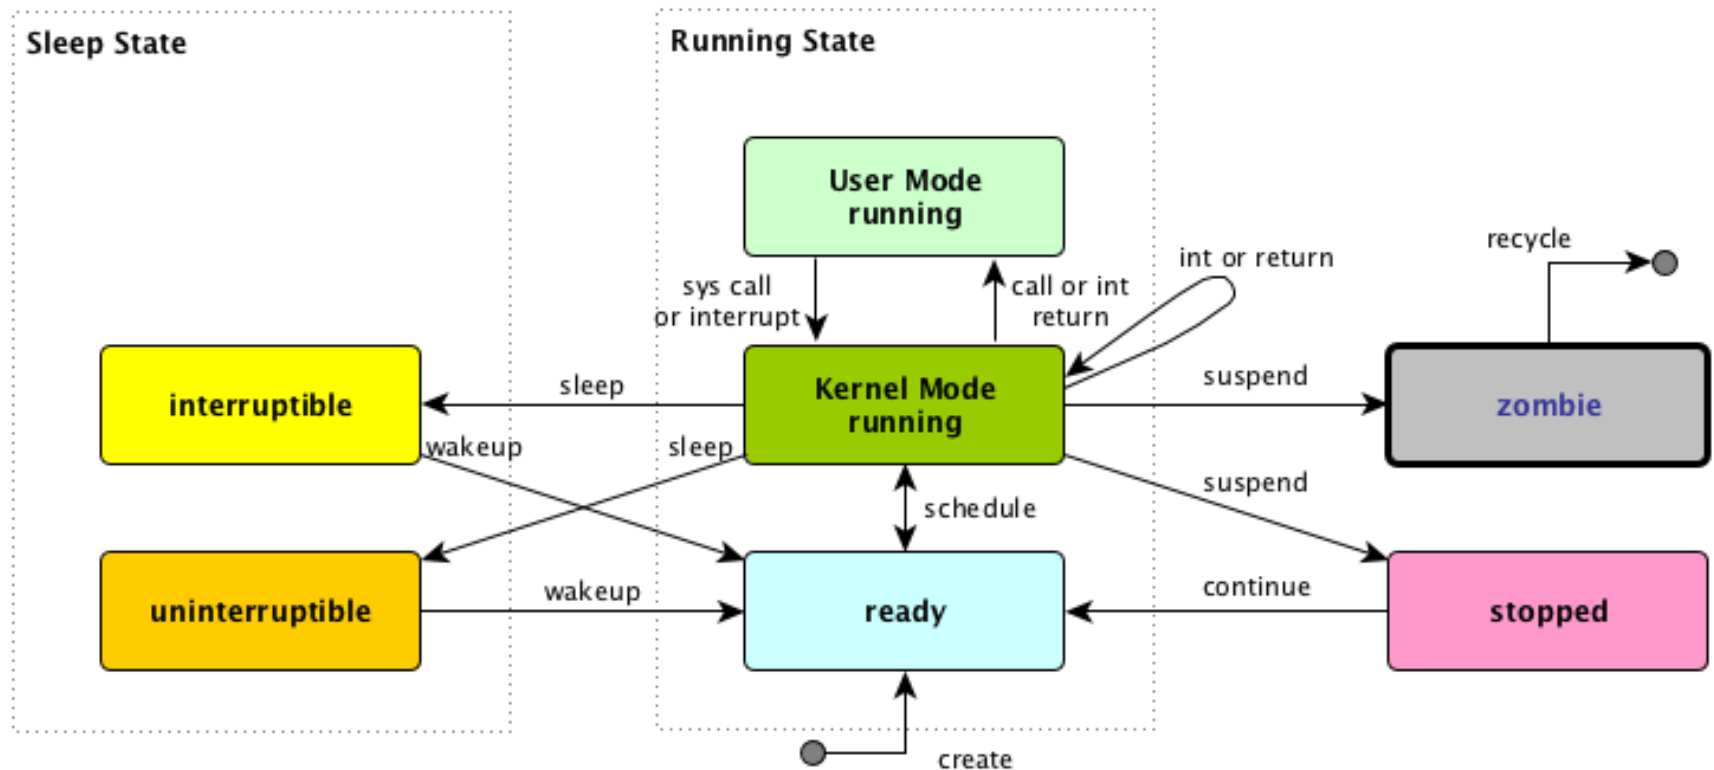
■ Process Representation in Linux

■ Linux Process State.



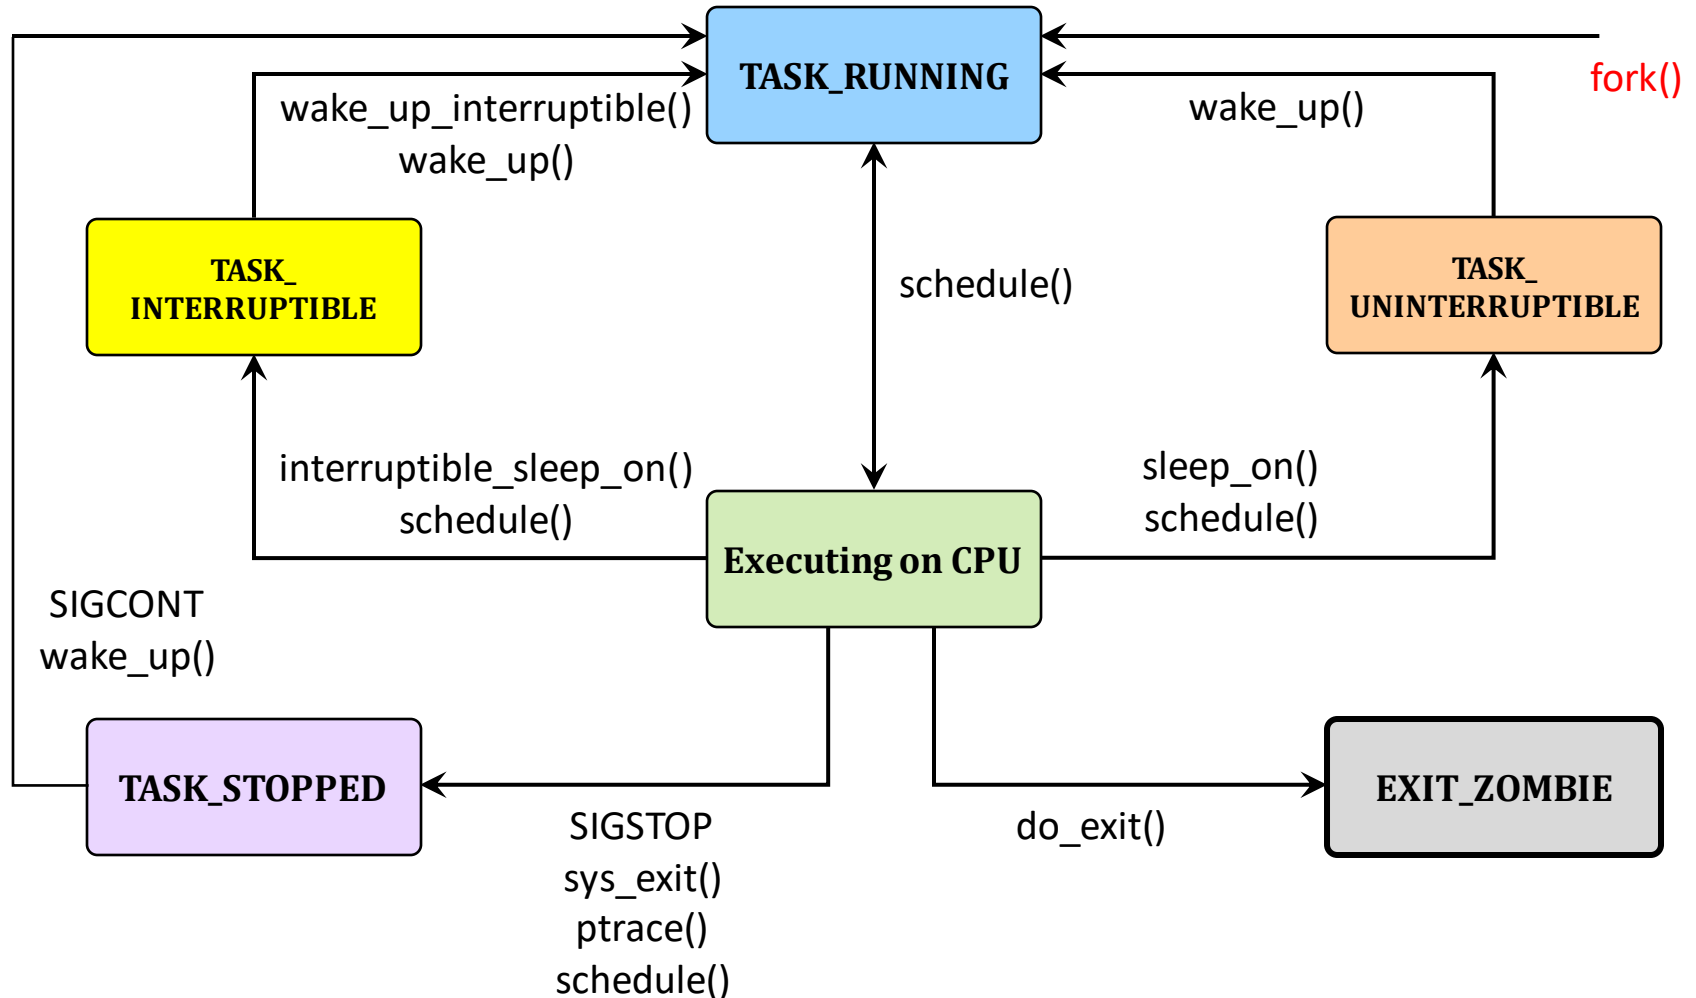
■ Process Representation in Linux

■ Linux Process State.



Process Representation in Linux

Linux Process State.





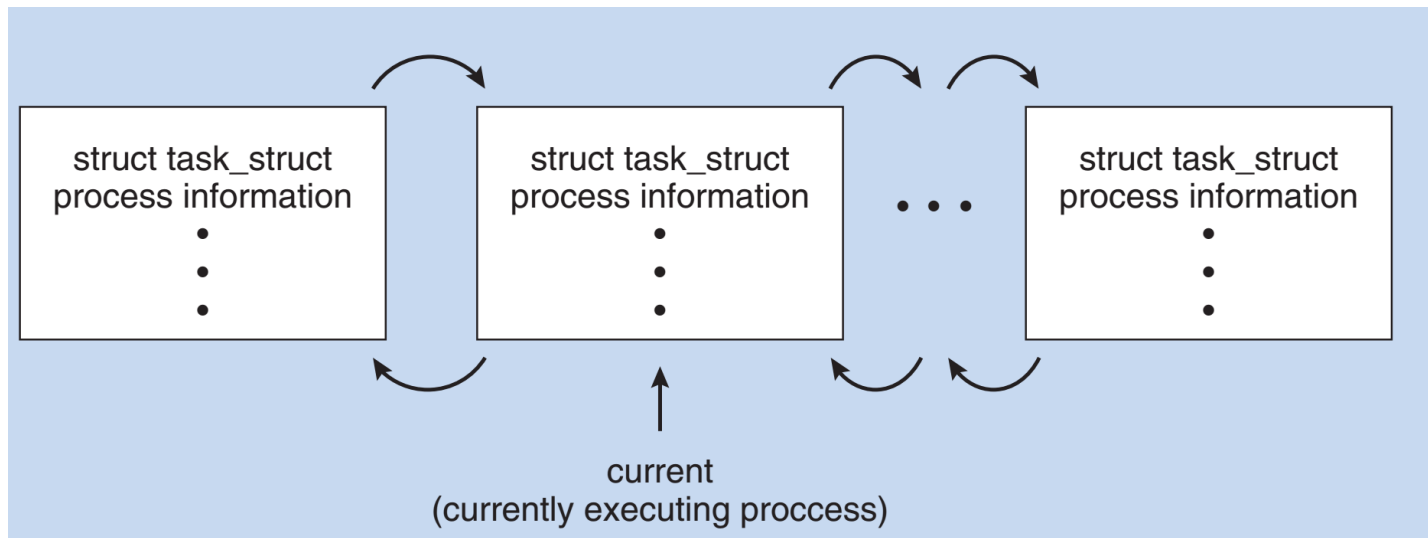
■ Process Representation in Linux

- The PCB in Linux is represented by the C structure **task_struct** included in `<linux/sched.h>`.

- **task_struct** is declared in `linux/sched.h`.

```
$ sudo apt-get install linux-source
$ tar ...
$ vim /usr/src/linux-source-X.XX.X/include/linux/sched.h
```

- The Linux kernel uses a *circular doubly-linked list* of `struct task_struct` to store these process descriptors.





■ Process Representation in Linux

- **task_struct** contains all the necessary information for representing a process, including
 - PID
 - the state of the process
 - processor registers
 - scheduling and memory-management information
 - list of open files
 - pointers to the process's parent and a list of its children and siblings.
- Some of these fields include:

```
long state; /* state of the process */
struct sched_entity se; /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



■ Process Representation in Linux

- Here are a few fields from kernel 2.6.15-1.2054_FC5, starting at line 701: (what about these fields in your current Ubuntu version?)

```
701.struct task_struct {
702.    volatile long state; /*-1 unrunnable,0 runnable,>0 stopped */
703.    struct thread_info *thread_info;
704.    . . .
767.    /* PID/PID hash table linkage. */
768.    struct pid pids[PIDTYPE_MAX];
769.    . . .
798.    char comm[TASK_COMM_LEN]; /* executable name excluding path
```

- Type definition of volatile long.

```
#define TASK_RUNNING 0
#define TASK_INTERRUPTIBLE 1
#define TASK_UNINTERRUPTIBLE 2
#define TASK_STOPPED 4
#define TASK_TRACED 8 /* in tsk->exit_state */
#define EXIT_ZOMBIE 16
#define EXIT_DEAD 32 /* in tsk->state_again */
#define TASK_NONINTERACTIVE 64
```



■ Overview

- The objective of multiprogramming is to have some process running at all times, to maximize CPU utilization.
 - The objective of time sharing is to switch the CPU among processes so frequently that users can interact with each program while it is running.
- *Process scheduler* selects among available processes for next execution on CPU.
 - For a single-processor system, there will be only one running process and the rest will have to wait until the CPU is free and can be rescheduled.



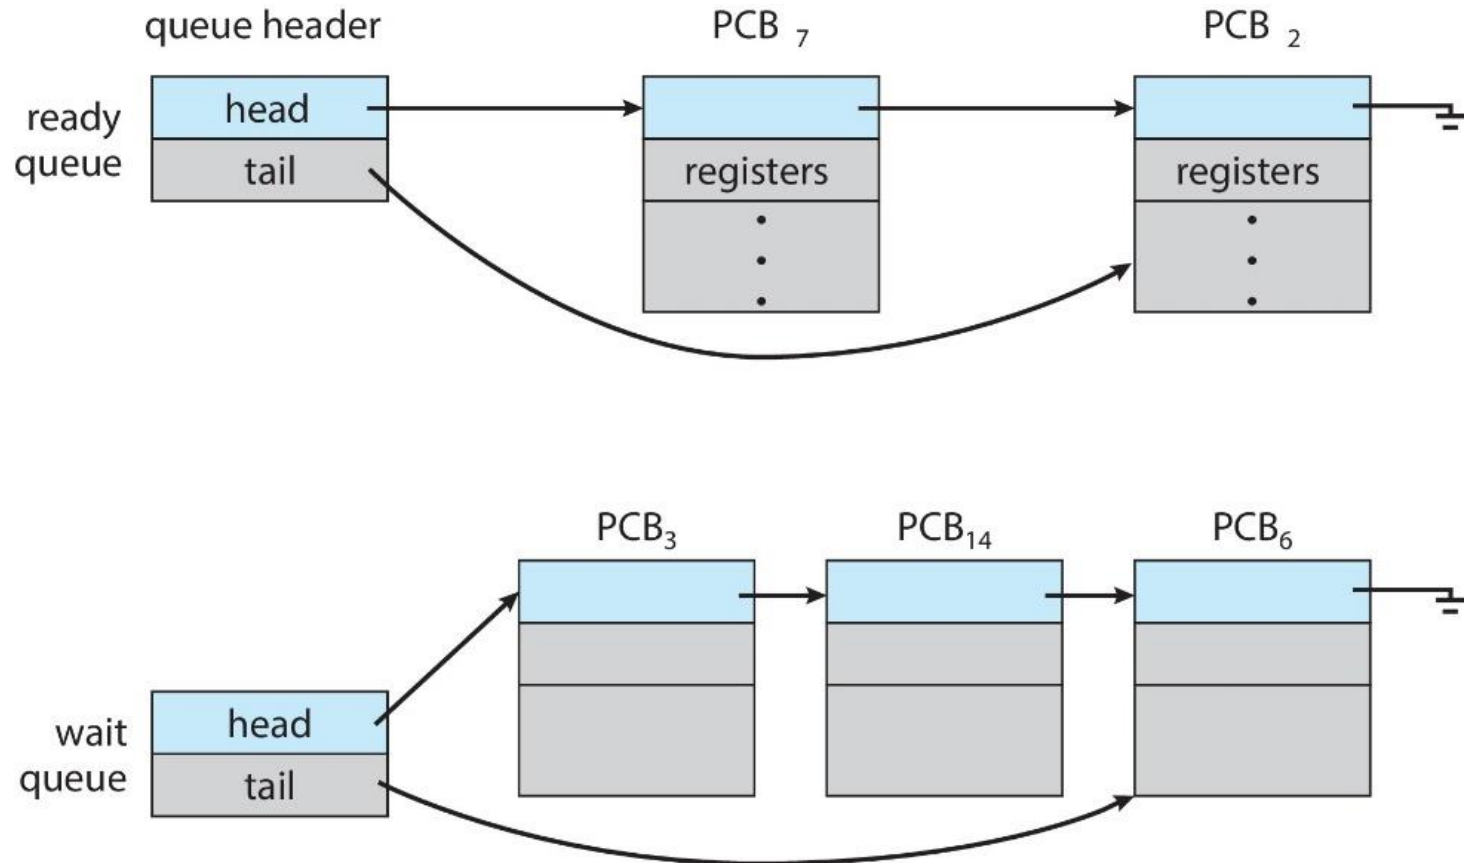
■ Overview

- *Scheduling queues* of processes are to be maintained:
 - Job Queue
 - set of all processes in the system
 - Ready Queue
 - set of all processes residing in main memory, ready and waiting to execute.
 - Device/Waiting Queues
 - set of processes waiting for an I/O device
- Processes *migrate* among the various queues.



Overview

- The Ready and wait Queues.





■ Overview

- Most processes can be described as either I/O-bound or CPU-bound:
 - I/O-bound process
 - spends more time doing I/O than computations
 - short CPU bursts
 - CPU-bound process
 - spends more time doing computations
 - long CPU bursts



■ Types of Process Schedulers

- There are three types/levels of Process Schedulers
 - Long-term Scheduler
 - High-level Scheduler, or
 - Jobs Scheduler.
 - Medium-term Scheduler
 - Medium-level Scheduler,
 - Swapping Scheduler, or
 - Emergency Scheduler.
 - Short-term Scheduler
 - Low-level Scheduler,
 - CPU Scheduler,
 - Micro Scheduler, or
 - Process/Thread Scheduler on narrow sense (狭义上的进程/线程调度).



■ Long-term Schedulers

- Long-term process scheduler selects which programs/processes should be brought into the *ready queue*.
 - determines which **programs** are admitted to the system for processing
 - controls the degree of multiprogramming
 - strives for good process mix of I/O-bound and CPU-bound processes
 - Long-term scheduler is invoked infrequently.
 - seconds, minutes
 - may be slow
- If more processes are admitted for processing:
 - less likely that all processes will be blocked
 - bringing better CPU usage
 - each process has less fraction of the CPU time



■ Short-term Schedulers

- Short-term process scheduler selects which process should be *executed* next and allocates CPU — also called CPU scheduling (处理机调度).
 - CPU scheduling determines which *process* is going to execute next according to a scheduling algorithm.
 - Short-term scheduler is also known as the *dispatcher* (which is part of it) that moves the processor from one process to another, and prevents a single process from monopolizing (独占) processor time.
 - Short-term scheduler is invoked on an event that may lead to choose another process for execution:
 - Clock interrupts
 - I/O interrupts
 - Operating system calls and traps
 - Signals.
 - Short-term scheduler is invoked very frequently
 - milliseconds
 - must be fast.



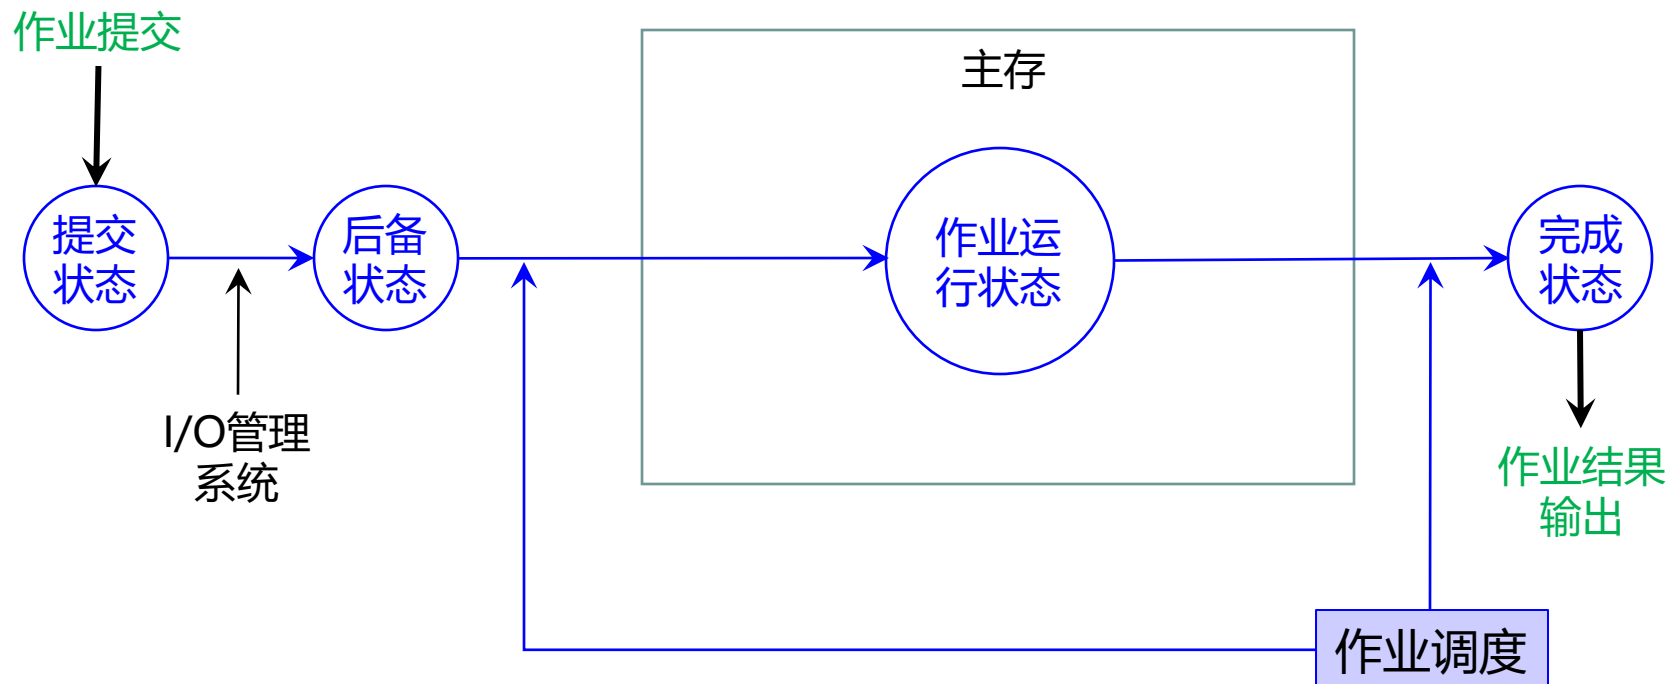
■ Medium-term Schedulers

- Medium-term process scheduler selects which job/process should be *swapped* out if system is overloaded.
 - So far, all processes have to be (at least partly) in main memory.
 - Even with virtual memory, keeping too many processes in main memory will deteriorate the system's performance.
 - OS may need to swap out some processes **to disk**, and then later swap them back in.
 - Swapping decisions is based on the need of multiprogramming management.



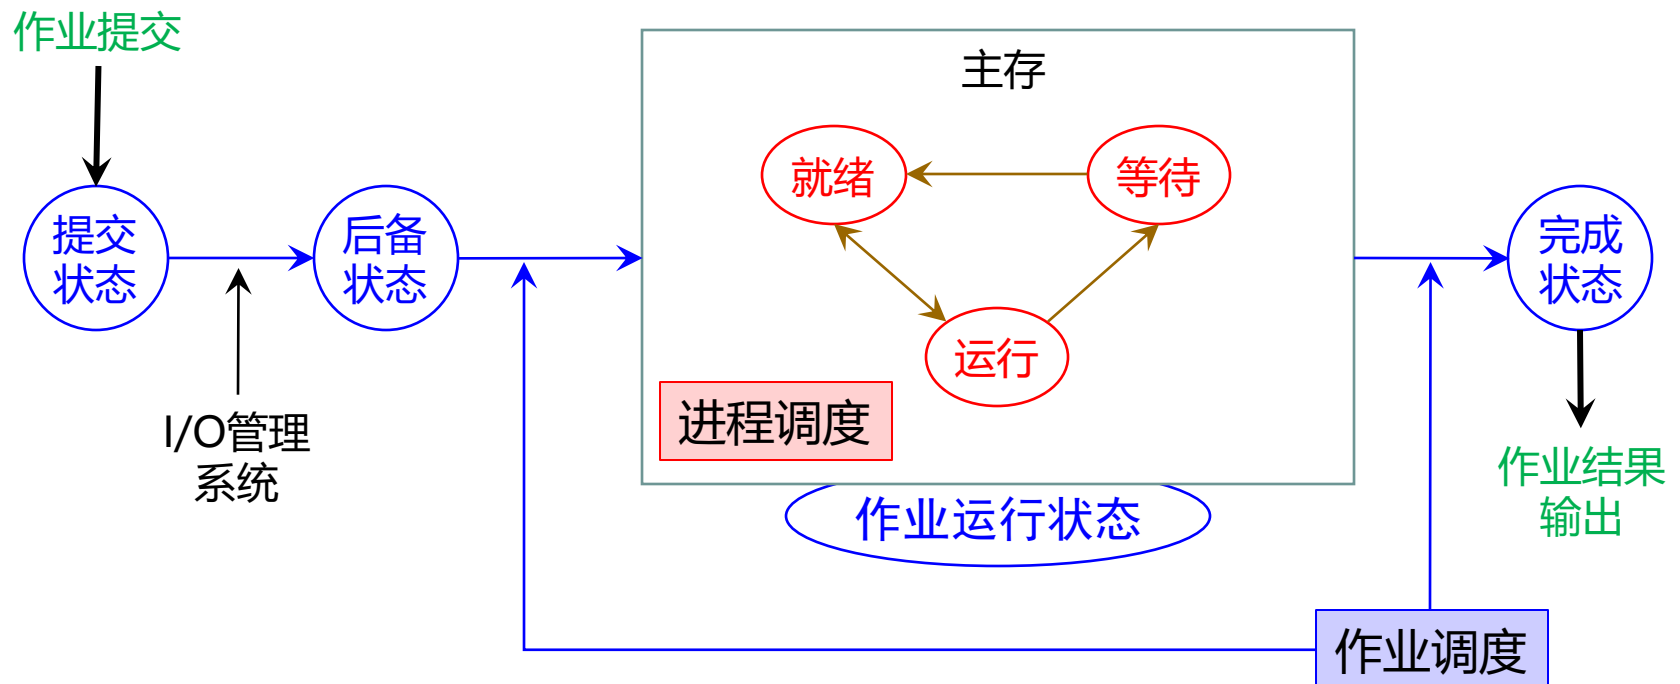
■ Schematic View of Schedulers

- Job scheduling, Swapping scheduling and Process/Thread scheduling.



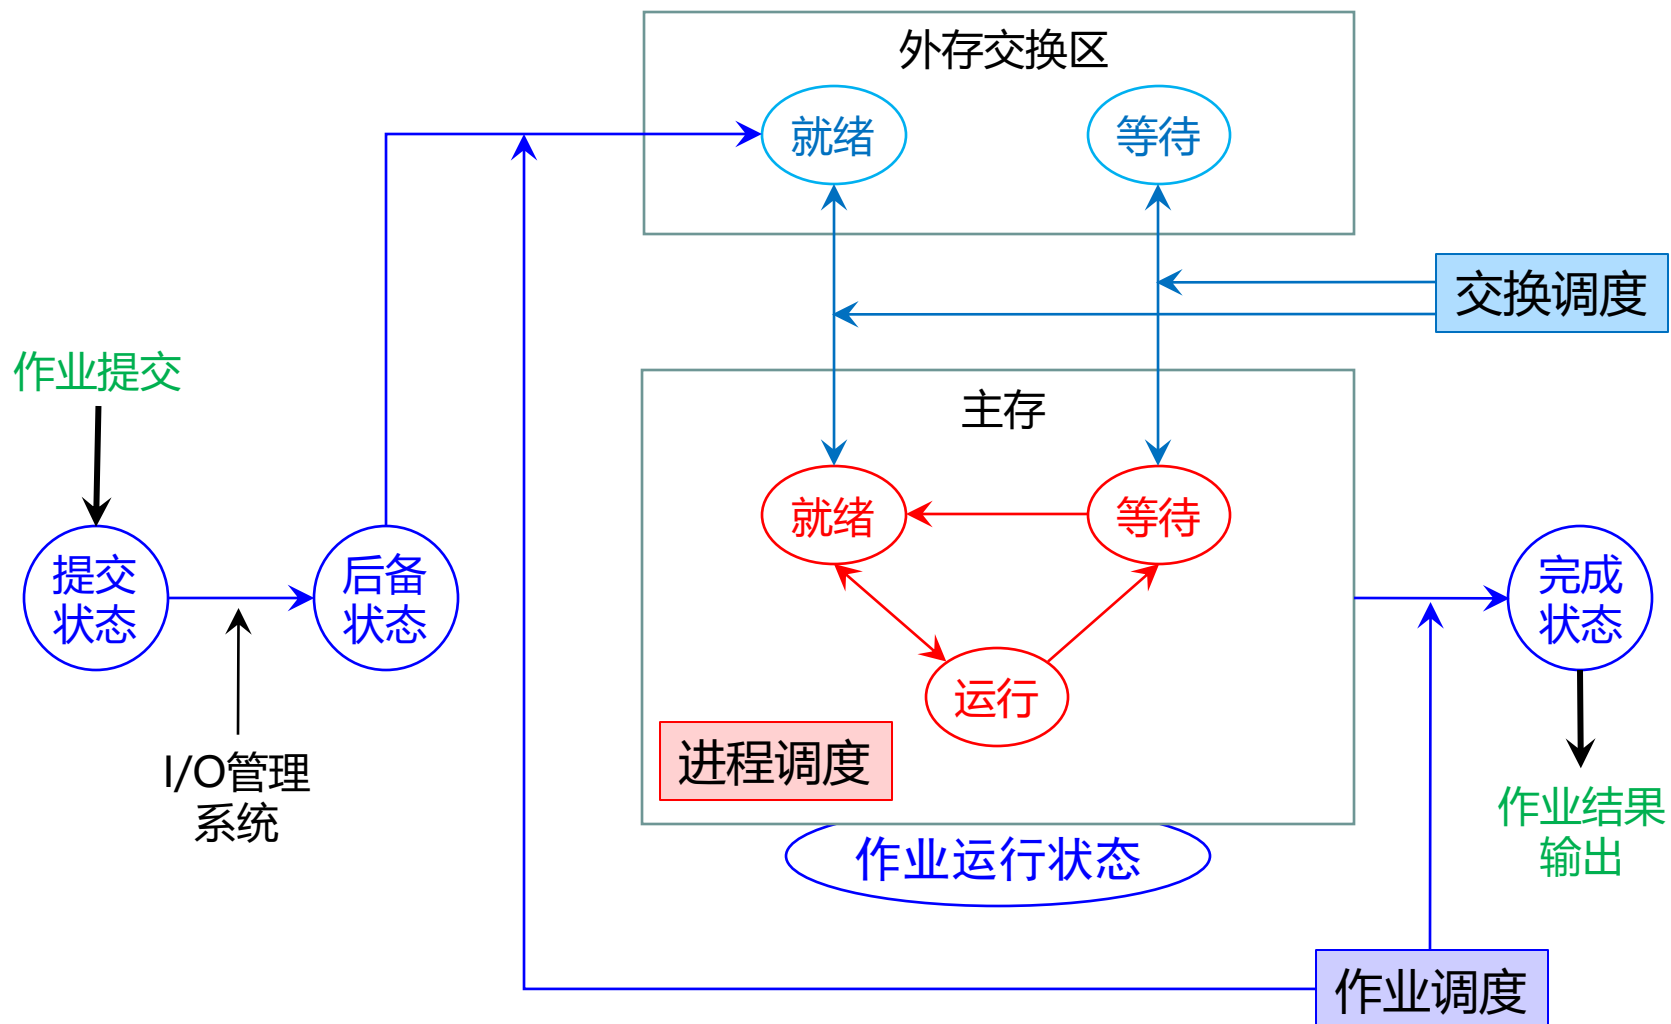
■ Schematic View of Schedulers

- Job scheduling, Swapping scheduling and Process/Thread scheduling.



■ Schematic View of Schedulers


- Job scheduling, Swapping scheduling and Process/Thread scheduling.





■ Process Swapping

■ Schematic View of Swapping.





■ Process Swapping


■ Dynamics of Swapping

- A process can be swapped temporarily out of memory to a backing storage, and then brought back into memory for continued execution.
- Backing storage
 - fast **disk** large enough to accommodate copies of all memory images for all users
 - must provide direct access to these memory images
- Roll out, roll in
 - swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped.
- Modified versions of swapping are found on many systems.
 - e.g., UNIX, Linux, and Windows.
- System maintains a ready queue of ready-to-run processes which have memory images on disk.



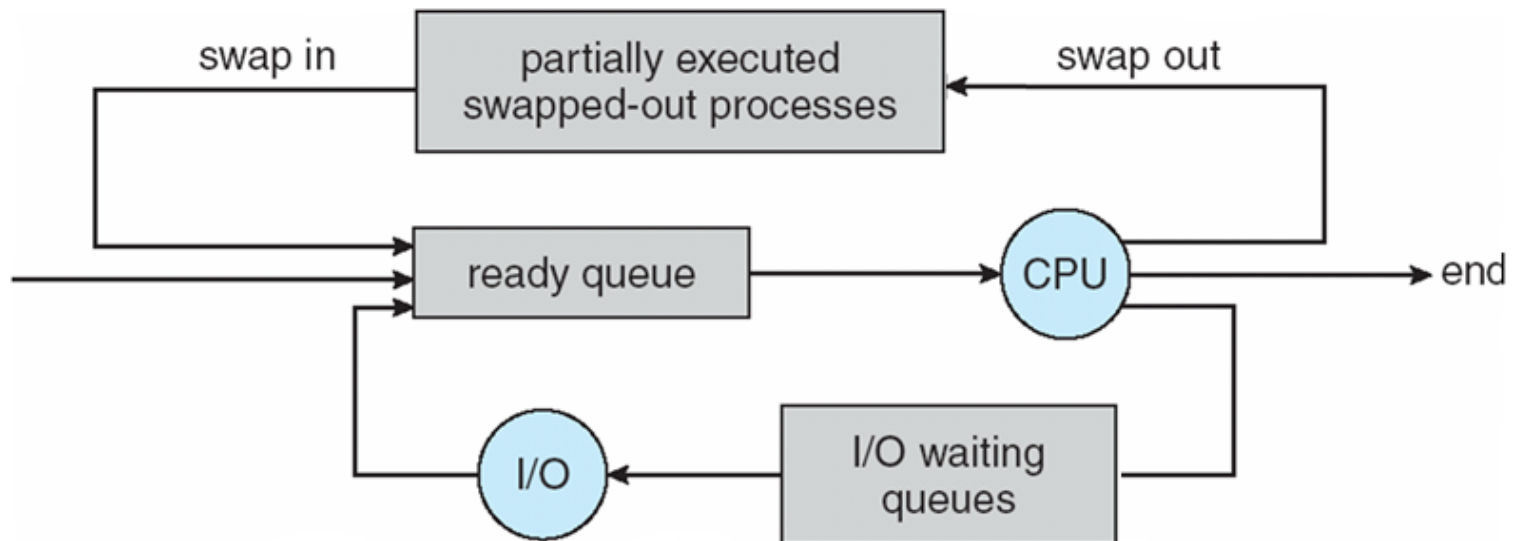
■ Process Swapping

■ Swapping Example.



■ Process Swapping

- Addition of Medium-term Scheduling.





■ Process Swapping

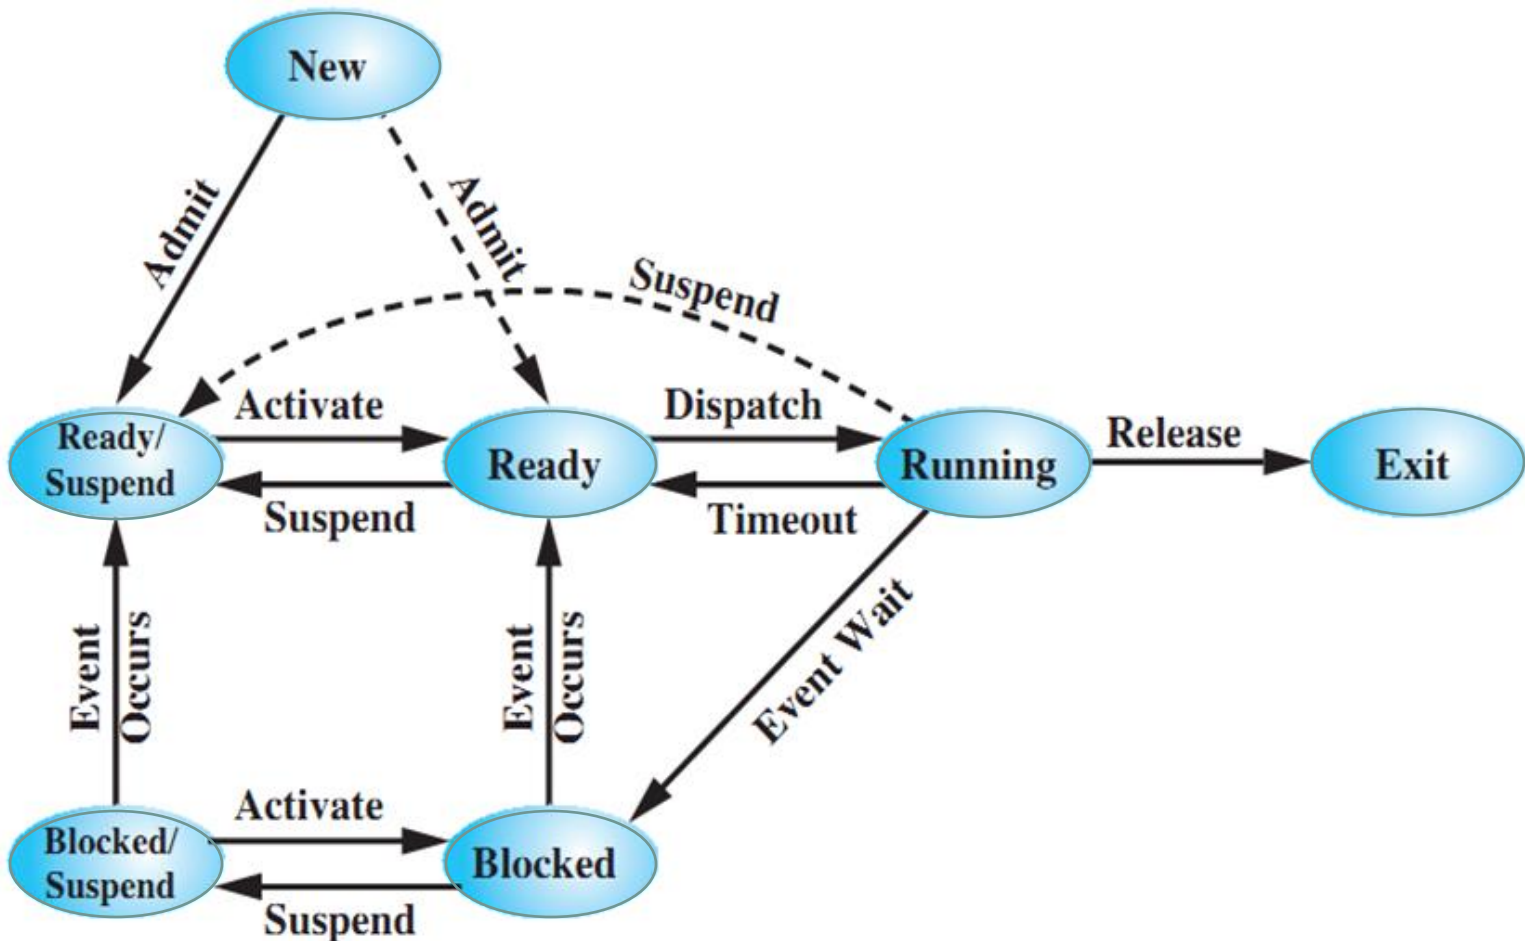
■ Support of Swapping

- OS may need to suspend some processes, i.e., to swap them out to disk and then swap them back in.
- Two new states might need to be added:
 - Blocked Suspend: blocked processes which have been swapped out to disk
 - Ready Suspend: ready processes which have been swapped out to disk



■ Process Scheduling Activities

- A Seven-state Process Model.





■ Process Scheduling Activities

■ A Seven-state Process Model

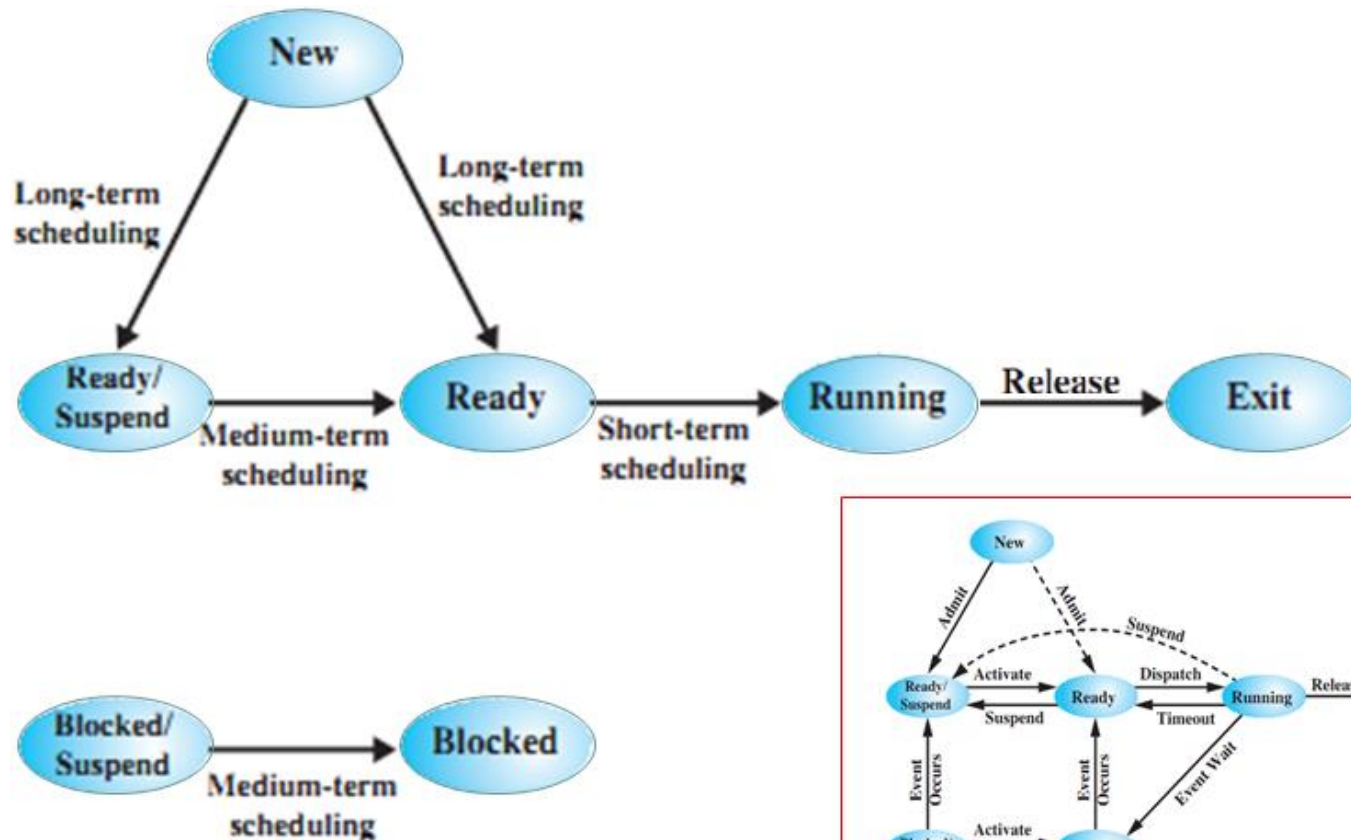
■ Additional state transitions

- Blocked → Blocked Suspend
 - When all processes are blocked, the OS will make room to bring a ready process in memory.
- Blocked Suspend → Ready Suspend
 - when the event for which it has been waiting occurs (state info is available to OS).
- Ready Suspend → Ready
 - when no more ready processes in main memory
- Ready → Ready Suspend (unlikely)
 - when there are no blocked processes and memory must to be freed for adequate performance



■ Process Scheduling Activities

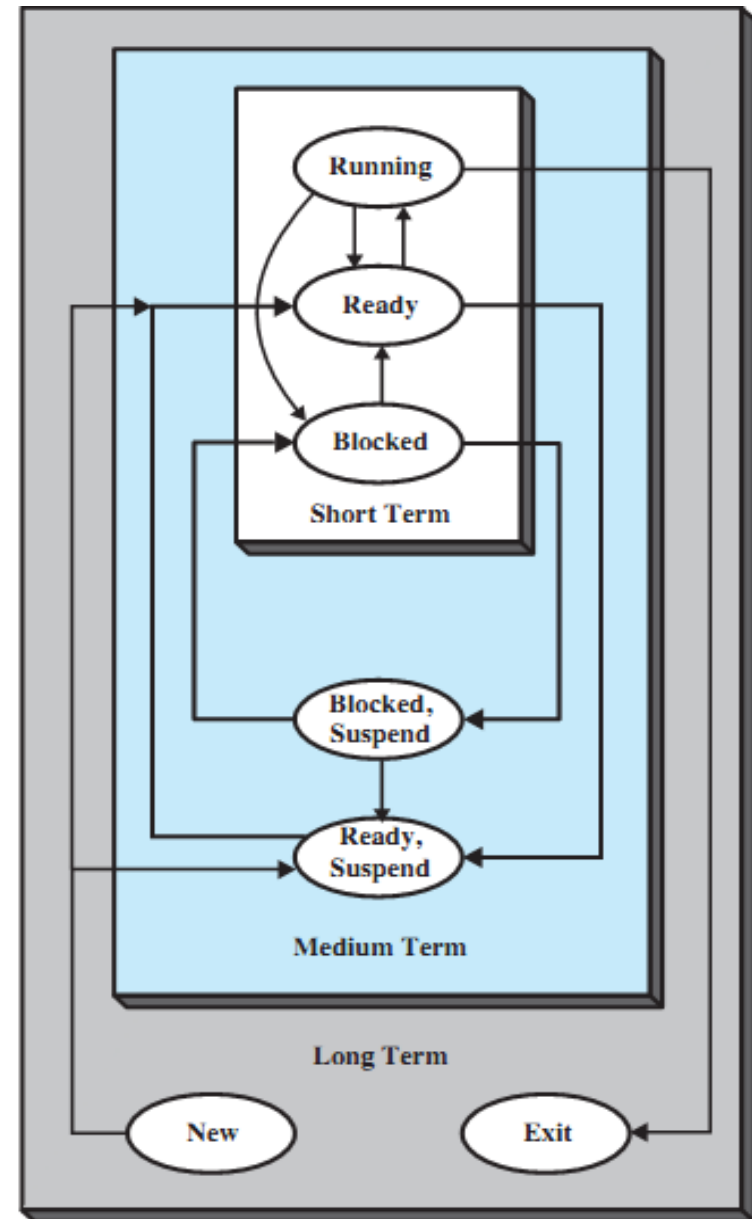
- A Seven-state Process Model
 - Scheduling Levels and Scheduling Activities.



■ Process Scheduling Activities

- A Seven-state Process Model.

- Another view of the three levels of scheduling.





■ Process Scheduling Queues

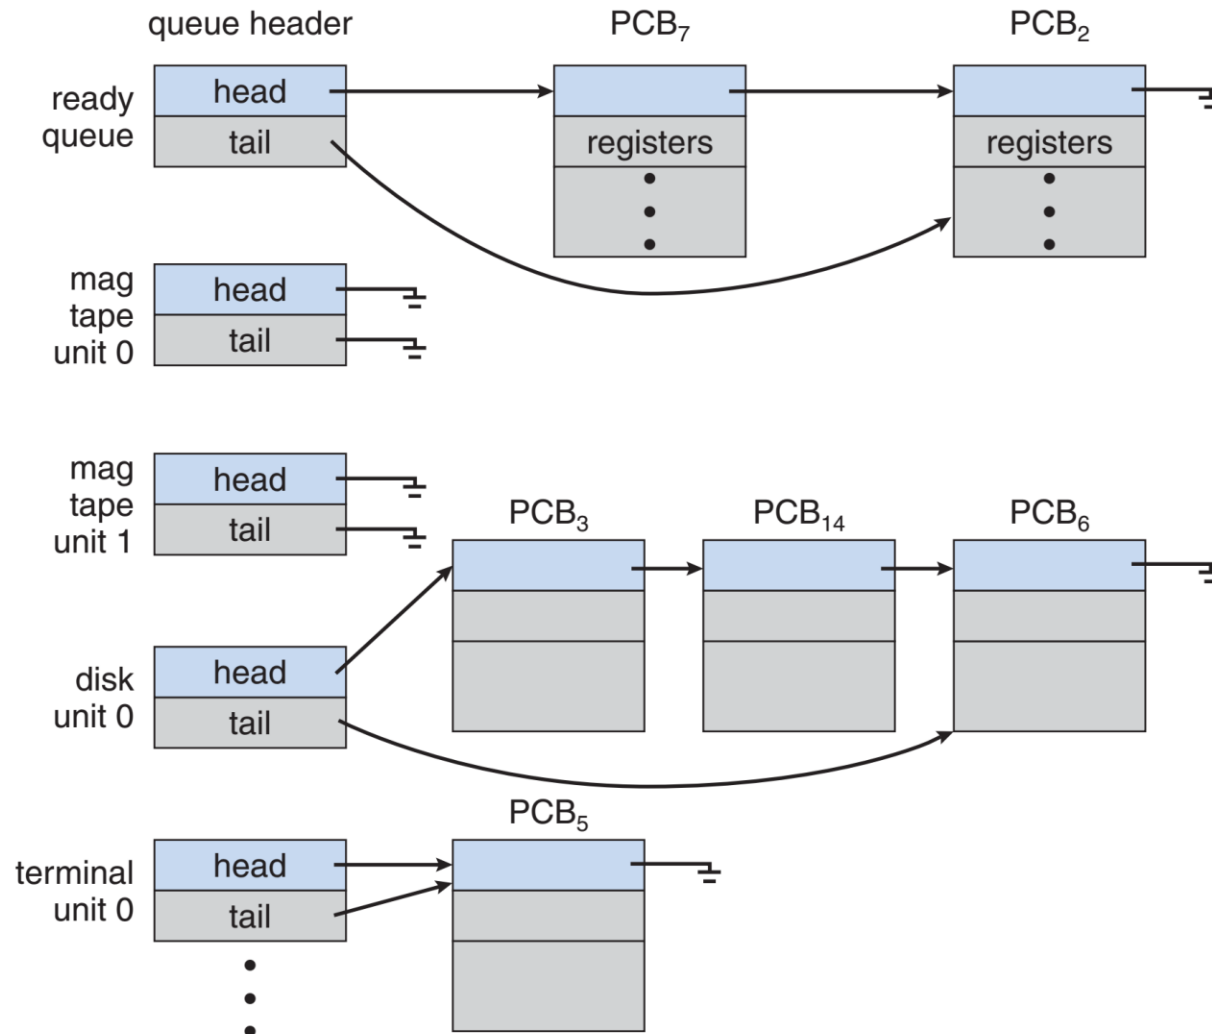
■ Job Queue, Ready Queue and Device Queue

- As processes enter the system, they are put into a *job queue*, which consists of all processes in the system.
- The processes that are residing in main memory and are ready and waiting to execute are kept on a list called the *ready queue*.
 - The ready queue is generally stored as a linked list. The header contains pointers to the first and final PCBs in the list. Each PCB includes a pointer field that points to the next PCB in the ready queue.
- The list of processes waiting for a particular I/O device is called a *device queue*. Each device has its own device queue.
 - Suppose A process makes an I/O request to a shared device such as a disk, but the disk may be busy with the I/O request of some other process. The process therefore may have to wait for the disk.
 - A device queue is not a queue of devices, but a queue of process waiting for a particular I/O device.



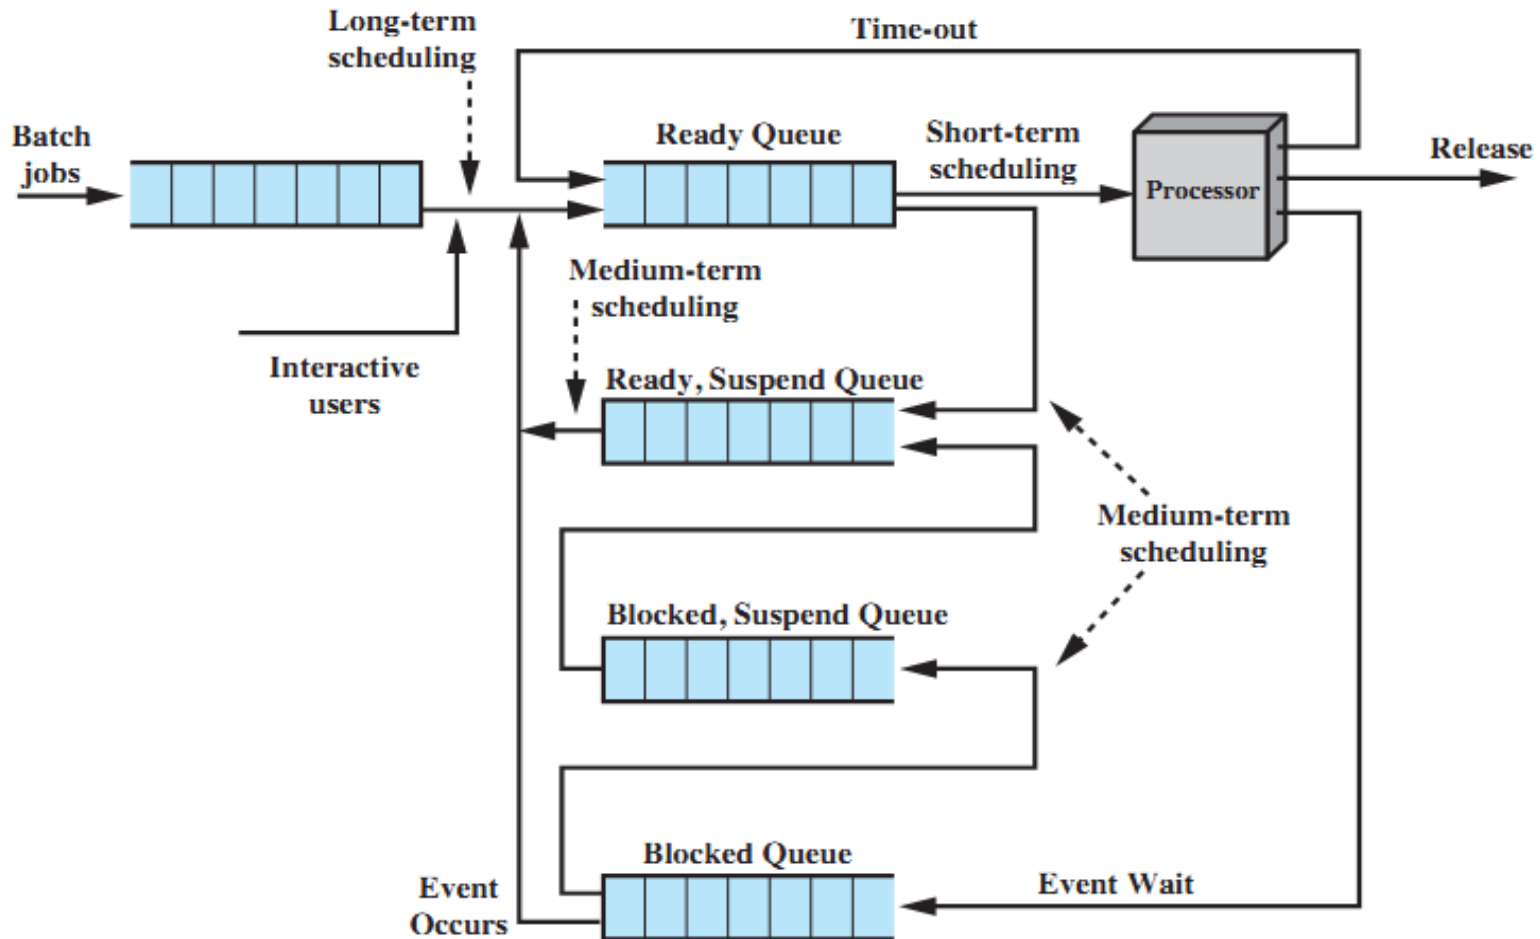
■ Process Scheduling Queues

- The ready queue and various I/O device queues.



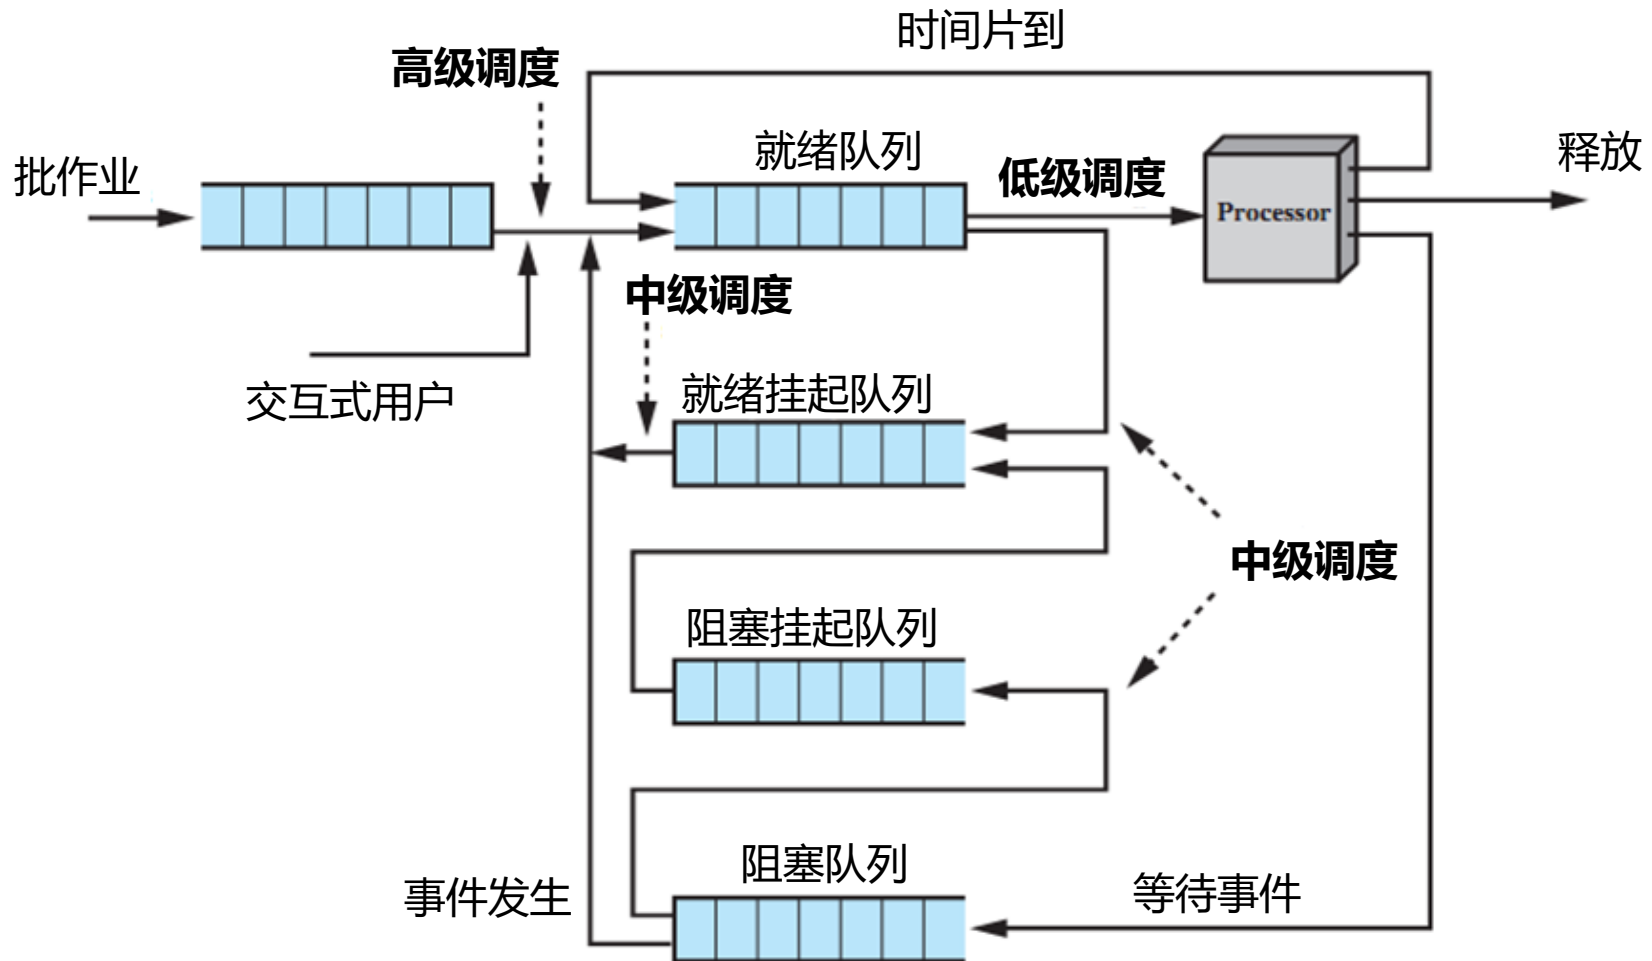
■ Process Scheduling Queues

- Queuing diagram for process scheduling
 - Processes migrate among the various queues.



Process Scheduling Queues

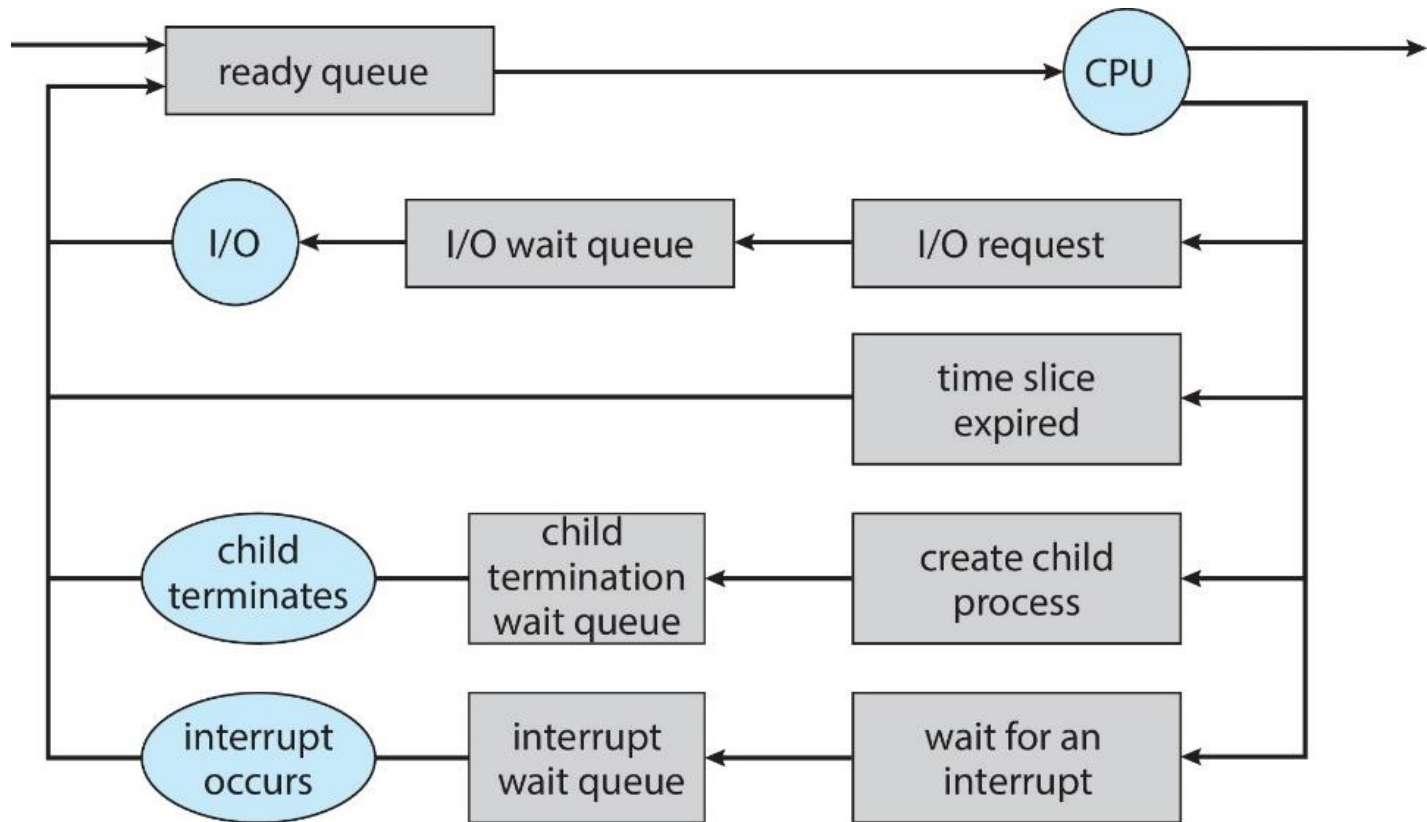
- Queuing diagram for process scheduling
 - Processes migrate among the various queues.





■ Process Scheduling Queues

- Queuing diagram for process scheduling.

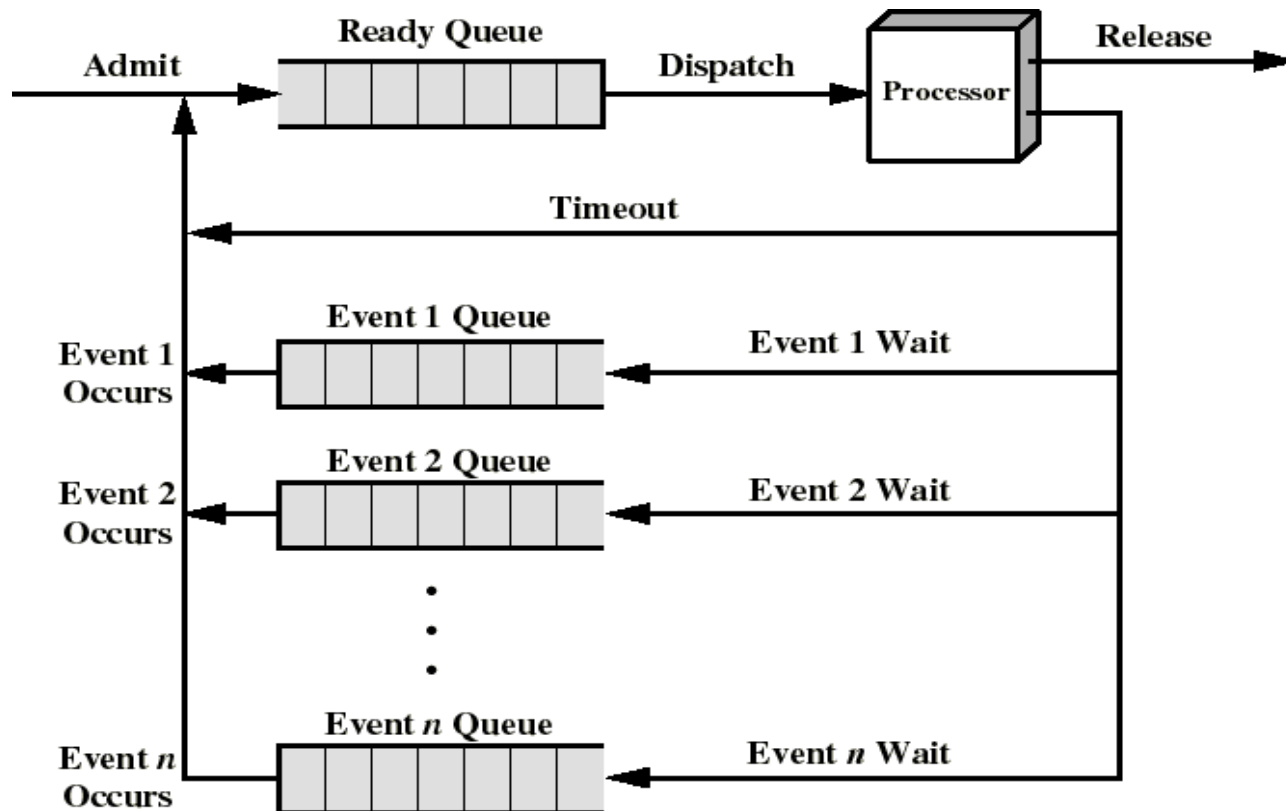




■ Process Scheduling Queues

■ A Queuing Discipline

- When some event occurs, the corresponding process is moved into the ready queue.






■ Process Switching


■ Multiprogramming Scheduling.

One program counter




(a)

Four program counters



(b)



(c)



■ Process Switching

- A process switch (进程切换) may occur whenever the OS has gained control of CPU. i.e., when:
 - Supervisor Call
 - explicit request by the program (example: file open)
 - The process will probably be blocked.
 - Trap
 - an error resulted from the last instruction
 - It may cause the process to be moved to terminated state.
 - Interrupt
 - the cause is external to the execution of the current instruction
 - Control is transferred to Interrupt Handler.



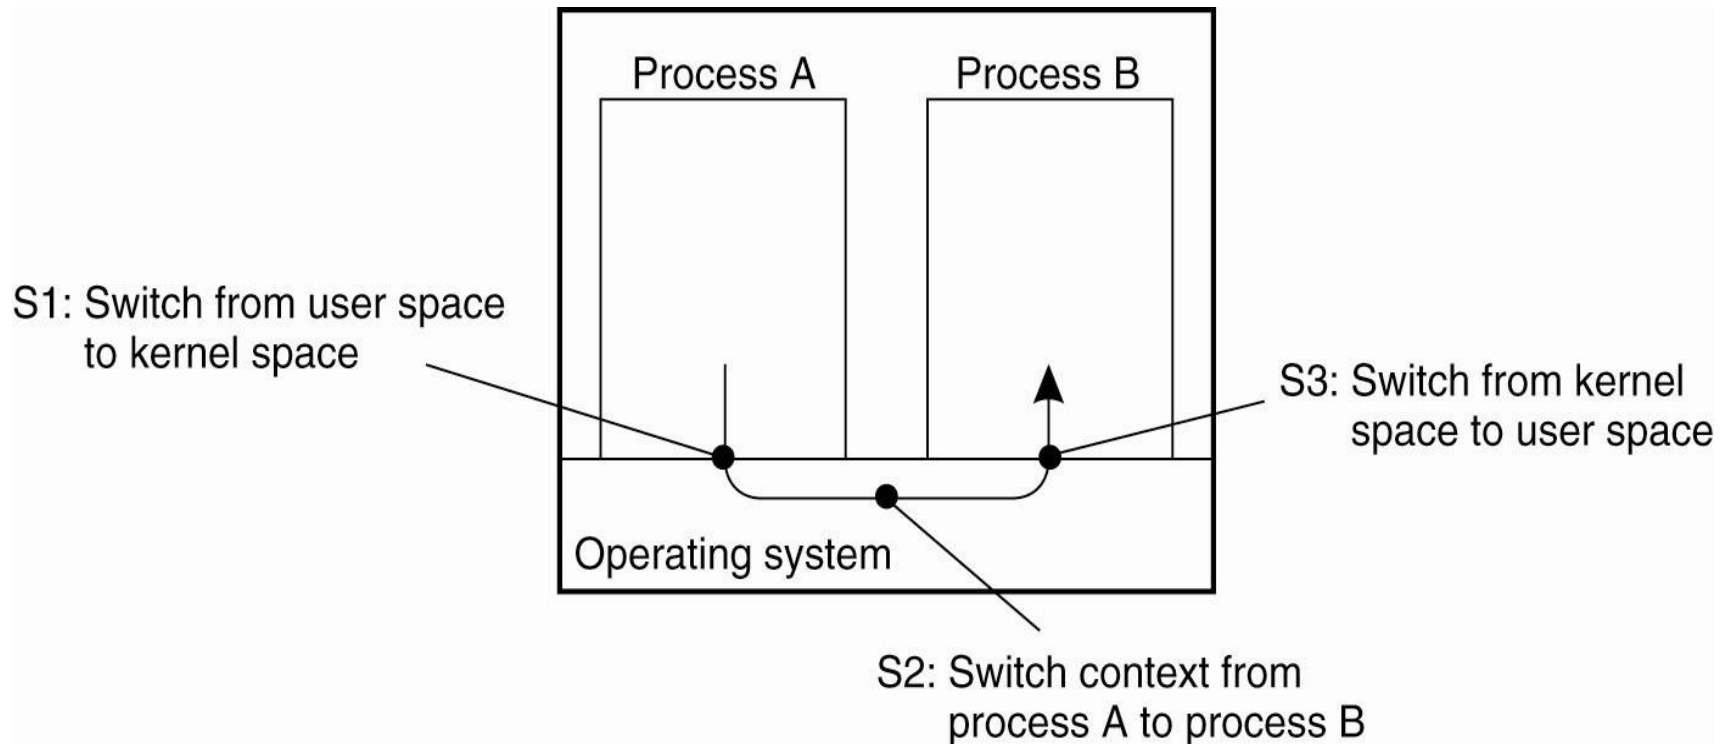
■ Context Switching

- When CPU switches to another process, the system must save the state of the rolling out process and load the saved state of the rolling in process.
 - This is called *Context Switch* (上下文切换).
 - Context of a process is represented in the PCB.
- The time it takes is dependent on hardware support.
- Context-switch time is overhead (开销).
 - The system does no useful work while switching.



■ Context Switching

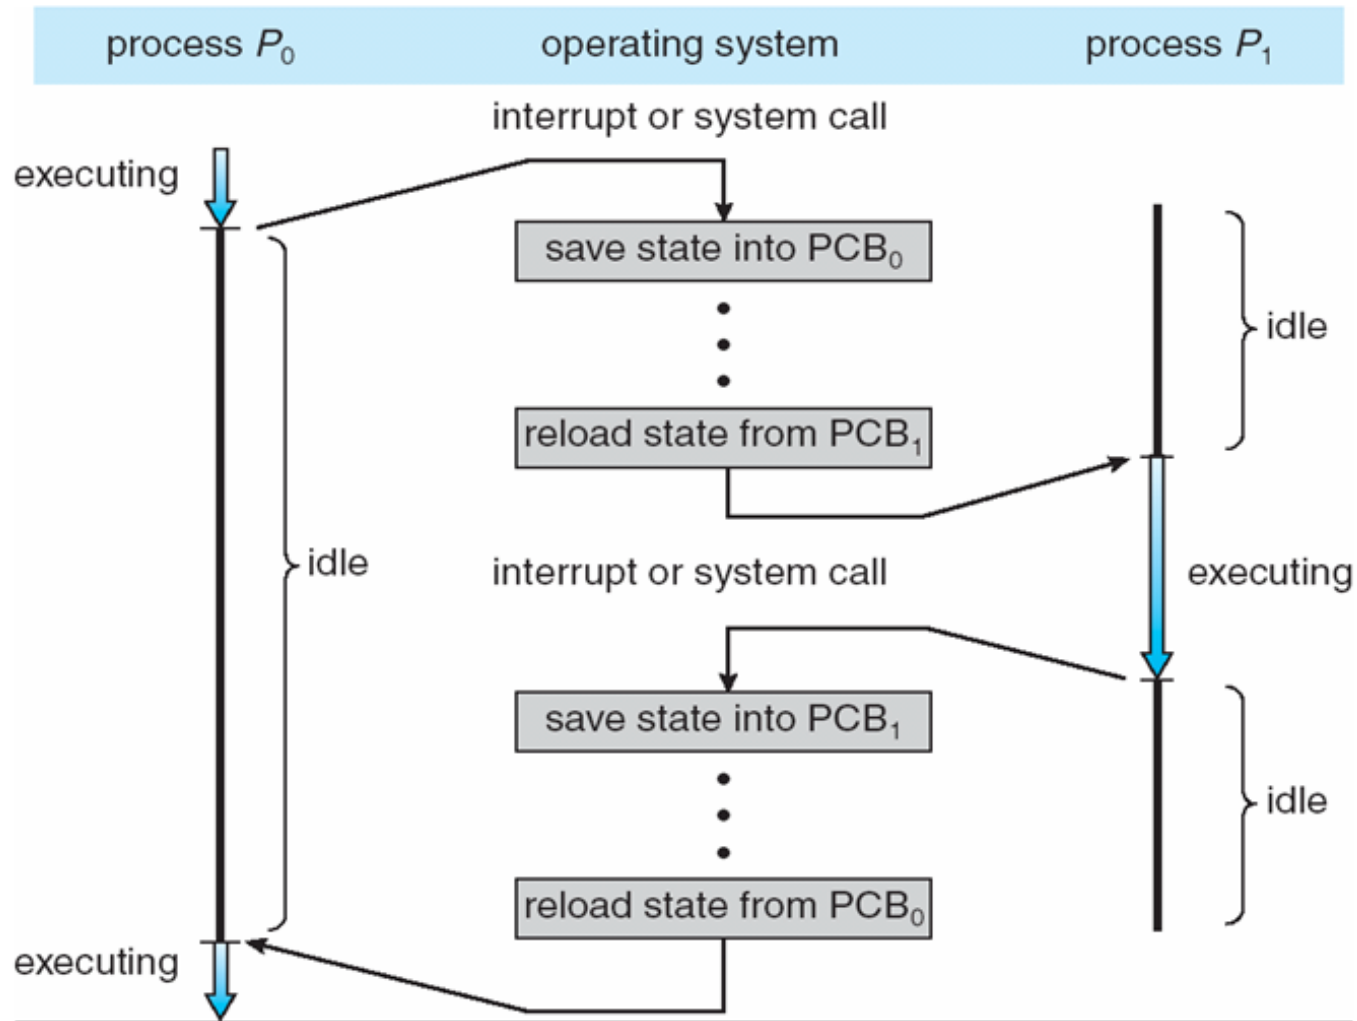
- Context Switch between Processes.





Context Switching

Context Switch between Processes.





■ Context Switching

■ Steps in Context Switch

- Save **CPU context** including program counter and other registers.
- Update the **PCB** of the running process with its new state and other associate information.
- Move PCB to an appropriate **queue** – ready, blocked,
- **Select** another process (for next execution).
- Update **PCB** of the selected process.
- Restore **CPU context** from that of the selected process.



■ Mode Switching

- An interrupt may not produce a context switch.
 - The control can just return to the interrupted program.
- Then only the processor state information needs to be saved on stack.
- This is called *Mode Switch* (模式切换)
 - user mode to kernel mode when going into Interrupt Handler.
- Less overhead
 - no need to update the PCB like for context switch.