
Introduction to Mem. Management

Operating Systems

School of Data & Computer Science
Sun Yat-sen University

Lecture Notes: os_sysu@163.com
Instructor: Guoyang Cai
email: isscg@mail.sysu.edu.cn





■ Contents

- Basic Concepts
- Memory Management Requirements
 - Relocation
 - Protection
 - Sharing
 - Logical Organization
 - Physical Organization
- Memory Partitioning
 - Fixed Partitioning
 - Variable Partitioning
 - Buddy System

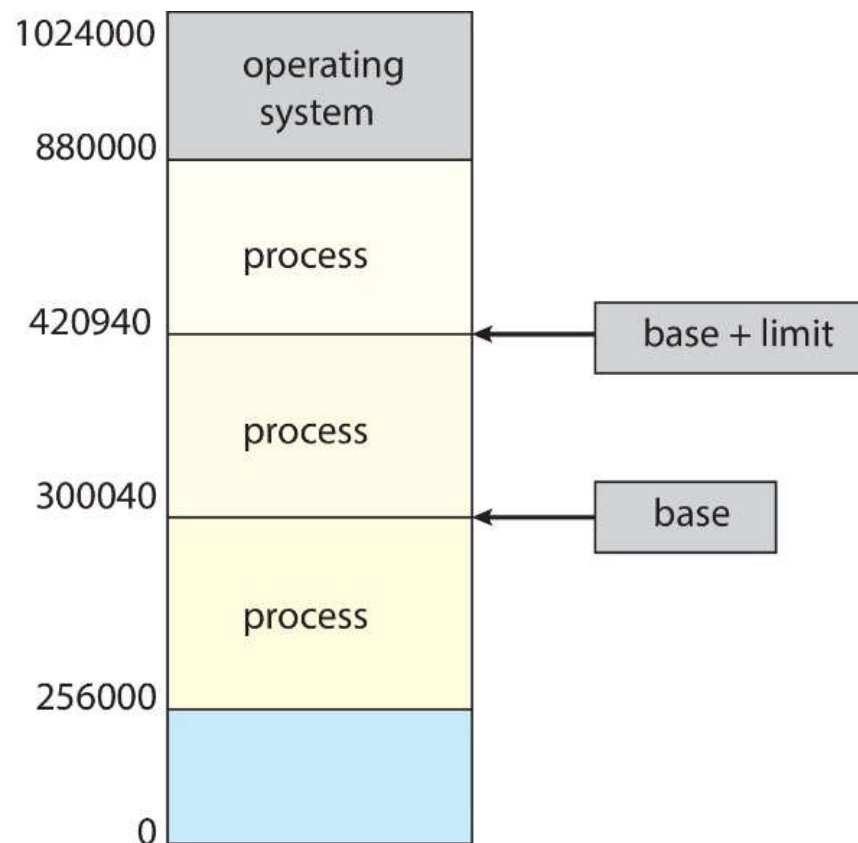


■ Background

- User programs must be brought (from disk) into memory and placed within a process for it to be run.
 - **Memory Unit** can sense only a stream of <addresses, read> requests, or <address, data, write> requests.
 - Main memory and registers are only storage that CPU can access directly.
 - Registers are accessed in one CPU clock (or less)
 - Main memory access can take many cycles, causing a stall.
 - Cache sits between main memory and CPU registers.
- Memory management is the task carried out by OS and hardware to accommodate multiple processes in main memory.
 - Protection of memory is required to ensure correct operation.

■ Hardware Address Protection

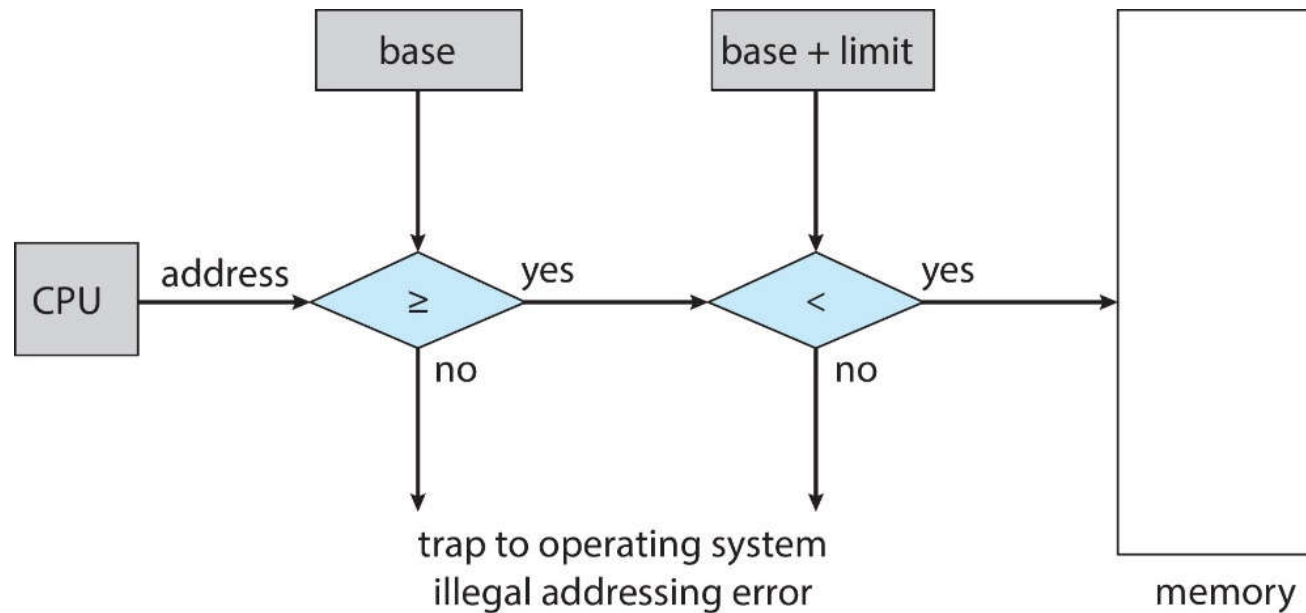
- A pair of *base register* and *limit register* define the logical address space for a process.





■ Hardware Address Protection

- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user.

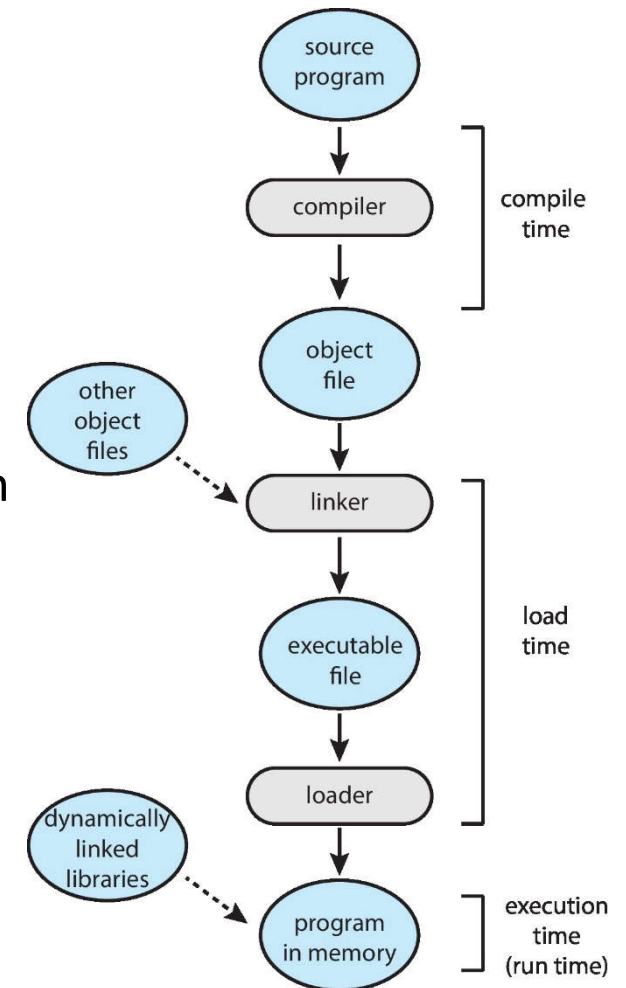


■ Address Binding

- Programs on disk form an *input queue* when they are ready to be brought into memory to execute.
 - How to allocate the first physical address of a user process? Always at 0...00?
- Further more, addresses are represented in different ways at different stages of a program's life.
 - Source code addresses are usually symbolic.
 - Compiled code addresses are bound to relocatable addresses.
 - e.g., “14 bytes from beginning of this module”.
 - Linker or loader will bind relocatable addresses to absolute addresses.
 - e.g., $14 + 71000 = 71014$.
 - Each binding maps one address space to another.

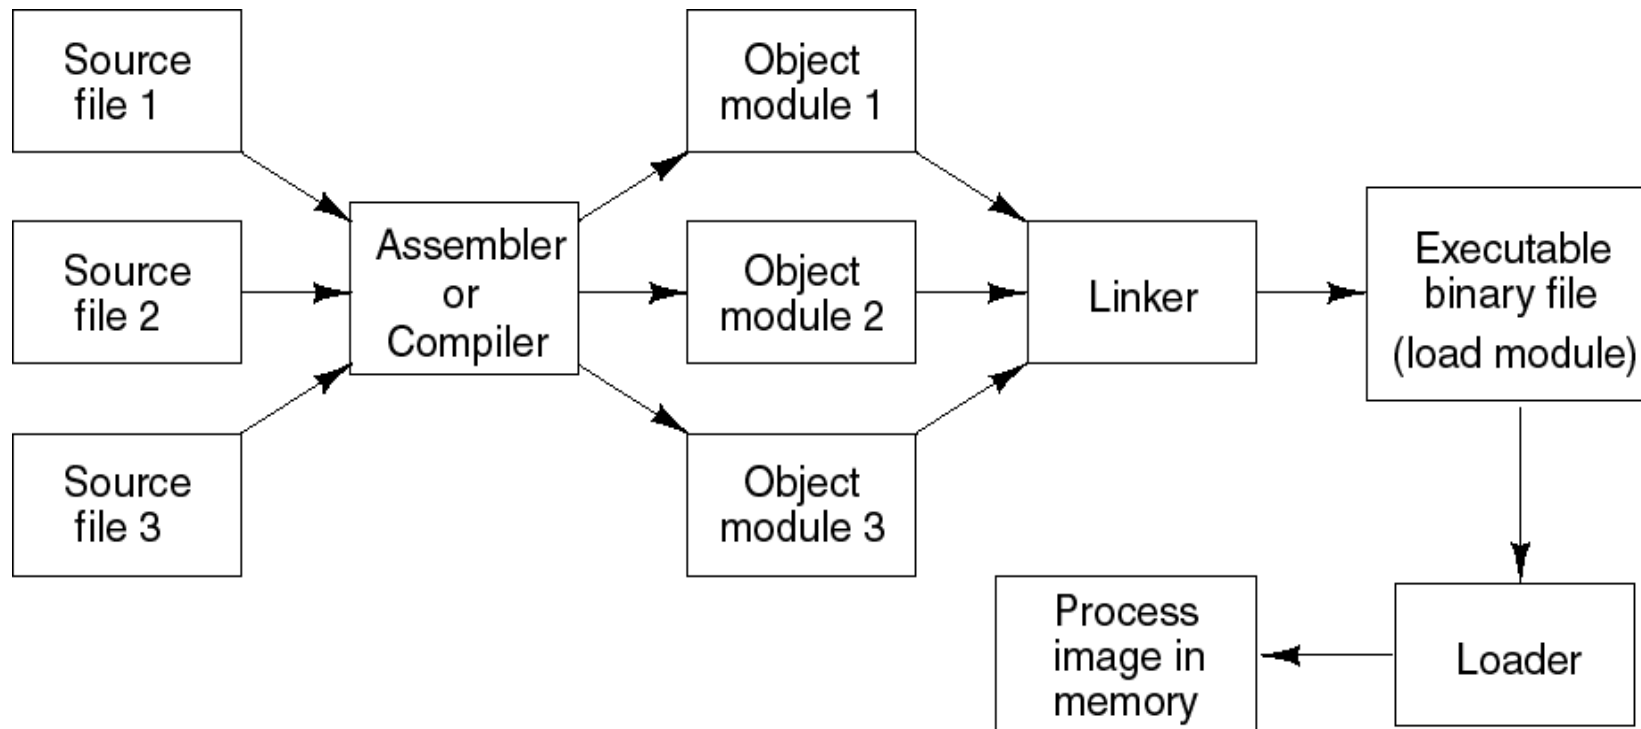
■ Address Binding

- Address binding of instructions and data to memory addresses can happen at three different stages of *multistep Processing of a User Program*.
 - **Compile time:** If memory location known a priori, *absolute codes* can be generated; recompiling needed if starting location changes.
 - **Load time:** If memory location is not known at compile time, *relocatable codes* are generated at load time.
 - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another.
 - Need hardware support for address maps (e.g., base and limit registers).



■ Address Binding

- Multi-step Processing of User Program.





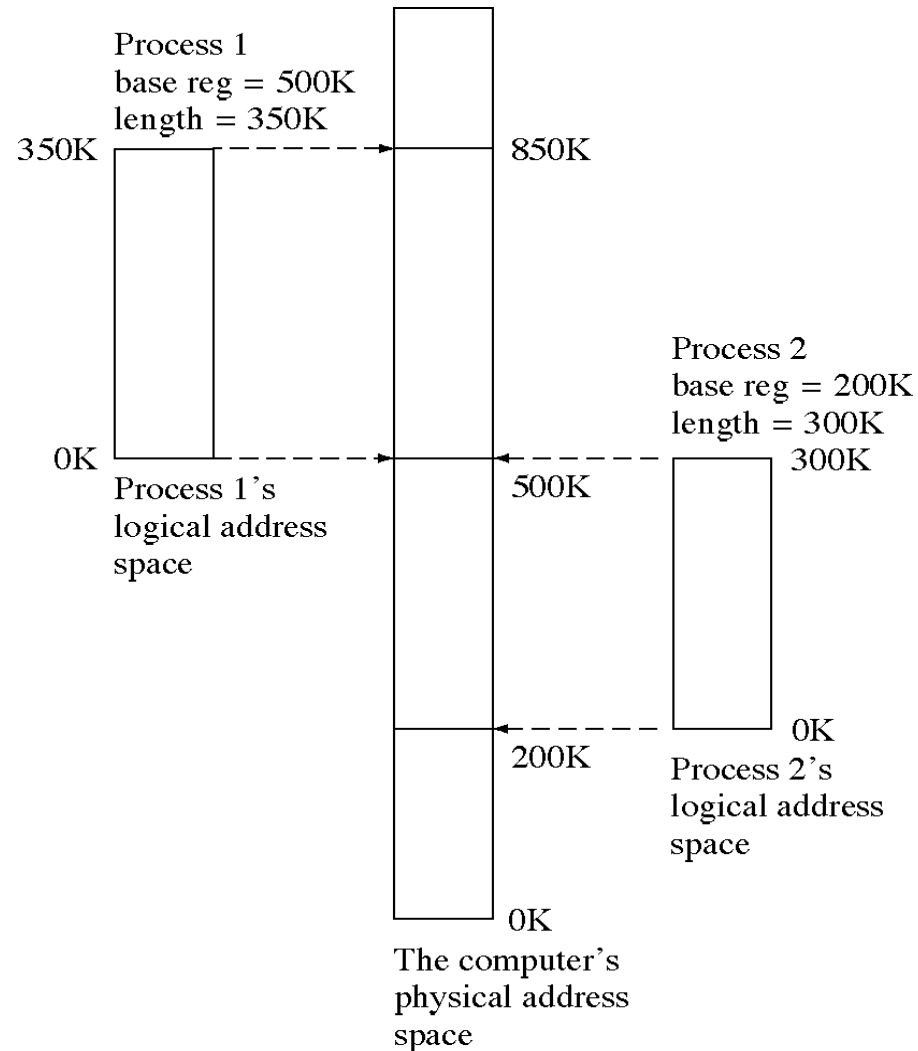
■ Logical Address Space vs. Physical Address Space

- The concept of a logical address space that is bound to a separate *physical address space* is central to proper memory management.
 - *Logical address* – generated by the CPU; also referred to as *virtual address*, is a reference to a memory location that is independent of the physical organization of memory. *Logical address space* is the set of all logical addresses generated by a program.
 - *Physical address* – absolute address seen by the memory unit, is a physical location in main memory. *Physical address space* is the set of all physical addresses generated by a program.
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme.
 - Compilers produce code in which all memory references are logical addresses.
 - A relative address is an example of logical address in which the address is expressed as a location relative to some known point in the program (e.g., the beginning).



■ Logical Address Space vs. Physical Address Space

- Logical and Physical address Spaces.






■ Memory-Management Unit (MMU)

- Memory-Management Unit (MMU) is a **hardware device** that maps virtual addresses to physical addresses **at run time**.
- In MMU scheme, the value in the base register is added to every logical (virtual) address generated by a user process at the time it is sent to memory.
 - The base register is now called the **relocation register** (重定位寄存器).
 - MS-DOS on Intel 80x86 used 4 relocation registers.
- The user program deals with *logical (virtual)* addresses; it never sees the *real (physical)* addresses.
 - Execution-time binding occurs when reference is made to location in memory.
 - Logical addresses are bound to physical addresses.



■ Memory-Management Unit (MMU)


- CPU, MMU and Memory.





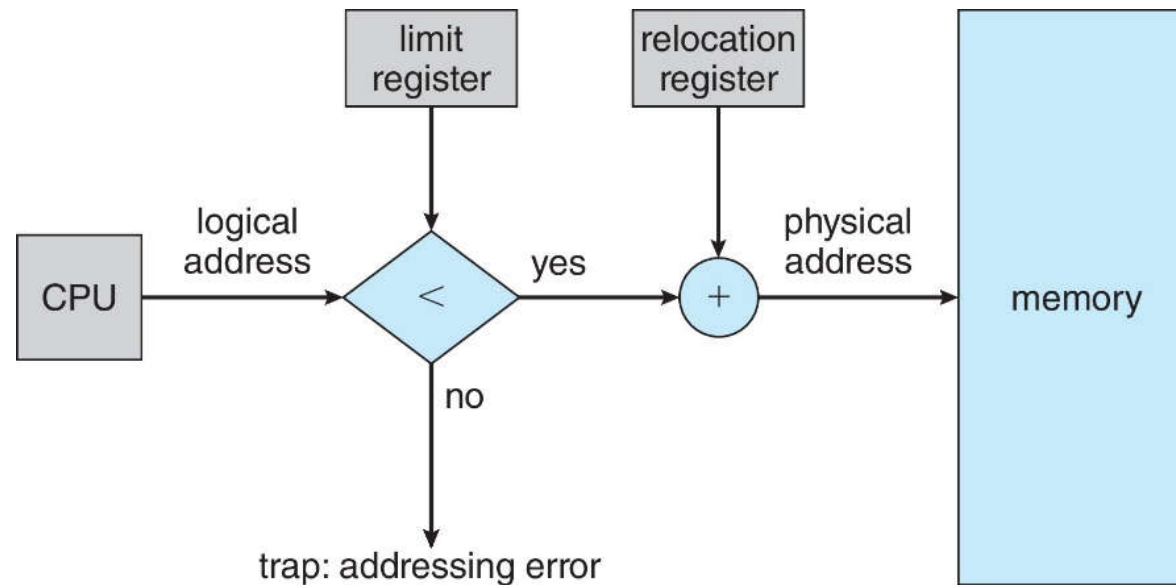
■ Dynamic Relocation Using a Relocation Register

- A routine is not loaded until it is called.
 - All routines kept on disk in relocatable load format.
 - unused routine never loaded - Better memory-space utilization
 - useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required.
 - implemented through program design
 - OS can help by providing libraries to implement dynamic loading.



■ Hardware Support for Relocation and Limit Registers

- The base register is now called the relocation register.





■ Dynamics of Hardware Translation of Addresses

- When a process is assigned to the running state, a relocation/base register gets loaded with the starting physical address of the process.
- A limit/bounds register gets loaded with the process's ending physical address.
- When a relative addresses is encountered, it is added with the content of the base register to obtain the physical address which is compared with the content of the limit/bounds register.
- This provides hardware protection: each process can only access memory within its process image.



■ Dynamic Linking

- *Static linking* – system libraries and program code combined by the **loader** into the binary program image.
- *Dynamic linking* – linking postponed until execution time.
 - Small piece of code, *stub*, used to locate the appropriate memory-resident library routine.
 - Stub replaces itself with the address of the routine, and executes the routine.
 - **Operating system** checks if routine is in processes' memory address.
 - if not, add it to address space
 - Dynamic linking is particularly useful for shared/common libraries – here full OS support is needed.
 - like standard C language library
- Consider applicability to patching system libraries
 - Versioning may be needed.



■ Swapping

- A process can be *swapped* temporarily out of memory to a backing store, and then brought back into memory for continued execution.
 - With support of swapping, total memory space of processes can exceed the real physical memory size.
- *Backing store* – fast disk large which is enough to accommodate copies of all memory images for all users; must provide direct access to these memory images.
- *Roll out, roll in* – swapping variant which is used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed.
- Major part of swap time is transfer time.
 - Total transfer time is directly proportional to the amount of memory swapped.
- System maintains a *ready queue* of ready-to-run processes which have memory images on disk.




■ Swapping

- Does the swapped out process need to swap back in **to the same** physical addresses as before?
 - It depends on address binding method.
 - considering pending I/O to or from process memory space
- Modified versions of swapping are found on many systems (e.g., UNIX, Linux, and Windows).
 - Swapping normally disabled
 - Started if more than threshold amount of memory allocated
 - Disabled again once memory demand reduced below threshold



■ Swapping

- Schematic View of Swapping.





■ Swapping

- Context Switch Time including Swapping.
 - If the next process (the target process) to be put on CPU is not in memory, it needs to swap out some process and swap in the target process.
 - Context switch time can then be very high.
 - Consider 100MB process swapping to hard disk with transfer rate of 50MB/sec:
 - Swap out time of $100/50 \text{ (sec)} = 2000 \text{ (ms)}$
 - Plus swap in of same sized process
 - Total context switch swapping component time of 4000ms (4 seconds)
 - Such swapping time can be reduced if we reduce the size of memory swapped – by knowing how much memory really being used.
 - System calls to inform OS of memory use via
 - `request_memory()`
 - `release_memory()`



■ Swapping

- Context Switch Time including Swapping.
 - Other constraints as well on swapping
 - Pending I/O – can't swap out as I/O would occur to wrong process.
 - Or always transfer I/O to kernel space, then to I/O device
 - known as *double buffering*, adds overhead
 - Standard swapping is not used in modern operating systems.
 - But modified version is in common use.
 - swap only when free memory extremely low



■ Swapping

■ Swapping on Mobile Systems

■ Not typically supported

● Flash memory based

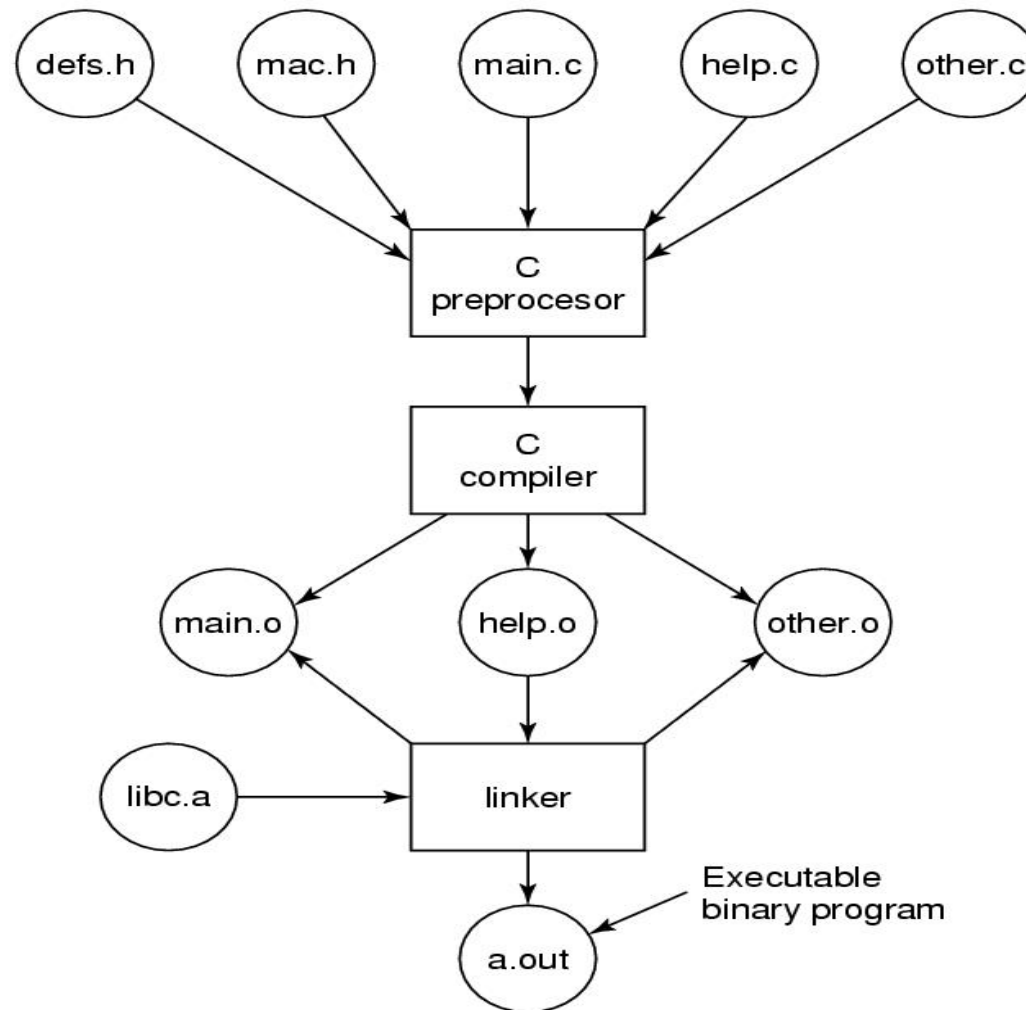
- Small amount of space
- Limited number of write cycles
- Poor throughput between flash memory and CPU on mobile platform

■ Instead use other methods to free memory if low.

- iOS **asks** apps to voluntarily relinquish allocated memory
 - Read-only data thrown out and reloaded from flash if needed
 - Failure to free can result in termination
- Android terminates apps if low free memory, but first writes *application state* to flash for fast restart.
- Both OSes support paging as discussed later.



■ A C Compilation Example





■ A C Compilation Example

- Public names are usable by other object modules.
- External names are defined in other object modules.
 - includes the list of instructions having these names as operands
- Relocation dictionary
 - has the list of instructions who's operands are addresses (since they are relocatable)
- Only code and data will be loaded in physical memory.
 - The rest is used by the linker and then removed.
- The stack is allocated only at load time.

End of module
Relocation dictionary
Data
Machine code
External names table
Public names table
Module identification



■ A C Compilation Example

- Initially, each object module has its own address space.
- All addresses are relative to the beginning of the module.

Object module A

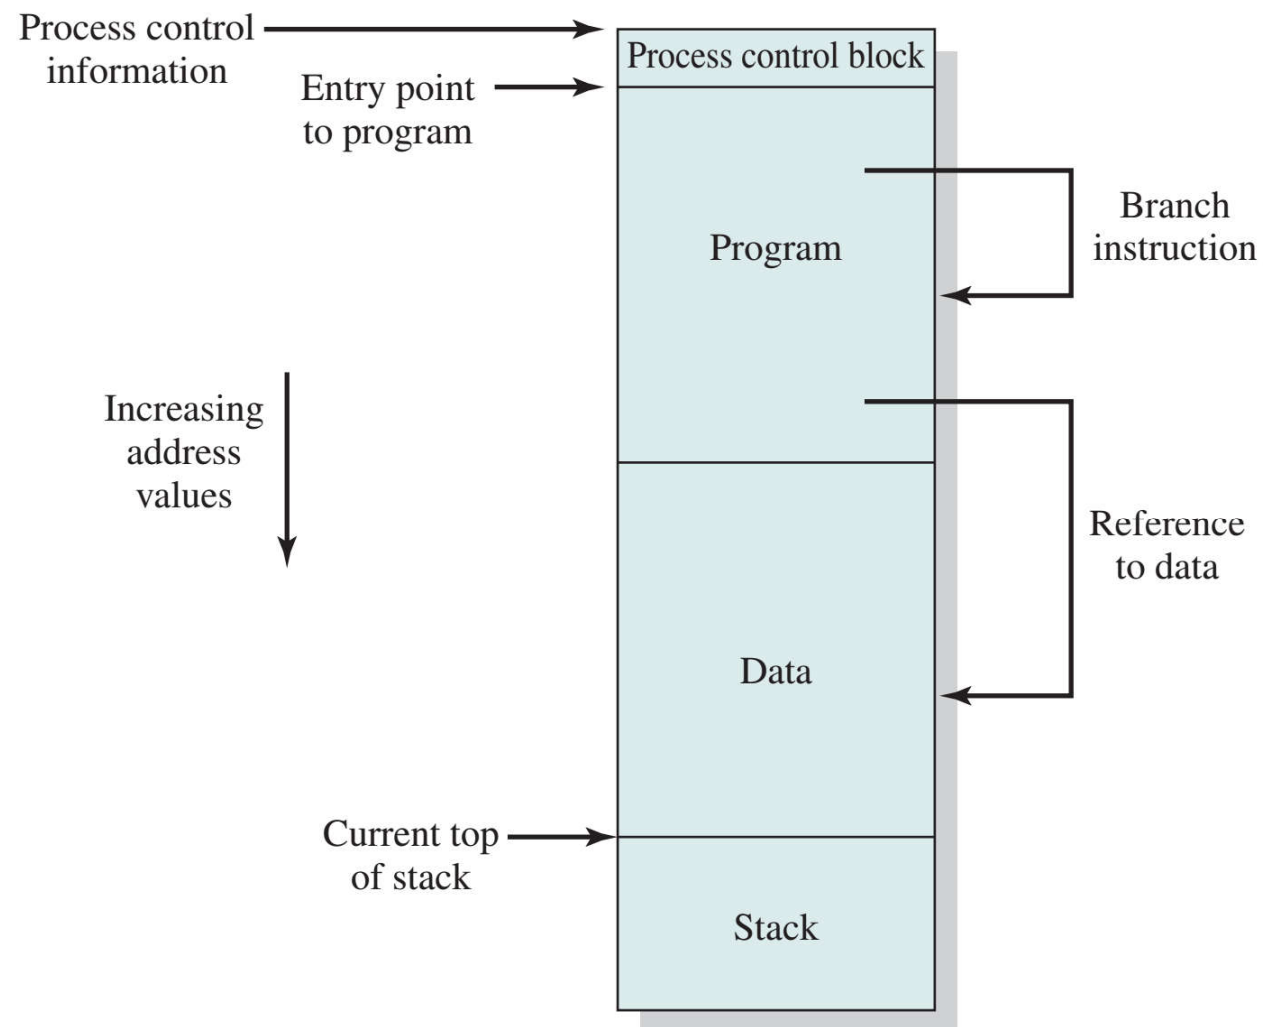
400	
300	CALL B
200	MOVE P TO X
100	
0	BRANCH TO 200

Object module B

500	
400	CALL A
300	
200	MOVE R TO X
100	
0	BRANCH TO 200

■ A C Compilation Example

■ Addressing Requirements for Process.

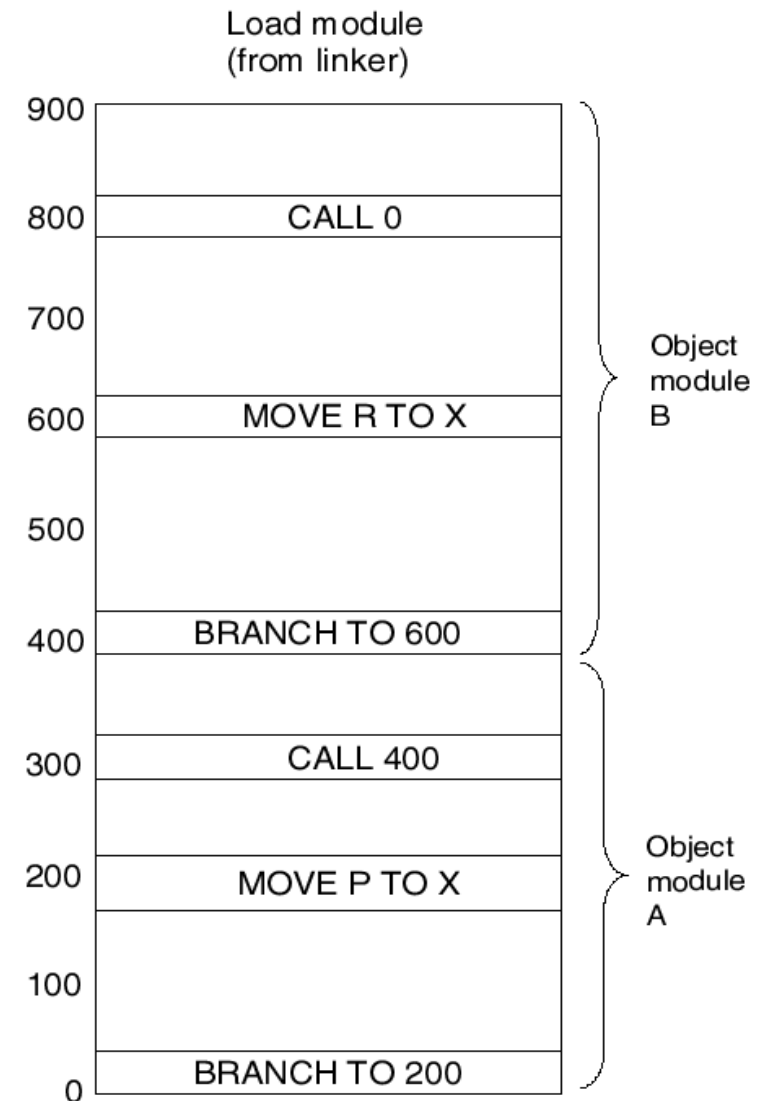




■ A C Compilation Example

■ Static Linking

- The linker uses tables in object modules to link modules into a single linear addressable space.
- The new addresses are addresses relative to the beginning of the load module.





■ A C Compilation Example

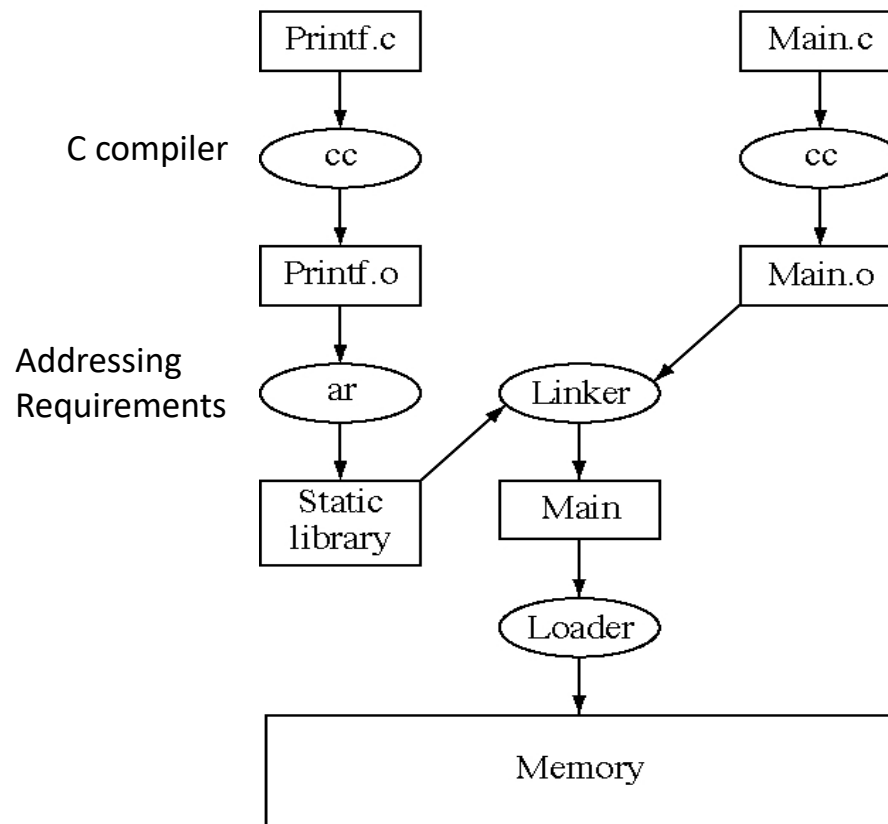
■ Dynamic Linking

- The linking of some external modules is done after the creation of the load module (executable file).
- Load-time dynamic linking:
 - The load module contains references to external modules which are resolved at load time.
- Run-time dynamic linking:
 - References to external modules are resolved when a call is made to a procedure defined in the external module.
 - Unused procedure is never loaded.
 - Process starts faster.




■ A C Compilation Example

■ Static Linking vs. Dynamic Linking



(a) Static Linking



(b) Dynamic Linking



■ Program Size vs. Memory Size

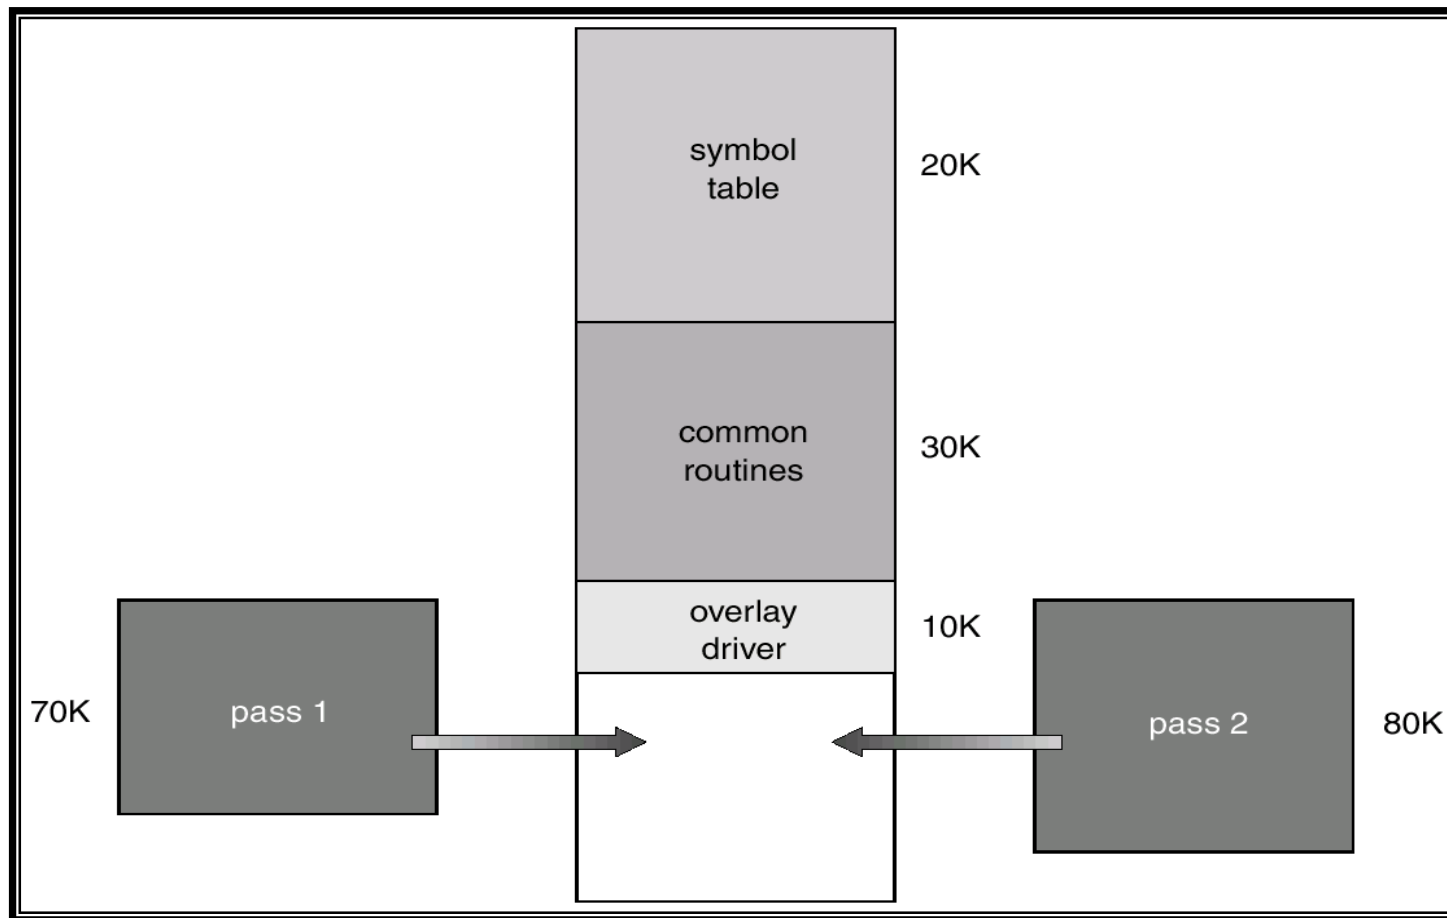
- What to do when program size is larger than the amount of memory or partition (that exists or can be) allocated to it?
- There are two basic solutions within real memory management:
 - Overlays
 - Dynamic Linking (Libraries – DLLs)
- Overlays
 - keep in memory only the part of the program whose instructions and data are needed *at any given phase/time*.
 - Overlays can be used only for programs that fit this model, e.g., multi-pass programs (多道程序) like compilers.
 - Overlays are designed/implemented *by programmer*.
 - Overlays need an overlay driver.
 - No special support needed from operating system
 - But program design of overlays structure is complex.

西復
羊西



■ Program Size vs. Memory Size

- Overlays.



Overlays for a Two-Pass Assembler



■ Program Size vs. Memory Size

■ Dynamic Linking

- Dynamic linking is useful when large amounts of code are needed to handle infrequently occurring cases.
 - routine not loaded unless/until it is called.
- Better memory-space utilization
 - unused routine never loaded
- Not much support from OS is required
 - It is implemented through program design.



■ Program Size vs. Memory Size

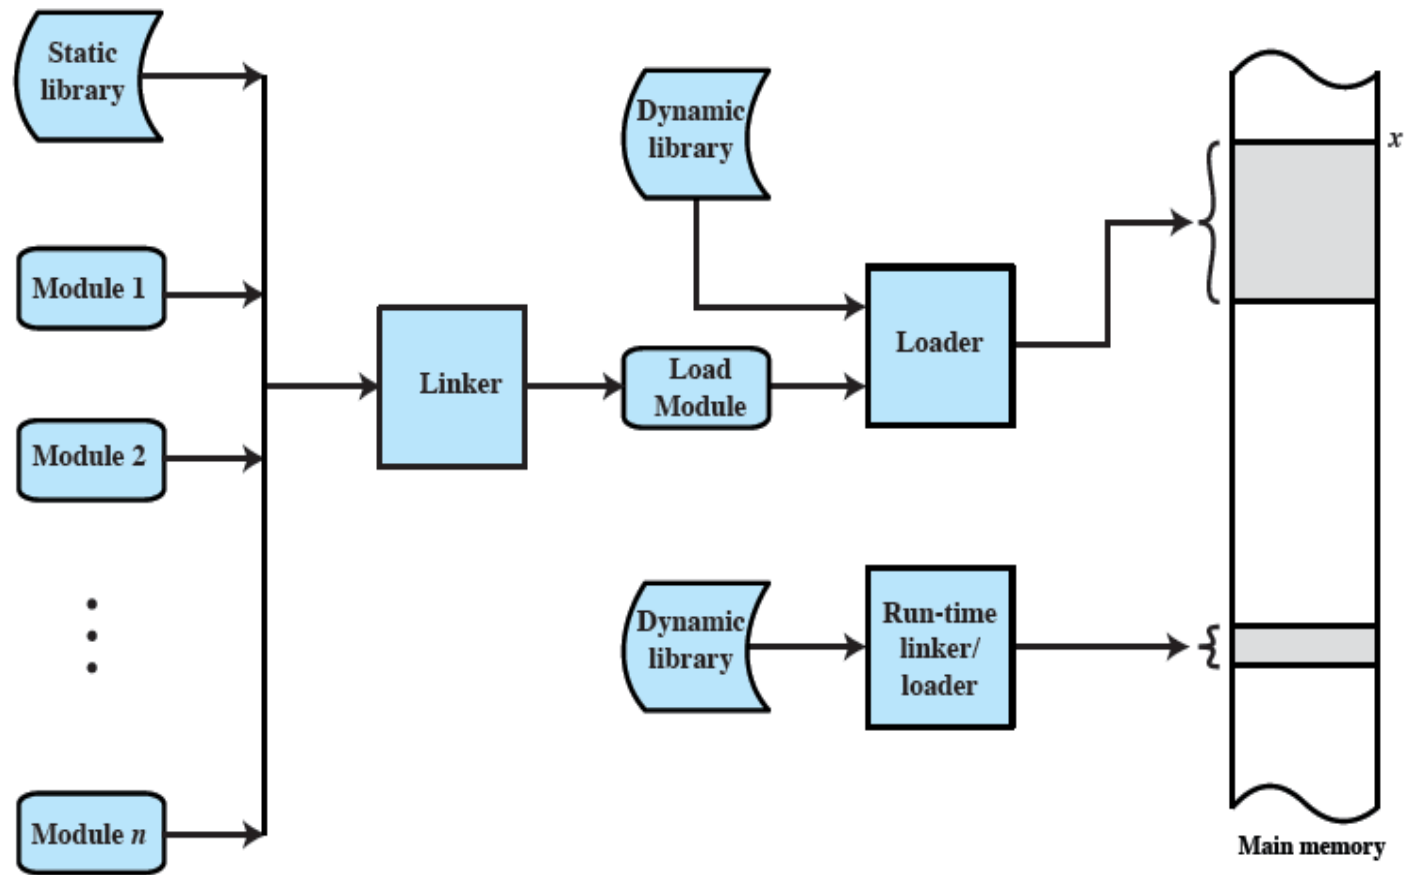
■ Advantages of Dynamic Linking

- Executable files, without the need of being modified, can use another version of the external module.
- Multiple processes are linked to the same external module.
 - saves disk space
- The same external module needs to be loaded in main memory only once.
 - Processes can share code and save memory.
- Examples:
 - Windows: .DLL files
 - Unix: .SO files (shared library)



■ Program Size vs. Memory Size

- Dynamic Linking/Loading Scenario.





■ **Memory Management Requirements**

- Memory needs to be allocated efficiently in order to pack as many processes into memory as possible to avoid the situation that all processes are waiting for I/O and the CPU is idle.
- It needs additional support for:
 - Relocation
 - Protection
 - Sharing
 - Logical Organization
 - Physical Organization



■ Memory Management Requirements

■ Relocation

- Programmer cannot know where the program will be placed in memory when it is executed.
- A process may be (often) relocated in main memory due to swapping/compaction:
 - Swapping enables the OS to have a larger pool of ready-to-execute processes.
 - Compaction (合并, 压紧) enables the OS to have a larger contiguous memory to place programs in.

■ Protection

- Processes should not be able to reference memory locations in another process without permission.
- It is impossible to check addresses in programs at compile/load-time since the program could be relocated.
- Address references must be checked at execution-time by hardware.



■ Memory Management Requirements

■ Sharing

- It must be allowed for several processes to access a common portion of main memory without compromising protection.
 - It is better to allow each process to access the same copy of the program rather than have their own separate copy.
 - Cooperating processes may need to share access to the same data structure.

■ Logical Organization

- Users write programs in modules with different characteristics.
 - Instruction modules are execute-only.
 - Data modules are either read-only or read/write.
 - Some modules are private and others are public.
- To effectively deal with user programs, OS and hardware should support a basic form of a module to provide the required protection and sharing.



■ Memory Management Requirements

■ Physical Organization

- External memory is the long term store for programs and data while main memory holds programs and data currently in use.
- Moving information between these two levels of the memory hierarchy is a major concern of memory management.
 - It is highly inefficient to leave this responsibility to the application programmer.

■ Contiguous Allocation

- In contiguous allocation (连续存储分配) an executing process must be loaded *entirely* in main memory (if overlays are not used).
- Main memory is usually split into two (Memory split) or more (Memory division) partitions:
 - resident operating system, usually held in low memory partition with interrupt vector
 - user processes then held in high memory partitions
- Relocation registers are used to protect user processes from each other, and from changing OS code and data.
 - Base register contains value of smallest physical address.
 - Limit register contains range of logical addresses – each logical address must be less than the limit register.
 - MMU maps logical address *dynamically*.


■ Real Memory Management Techniques

- Although the following simple (basic) memory management techniques are not used in modern operating systems, they lay the ground for a later proper discussion of virtual memory.
 - Fixed (Static) Partitioning (固定/静态分区分配)
 - Variable (Dynamic) Partitioning (可变/动态分区分配)
 - Simple (Basic) Paging (简单页式分配)
 - Simple (Basic) Segmentation (简单段式分配)




■ Fixed Partitioning 固定分区

- In fixed partitioning scheme main memory is divided into a set of *non-overlapping* memory regions called *partitions*.
- Fixed partitions (aka static partitions) can be of equal or unequal sizes.
- Leftover space in partition, after program assignment, is called *internal fragmentation* (内部碎片).




■ Placement Algorithm with Partitions

- Unequal-size partitions, use of multiple queues:
 - assign each process to the smallest partition within which it will fit
 - a queue exists for each partition size
 - tries to minimize internal fragmentation
 - problem: some queues might be empty while some might be loaded.



■ Placement Algorithm with Partitions

- Unequal-size partitions, use of a single queue:
 - When its time to load a process into memory, the smallest available partition that will hold the process is selected.
 - increasing the level of multiprogramming at the expense of internal fragmentation





■ Dynamics of Fixed Partitioning

- Any process whose size is less than or equal to a partition size can be loaded into the partition.
- If all partitions are occupied, OS can swap a process out of a partition.
- A program may be too large to fit in a partition. In this case the programmer must design the program with overlays.

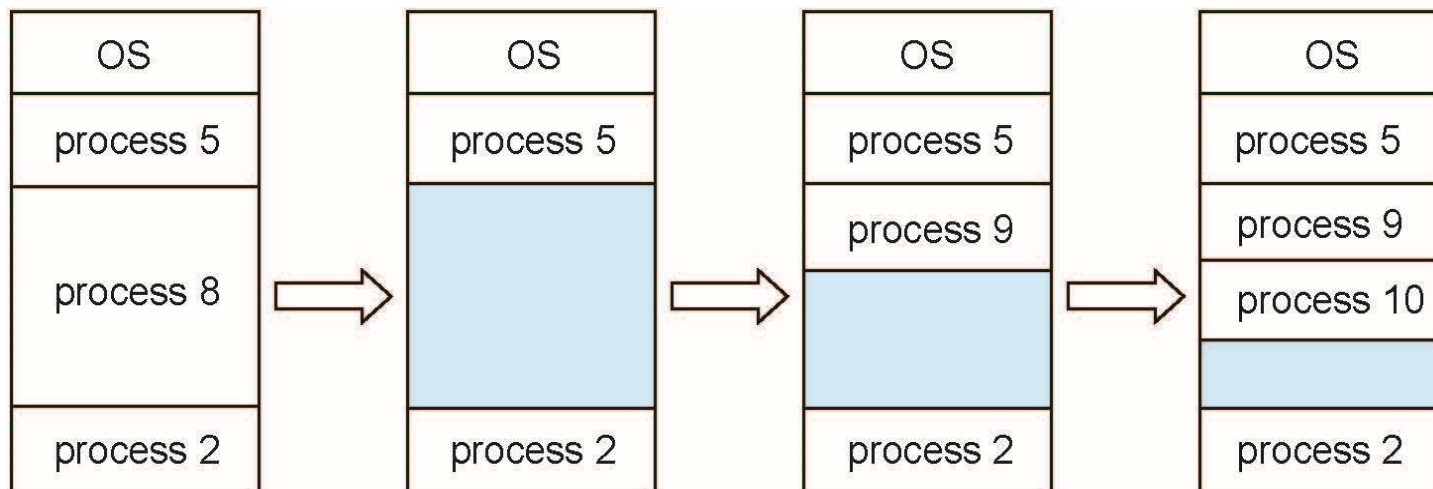
■ Comments on Fixed Partitioning

- Main memory use is inefficient. Any program, no matter how small, occupies an entire partition. This can cause *serious internal fragmentation*.
- Unequal-size partitions lessens these problems but they still remain ...
- Equal-size partitions was used in early IBM's OS/MFT (Multiprogramming with a Fixed number of Tasks).

Variable Partitioning

- Degree of multiprogramming is limited by number of partitions.
- Variable-partition sizes for efficiency (sized to a given process' needs).
- Hole – a free partition (空闲分区); holes of various size are scattered throughout memory.
- When a process arrives, it is allocated memory from a hole large enough to accommodate it.
- Process exiting frees its partition, adjacent free partitions combined.
- Operating system maintains information about:
a) allocated partitions. b) free partitions (holes).

从边缘开始分配




会产生大量外部碎片



Managing Allocated and Free Partitions


- Example: Memory with 5 processes A, B, C, D, E and 3 holes.
- Tick marks (刻度线) show memory allocation units.
- Shaded regions (0 in the bitmap) are free.






Managing Allocated and Free Partitions

- Example: Memory with 5 processes A, B, C, D, E and 3 holes.
- Tick marks (刻度线) show memory allocation units.
- Shaded regions (0 in the bitmap) are free.



Managing Allocated and Free Partitions

- Example: Free partitions combination.
 - Shaded regions are free.





■ Fragmentation

- Internal Fragmentation
 - Allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used.
- External Fragmentation
 - The total size of free partitions satisfies a request, but the available memory is not contiguous.
- First fit (首次适应) analysis reveals that given N blocks allocated, $0.5N$ blocks lost to fragmentation.
 - i.e., $1/3$ may be unusable -> *50-percent rule*.
- Reduce external fragmentation by doing compaction.
 - **Compaction** means to shuffle memory contents to place all free memory together in one large block (or possibly a few large ones).
 - Compaction is possible only if relocation is dynamic, and is done at execution time.
 - I/O problem:
 - Lock job in memory while it is involved in I/O.
 - Do I/O only into OS buffers.

■ Comments on Variable Partitioning

- Partitions are of variable length and number.
- Each process is allocated exactly as much memory as it requires.
- Eventually holes are formed in main memory. This can cause *external fragmentation*.
- *Compaction* is used to shift processes so they are contiguous; all free memory is in one block.
- Used in IBM's OS/MVT (Multiprogramming with a Variable number of Tasks).

■ Dynamic Storage - Allocation Problem


- To satisfy memory request of some size from the list of free holes, four basic placement algorithms can be applied:
 - First-fit (首次适应): Allocate the first hole that is big enough.
 - Next-fit (循环首次适应): Same logic as first-fit but starts search always from the last allocated hole (need to keep a pointer to this) in a wraparound fashion.
 - Best-fit (最佳适应): Allocate the smallest hole that is big enough; must search the entire list, unless ordered by size. It will produce the smallest leftover hole.
 - Worst-fit (最坏适应): Allocate the largest hole; must also search the entire list. It will produce the largest leftover hole.
- First-fit and best-fit are better than worst-fit in terms of speed and storage utilization.

Dynamic Storage - Allocation Problem

Example.

Which free block is selected to allocate a process of 16MB by the following algorithms.

- First-fit
- Next-fit
- Best-fit
- Worst-fit (to imagine)



■ Dynamic Storage - Allocation Problem

- Comments on the Placement Algorithms
 - **First-fit** favors allocation near the beginning of the list of free holes. It tends to create less fragmentation than Next-fit.
 - **Next-fit** often leads to allocation of the largest block at the end of memory.
 - **Best-fit** searches for smallest block. The fragment left behind is small as possible –
 - main memory quickly forms holes too small to hold any process: compaction generally needs to be done more often.
 - First/Next-fit and Best-fit better than **Worst-fit** (name is fitting) in terms of speed and storage utilization.



Buddy System

- Buddy System (*Harry Markowitz, 1963*) is a reasonable compromise to overcome disadvantages of both fixed and variable partitioning schemes.
 - Memory allocated using *power-of-2 allocation*; Satisfies requests in units sized as power of 2.
 - A request in units not appropriately sized is *rounded up* to the next highest power of 2. For example, a request for 11 KB is satisfied with a 16-KB segment.
- Memory blocks are available in size of 2^k where $l \leq k \leq u$ and where:
 - 2^l = the smallest size of block allocatable.
 - 2^u = the largest size of block allocatable
(generally, the entire memory available).
- Modified forms are used in Unix SVR4/Linux for kernel memory allocation.

Buddy System Allocation

- Buddies in the Buddy System.

两个块大小一样, 2^i

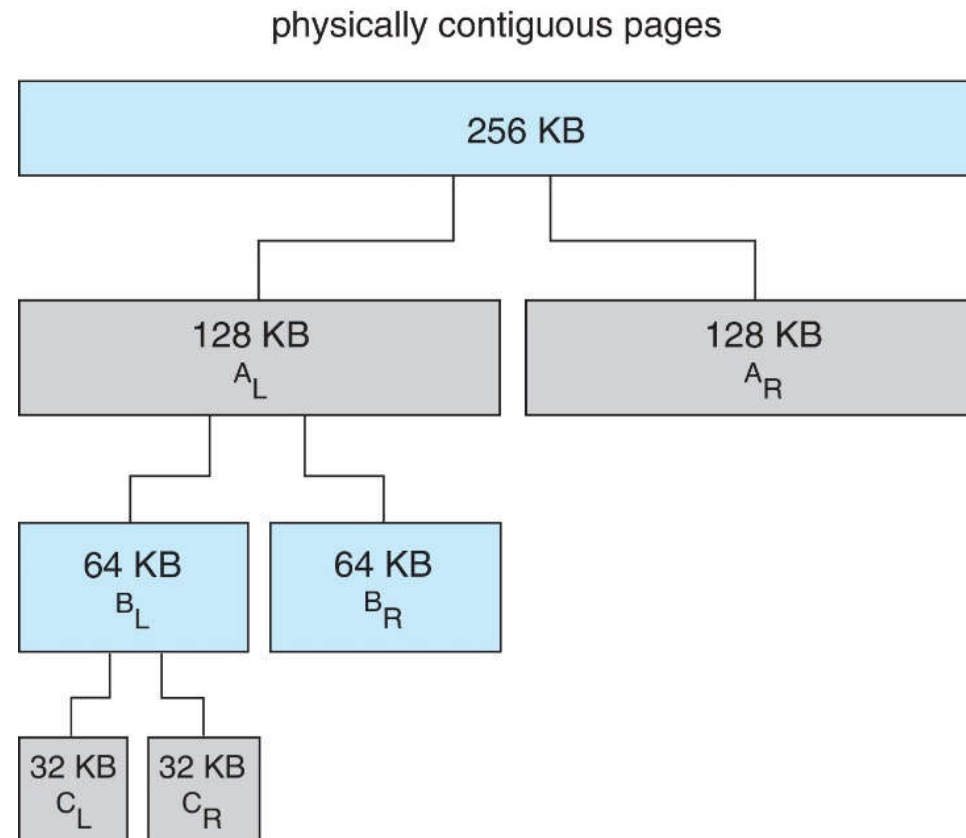
来源同一个父亲

可称为 buddies

A_L and A_R are buddies

B_L and B_R are buddies

C_L and C_R are buddies





Buddy System Allocation

■ Example of Buddy System.

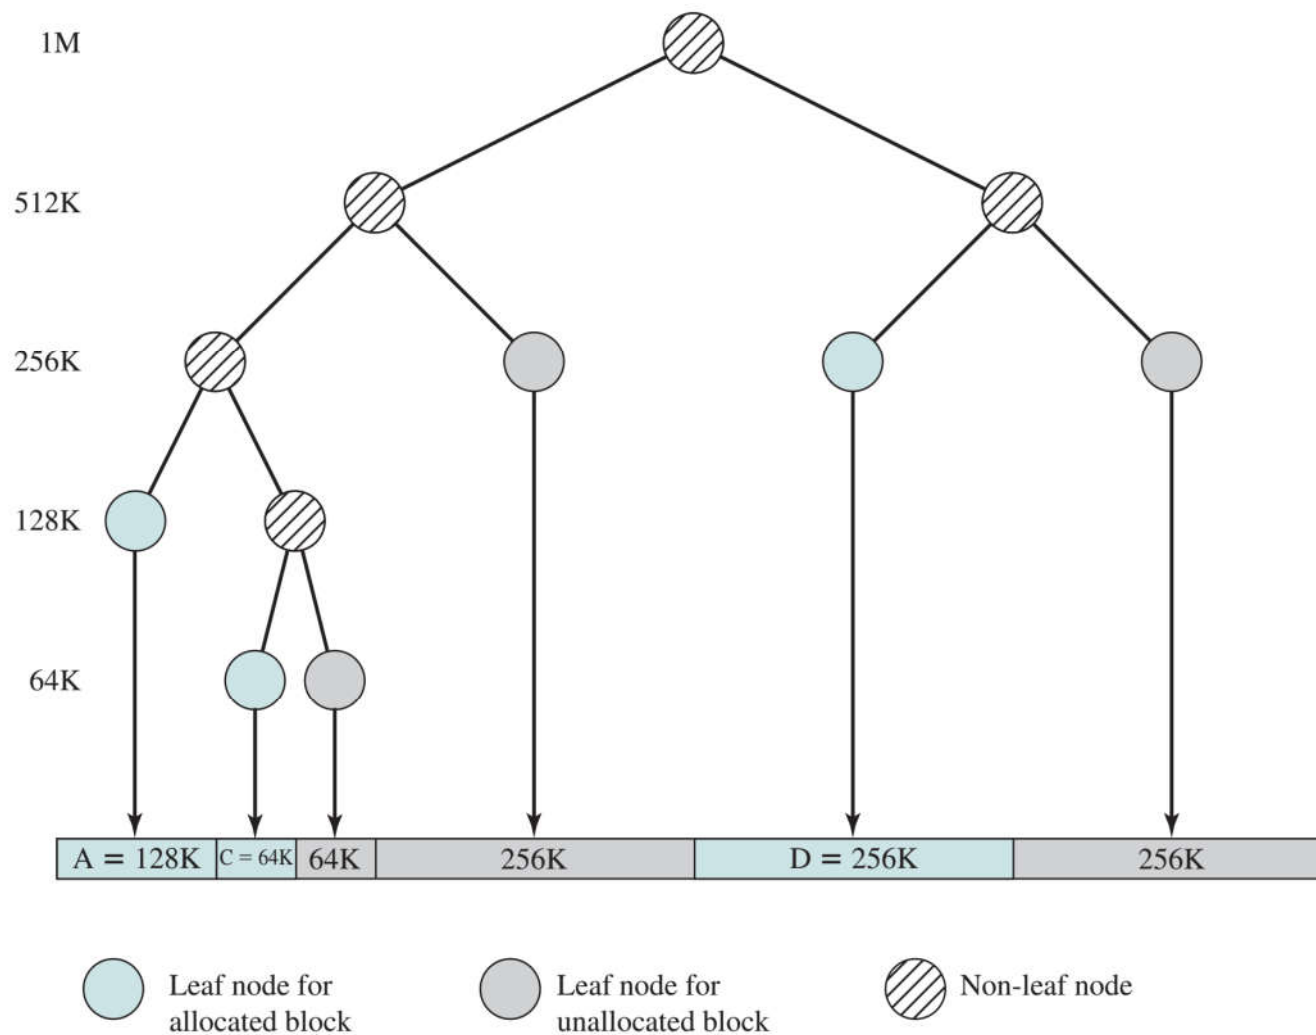
从开始分割

1-Mbyte block	1M					
Request 100K	A = 128K	128K	256K	512K		
Request 240K	A = 128K	128K	B = 256K	512K		
Request 64K	A = 128K	C = 64K	64K	B = 256K	512K	
Request 256K	A = 128K	C = 64K	64K	B = 256K	D = 256K	256K
Release B	A = 128K	C = 64K	64K	256K	D = 256K	256K
Release A	128K	C = 64K	64K	256K	D = 256K	256K
Request 75K	E = 128K	C = 64K	64K	256K	D = 256K	256K
Release C	E = 128K	128K	256K	D = 256K	256K	
Release E	512K				D = 256K	256K
Release D	1M					



Buddy System Allocation

- Tree Representation of the Buddy System.



■ Buddy System Allocation

■ Dynamics of Buddy System

- Suppose that we start with the entire block of size 2^u .
- When a request of size S , $S \leq 2^u$ is made:
 - If $2^{u-1} < S \leq 2^u$ then allocate the entire block of size 2^u to S .
 - Else split this block of size 2^u into two buddies, each of size 2^{u-1} .
 - If $2^{u-2} < S \leq 2^{u-1}$ then allocate one of the 2 buddies of size 2^{u-1} to S .
 - Otherwise one of the 2 buddies is split again.
- This process is repeated until the smallest block greater or equal to S is generated.
- Two buddies are coalesced whenever both of them become unallocated.



■ Buddy System Allocation

■ Dynamics of Buddy System

■ OS maintains several lists of holes:

- the i -list is the list of holes of size 2^i .
- whenever a pair of buddies in the i -list occur, they are removed from that list and coalesced into a single hole in the $(i+1)$ -list.

■ Presented with a request for an allocation of size S such that

$$2^{i-1} < S \leq 2^i$$

- the i -list is first examined.
- if the i -list is empty, the $(i+1)$ -list is then examined ...



■ Buddy System Allocation

■ Comments on Buddy System

- Mostly efficient when the size M of memory used by the Buddy System is a power of 2:
 - $M = 2^u$ “bytes” where u is an integer.
 - Then the size of each block is a power of 2.
 - The smallest block is of size 1.
- On average, internal fragmentation is 25%
 - Each memory block is at least 50% occupied.
- Programs are not shifted in memory:
 - Simplifies memory management.