
Memory Segmentation & Paging

Operating Systems

School of Data & Computer Science
Sun Yat-sen University

Lecture Notes: os_sysu@163.com
Instructor: Guoyang Cai
email: isscgymail@mail.sysu.edu.cn





■ Contents


- Simple Segmentation
- Simple Paging
- Address Translation
- Page Tables
- Examples

■ Segmentation Concepts

- Segmentation is memory-management scheme that supports **user view** of memory.
- A program is a collection of segments.
 - A segment is a logical unit such as:
 - main program
 - procedure
 - function
 - method
 - object
 - local variables, global variables
 - common block
 - stack
 - symbol table
 - arrays.

■ Segmentation Concepts

- User View of a Program.



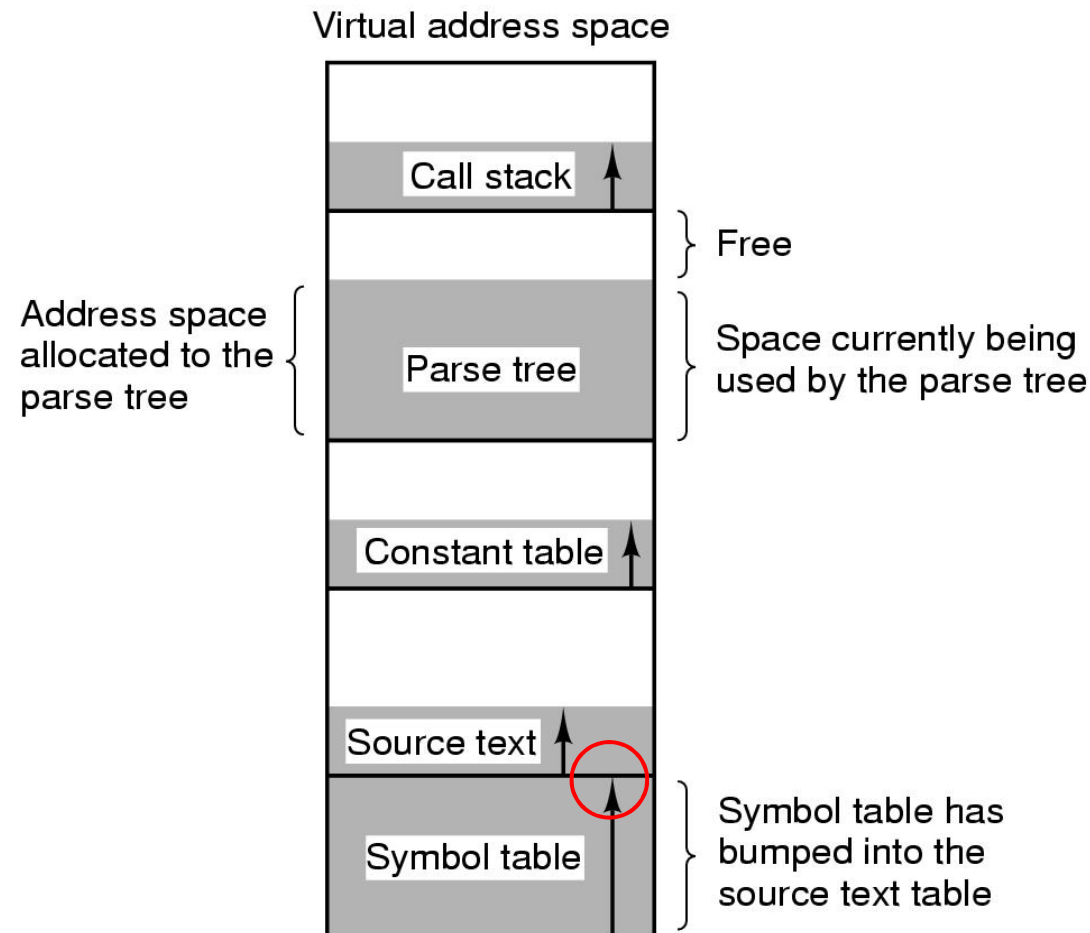
Logical address space.

■ Segmentation Concepts

- User View of a Program – Example.
 - A compiler has many tables that are built up as compilation proceeds, possibly including:
 - The source text being saved for the printed listing (on batch systems).
 - The symbol table – the names and attributes of variables.
 - The table containing integer, floating-point, constants used.
 - The parse tree (语法树), the syntactic analysis of the program.
 - The stack used for procedure calls within the compiler.

Segmentation Concepts

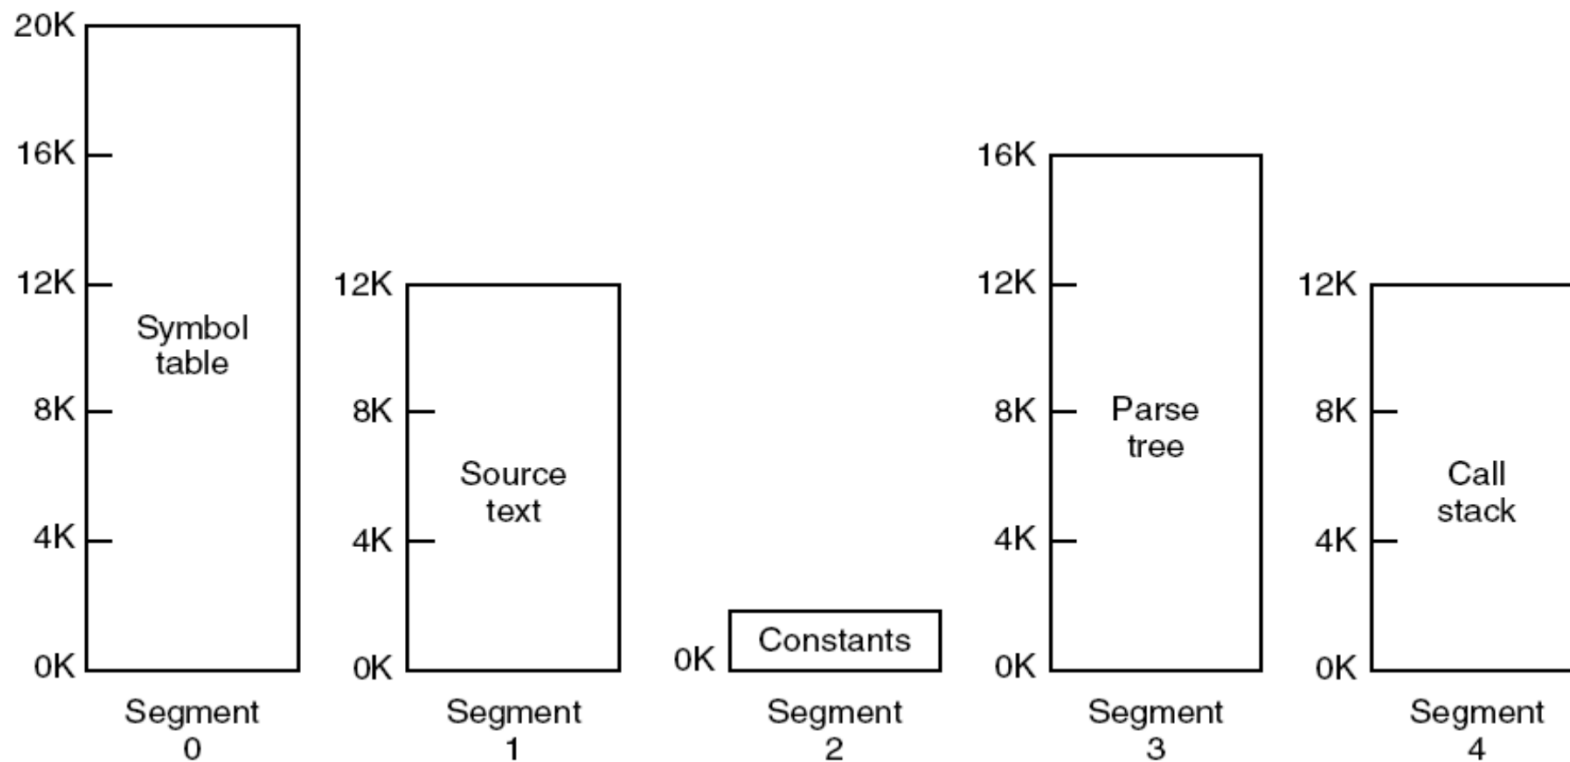
- One-dimensional address space
 - In a **one-dimensional address space** with growing tables, one table may bump into another.



■ Segmentation Concepts

■ Segmentation Solution


- Segmentation allows each table to grow or shrink independently of the other tables.




■ Dynamics of Simple Segmentation

- Each program is subdivided into blocks of non-equal size called segments.
- When a process gets loaded into main memory, its different segments can be located anywhere.
- Each segment is fully packed with instructions/data; **no internal fragmentation.**
- There is external fragmentation; it is reduced when using small segments.
- In contrast with paging, segmentation is **visible to the programmer.**
 - provided as a convenience to organize logically programs (e.g., data in one segment, code in another segment).
 - must be aware of segment size limit.
- OS maintains a **segment table** (段表) for each process. Each entry contains:
 - the starting *physical addresses* of that segment
 - the length of that segment for protection.

■ Logical view of simple segmentation



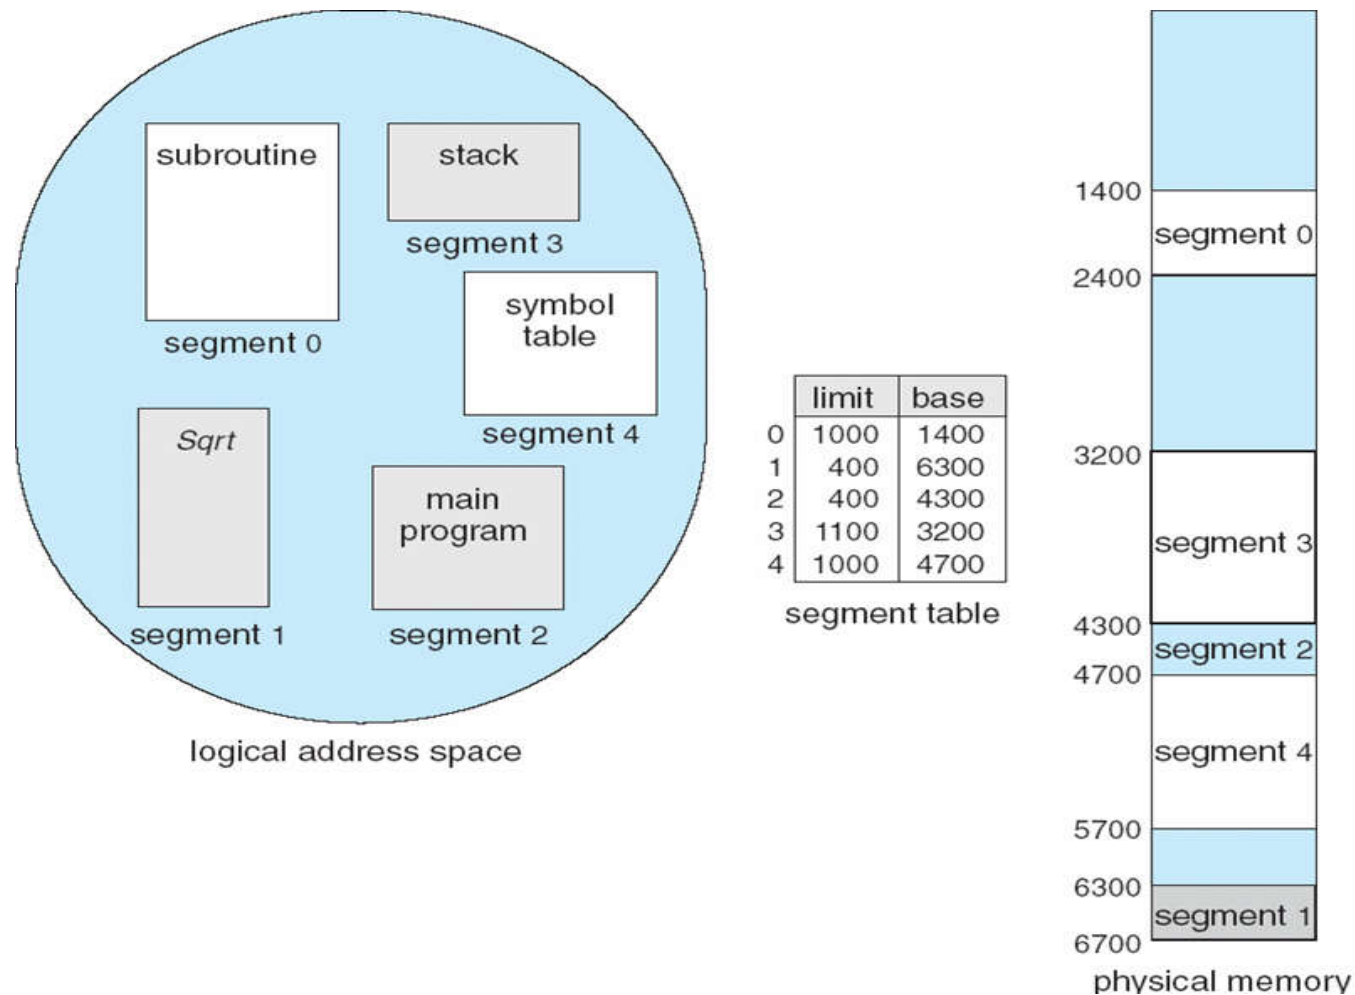
user space



physical memory space

Address Translation in segmentation

Example.



■ Segmentation Architecture

- A logical address consists of a two-tuple:

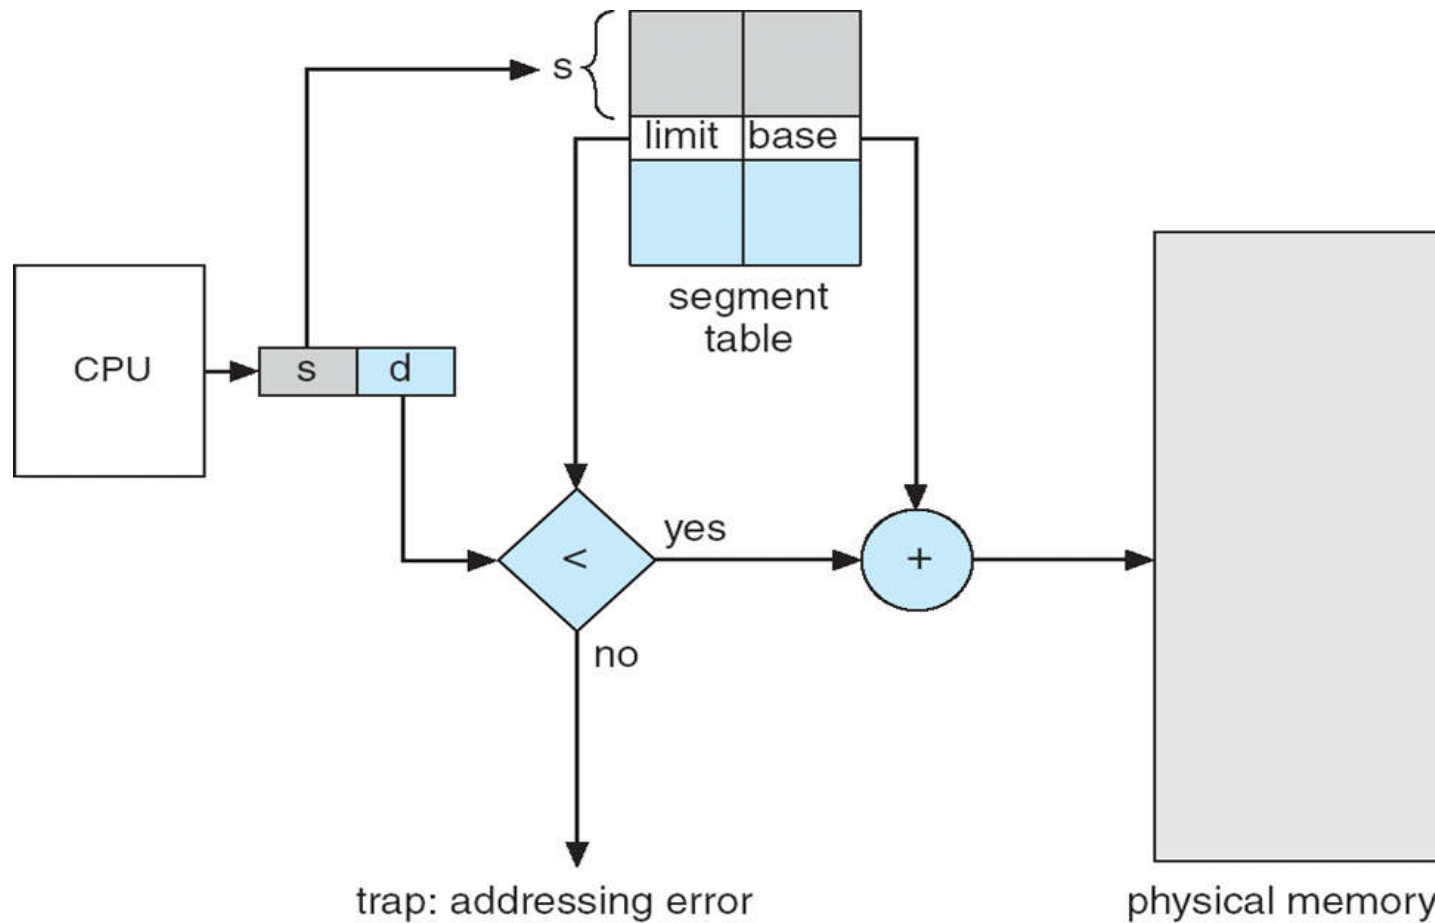
$\langle \text{segment_number}, \text{offset} \rangle$,

段号 偏移

- *Segment table* – maps two-dimensional physical addresses; each table entry has:
 - *base* – contains the starting physical address where the segments reside in memory.
 - *limit* – specifies the length of the segment.
- *Segment-table base register (STBR)* points to the location of the segment table in memory.
- *Segment-table length register (STLR)* indicates the *number of segments* used by a program
 - segment-number s is legal if $s < \text{STLR}$.

Segmentation Architecture

- Address Translation Architecture - Segmentation Hardware.



■ Segmentation Architecture

■ Address Translation in Segmentation

- When a process enters the Running state, a dedicated register gets loaded with the starting address of the process's segment table.
- Presented with a logical address

$$\langle \text{segment_number, offset} \rangle = \langle s, d \rangle,$$


the CPU indexes (with s) the segment table to obtain the starting physical address b and the length limit l of that segment.

- The physical address is obtained by adding d to b .
 - The hardware also compares the offset/displacement d with the limit l of that segment to determine if the address is valid.



Segmentation Architecture

- Address Translation in Segmentation
 - Example.



■ Segmentation Architecture

- Protection
 - With each entry in segment table associate:
 - validation bit = 0 \Rightarrow illegal segment
 - read/write/execute privileges.
- Protection bits associated with segments; code sharing occurs at segment level.
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem.

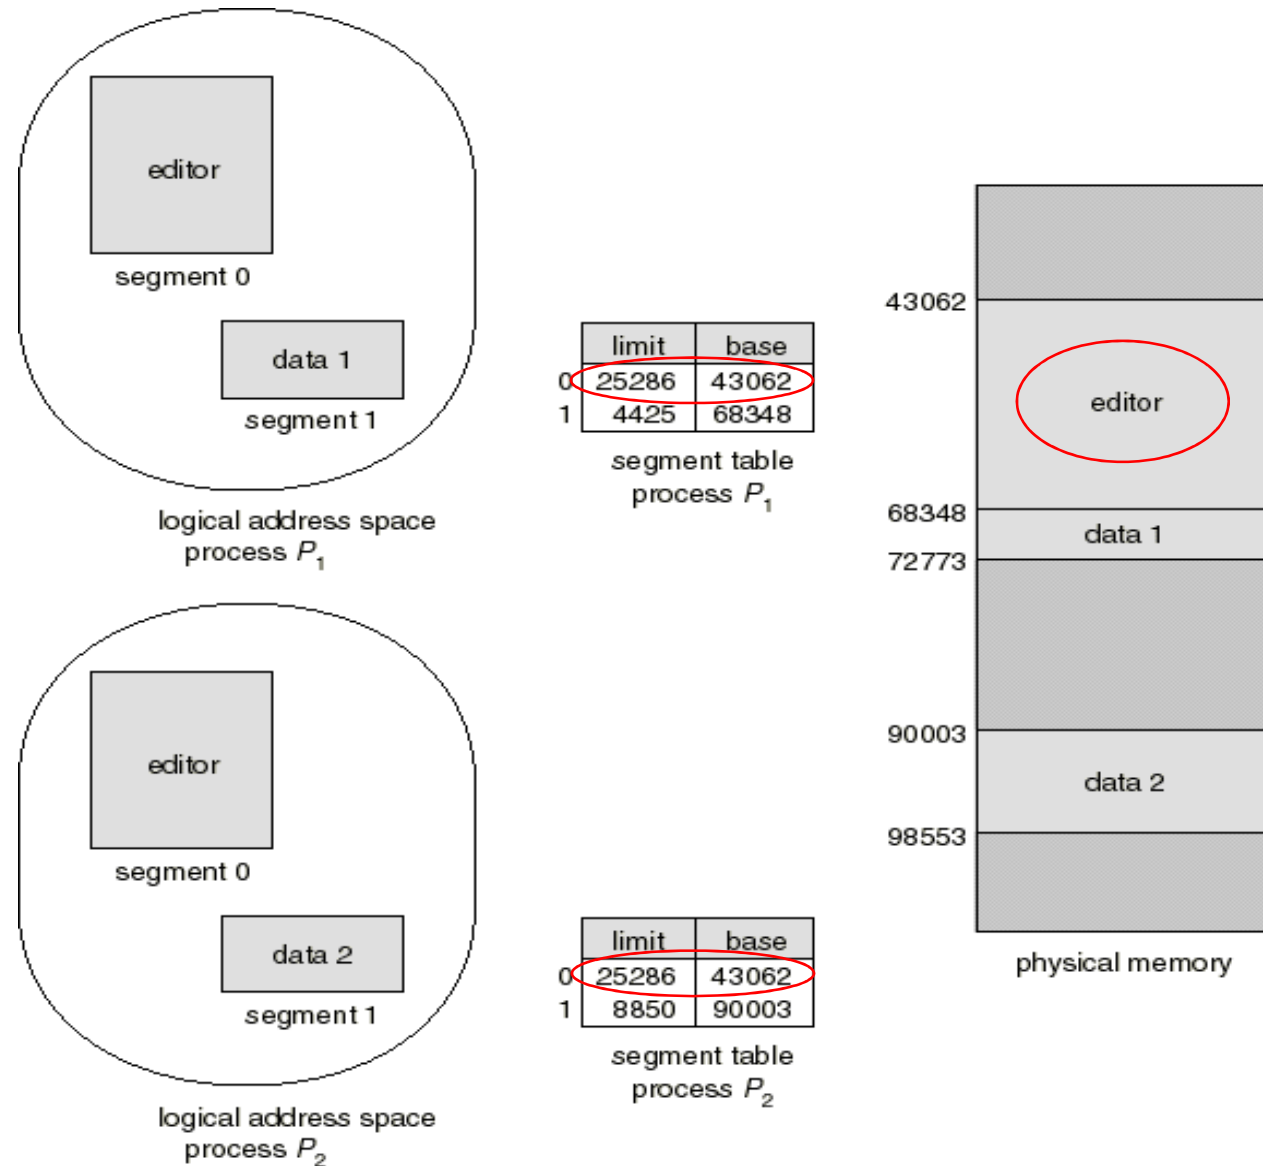
■ Shared Segments

■ Sharing in Segmentation Systems

- Segments are shared when entries in the segment tables of 2 different processes point to the same physical locations.
- Example.
 - The same code of a text editor can be shared by many users.
 - Only one copy is kept in main memory.
- But each user would still need to have its own private data segment.

Shared Segments

Example.

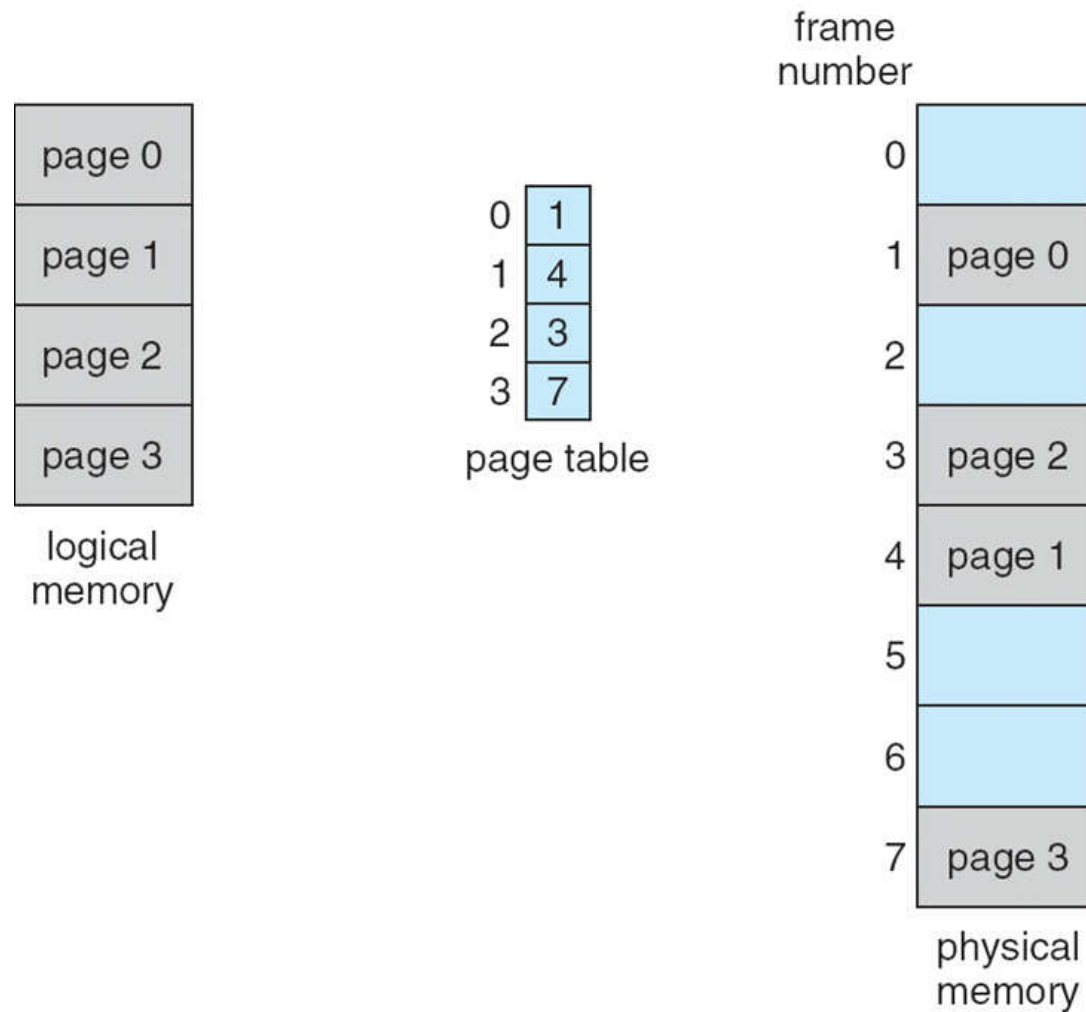


■ Simple Paging Concepts

- Idea: Physical address space of a process can be **noncontiguous**; process is allocated physical memory whenever the latter is available in size.
 - Avoids **external fragmentation**.
 - Avoids problem of varying sized memory chunks.
- Divide **physical memory** into fixed-sized chunks/blocks called **frames** (页帧, 页框, 物理页), size of power of 2, usually between 512 bytes and 16 MB.
- Divide **logical memory** into blocks of the same size as physical frames, called **pages** (页, 页面, 逻辑页).
- Process pages can thus be assigned to any free frames in main memory; a process does not need to occupy a contiguous portion of physical memory. To run a program of size n pages, we need to find n free frames of physical memory and load the program.
 - So we need to use a **free-frame list** to keep track of all free frames in physical memory.
 - And we need to set up a **page table** (页表) to translate logical pages to physical frames.
- **Internal fragmentation** possible only for the page at the end of program.


Simple Paging Concepts

Paging Example.



Simple Paging Concepts

- Free-frame list example.



■ Simple Paging Concepts

- OS now needs to maintain in main memory a *page table* for each process.
 - Each entry of a page table consists of the frame number where the corresponding page is physically located.
 - The corresponding page table is indexed by the page number to obtain the frame number.
- A *free frame table/list*, of available frames, is also maintained.
 - Example.

0	0
1	1
2	2
3	3

Process A
page table

0	—
1	—
2	—

Process B
page table

0	7
1	8
2	9
3	10

Process C
page table

0	4
1	5
2	6
3	11
4	12

Process D
page table

13
14

Free-frame
list

■ Calculating Internal Fragmentation

■ Example.

- Page/frame size: 2,048 bytes
- Suppose some process has size: 72,766 bytes
 - $72766 = 2048 \times 35 + 1,086$.
 - Need to allocate 36 frames to the process.
 - Internal fragmentation: $2,048 - 1,086 = 962$ (bytes)
- Worst case fragmentation: 2048 - 1 (byte)
 - For processes having size of $2048p + 1$ bytes.
- On average internal fragmentation: 1/2 frame size.

■ So small frame sizes desirable?

- But each page table entry takes memory to track.
- Page sizes growing over time:
 - Linux page sizes: 4KB / 2MB

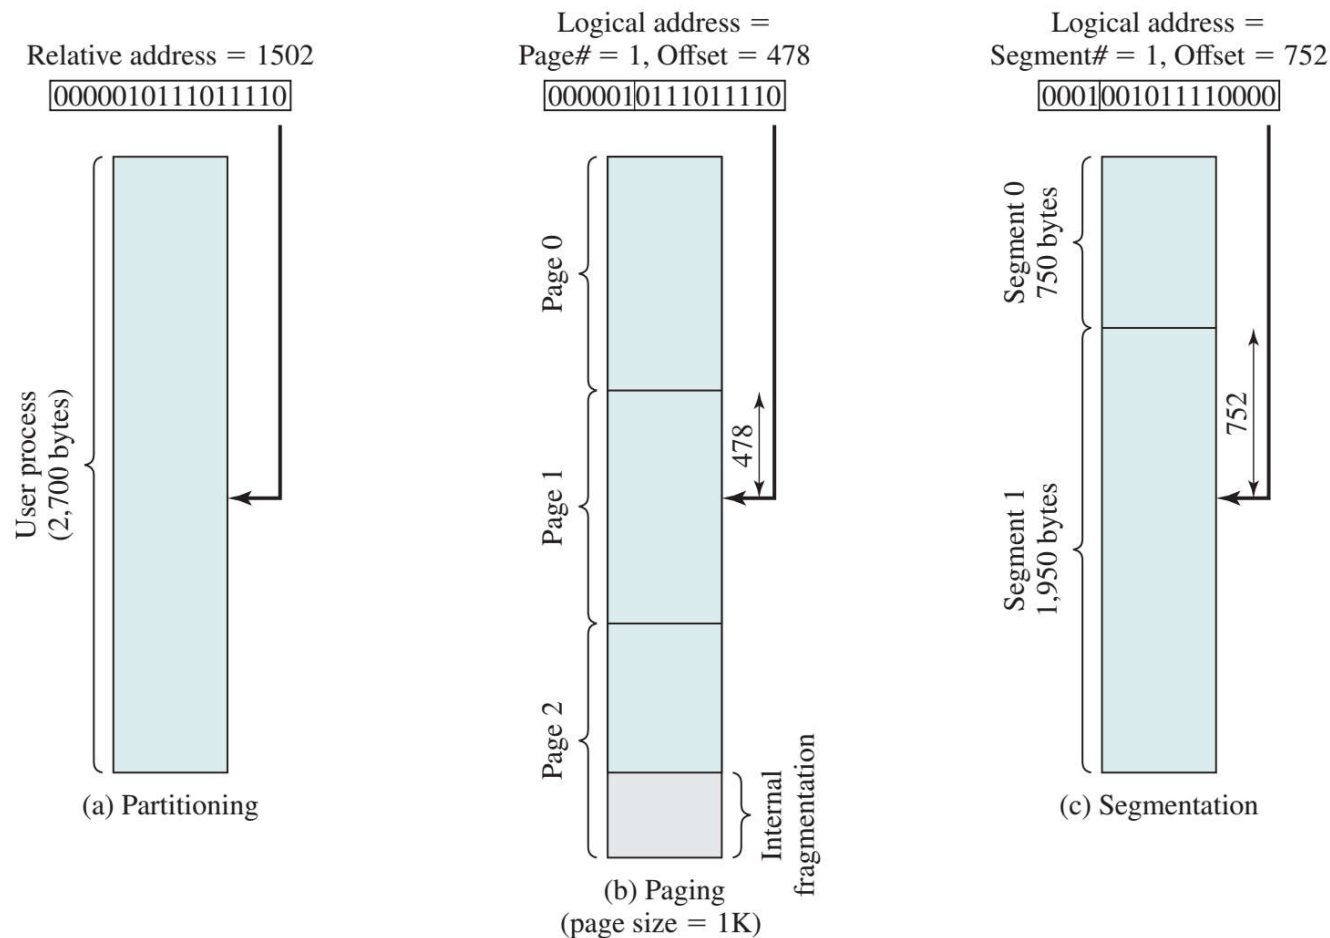
```
$getconf PAGESIZE
```

```
$cat /proc/meminfo | grep Huge
```
 - Solaris page sizes: 8KB / 4MB.



Logical Address in Paging

- The logical address becomes a relative address when the page size is a power of 2.
- Comparison of Partitioning, Paging and Segmentation.



■ Logical Address in Paging

- The logical address becomes a relative address when the page size is a power of 2.
 - Example.
 - If 16-bit addresses are used and page size = 1K, we need 10 bits for offset and have 6 bits available for page number.
 - Then the 16 bit address, obtained with the 10 least significant bits as offset and 6 most significant bits as page number, is a location relative to the beginning of the process.



■ Logical Address in Paging

- Within each program, each logical address must consist of a page number and an offset/displacement within the page

$$\langle \text{page_number}, \text{offset} \rangle = \langle p, d \rangle.$$

- A dedicated register always holds the starting physical address of the page table of the currently running process.
- Presented with the logical address

$$\langle \text{page_number}, \text{offset} \rangle = \langle p, d \rangle,$$


the processor accesses the page table to obtain the physical address

$$\langle \text{frame_number}, \text{offset} \rangle = \langle f, d \rangle.$$



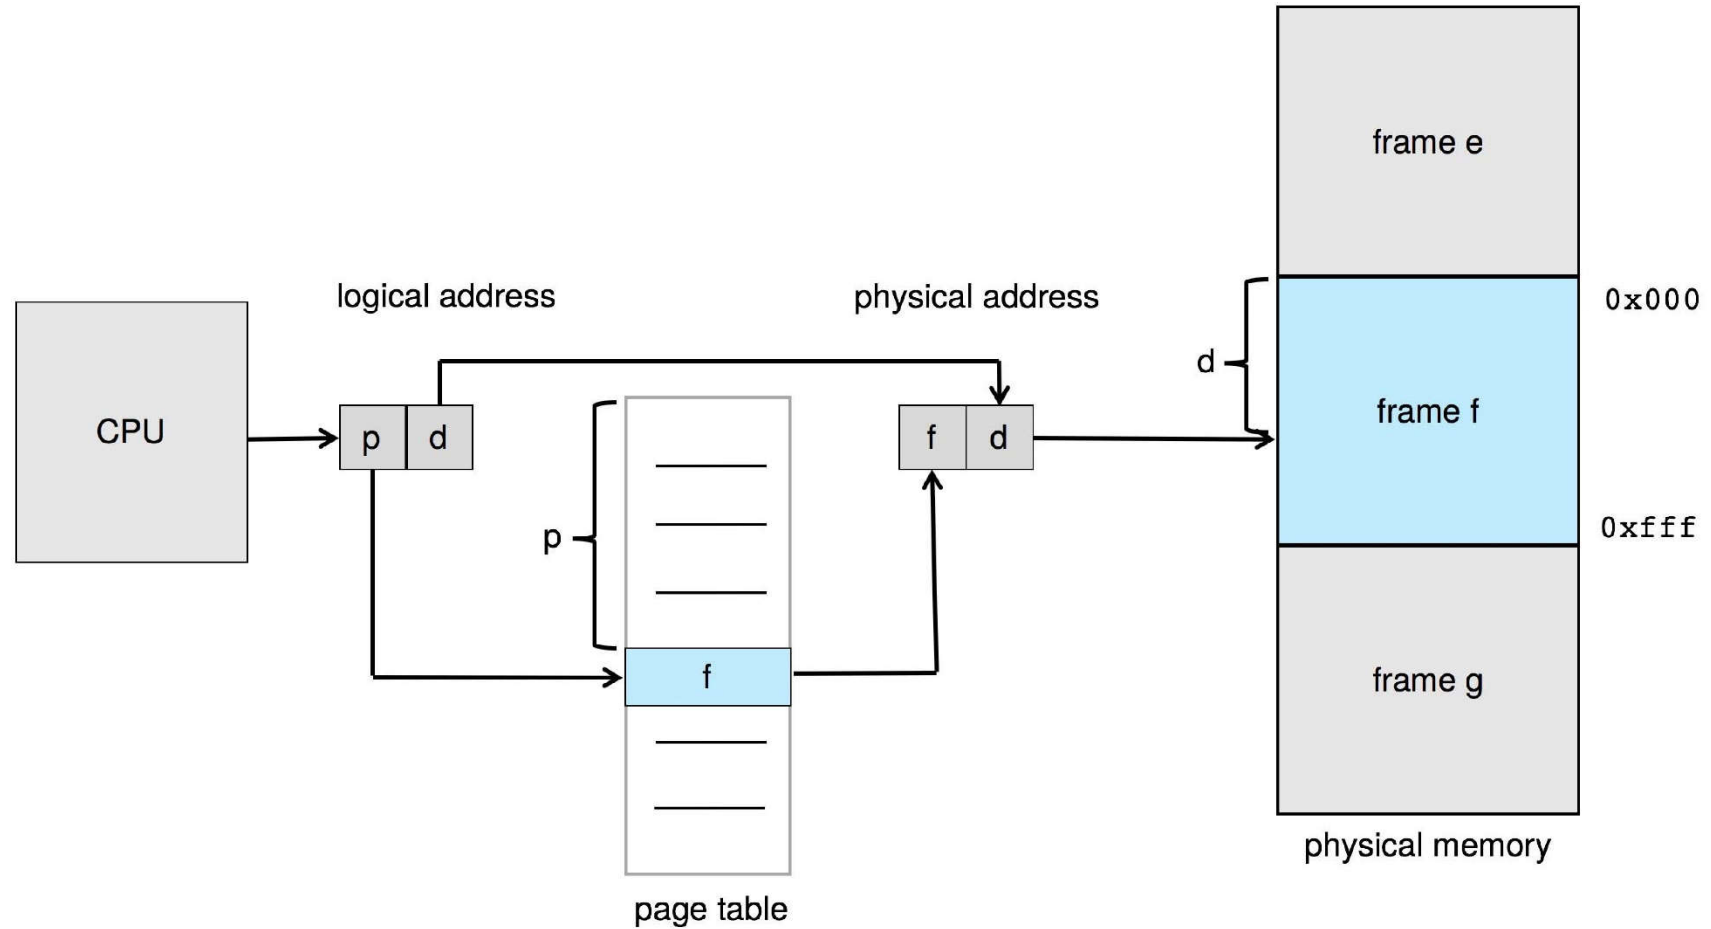
■ Address Translation Scheme

- Logical address generated by CPU is divided into two parts:
 - Page number (**p**) – used as an index into a page table which contains the base address of each page in physical memory.
 - Page offset/displacement (**d**) – combined with base address to define the physical memory address that is sent to the memory unit.
- For given logical address space 2^m and page size 2^n ($n < m$),



- By using a page size of a power of 2, the pages are *invisible* to the programmer, compiler/assembler, and the linker.
- Address translation at run-time is then easy to implement in hardware:
 - Logical address **<p, d>** gets translated to physical address **<f, d>** by indexing the page table with **p** and *appending* the same offset **d** to the frame of number **f**.


■ Address Translation Architecture



Address Translation Architecture

Example.

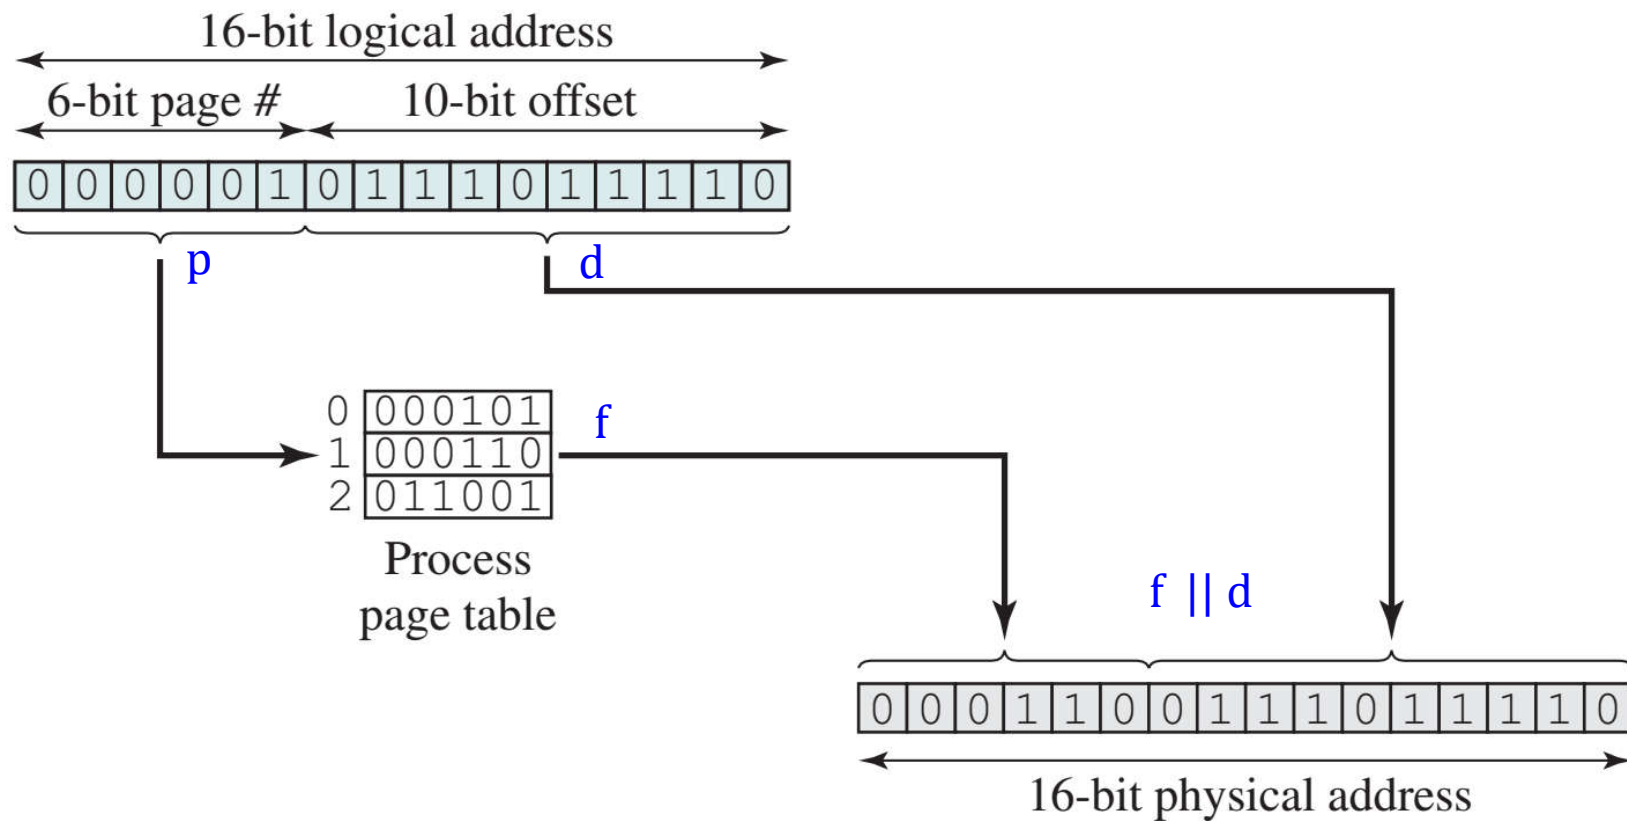
$m = 4, n = 2$, page size = $2^2 = 4$ (bytes), memory size = 32 (bytes).



Address Translation Architecture

Example.

$m = 16, n = 10$, page size = $2^{10} = 1024$ (bytes).






■ Implementing Page Table

- If we keep Page Table in main memory
 - *Page-table base register (PTBR)* points to the location of the page table in memory.
 - *Page-table length register (PTLR)* indicates the size of the page table.
 - Every data/instruction access requires *two memory accesses*.
 - one for the page table, and one for the data/instruction
- If we keep Page Table in hardware (in MMU?)
 - However, page table can be large – too expensive.
- The two memory accesses problem can be solved by combining these two mechanisms.
 - Use a special fast-lookup hardware cache called *Associative Memory* (Registers) or *Translation Look-aside Buffer (TLB, 快表)* – enables fast parallel search.
 - Address translation $\langle p, d \rangle$
 - If p is in associative register, get frame number out.
 - Otherwise get frame number from page table in memory.

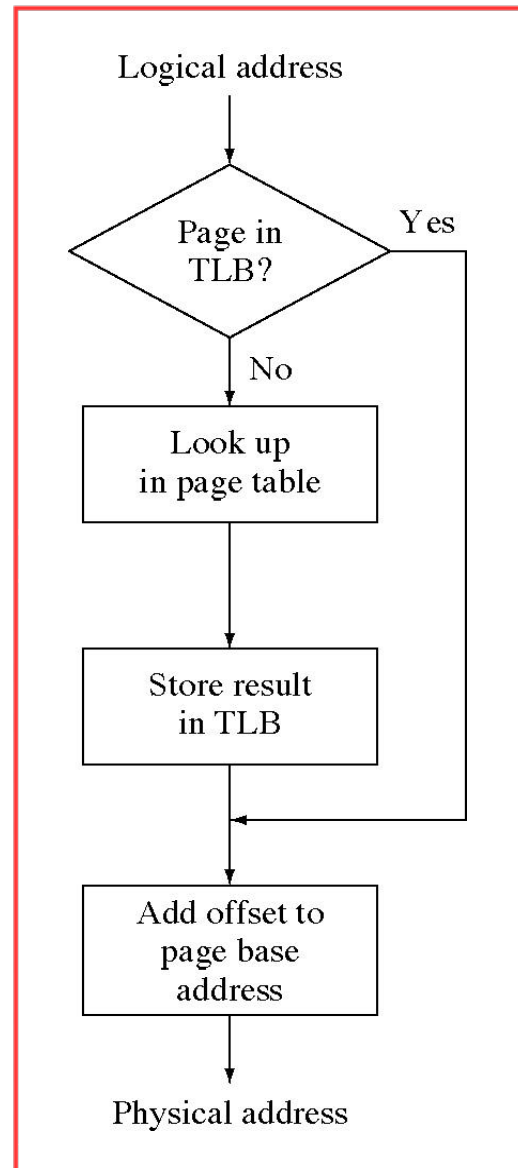
■ Implementing Page Table

■ Paging Hardware with TLB.



■ Implementing Page Table

- Paging Hardware with TLB.
- TLB Flow Chart.



■ Implementing Page Table

- Paging Hardware with TLB
 - TLB takes advantage of the *Locality Principle* (局部性原理).
 - TLB uses associative mapping hardware to simultaneously interrogate all TLB entries to find a match/hit on page number.
 - TLB hit rates are 90%+.
 - TLB must be flushed (erased) each time a new process enters the running state.
 - There is only one global TLB serving for all processes.
 - TLB information maybe kept/loaded in/from process context.

■ Implementing Page Table

■ Paging Hardware with TLB

- Each associative (memory) lookup time = ϵ time unit.

- time unit: memory access time
- can be < 10% of (far faster than) memory access time

- Hit ratio = α

- percentage of times that a page number is found in the associative memory;
- related to number of associative registers.

- Effective Access Time (EAT)

$$\text{EAT} = \alpha(\epsilon + 1) + (1 - \alpha)(\epsilon + 2) = 2 + \epsilon - \alpha.$$

Unit: memory access time.

- EAT is between $1 + \epsilon$ and $2 + \epsilon$ access times.

- should be closer to 1.

■ Implementing Page Table

■ Paging Hardware with TLB

■ Effective Access Time (EAT)

$$\text{EAT} = \alpha(\epsilon + 1) + (1 - \alpha)(\epsilon + 2) = 2 + \epsilon - \alpha.$$

Unit: memory access time.

■ Examples: Assume that memory cycle time is 100ns and $\epsilon = 20\text{ns}$ for TLB search.

- Consider $\alpha = 80\%$. Then

$$\text{EAT} = 0.80 \times (20 + 100) + 0.20 \times (20 + 200) = 140(\text{ns}), \text{ or}$$

$$\text{EAT} = 2 \times 100 + 20 - 1 \times 100 \times 0.8 = 140(\text{ns})$$

So it is $(140 - 100) / 100 = 40\%$ slowdown in memory access time.

- Consider more realistic hit ratio $\alpha = 98\%$. Then

$$\text{EAT} = 0.98 \times 120 + 0.02 \times 220 = 122(\text{ns}), \text{ or}$$

$$\text{EAT} = 2 \times 100 + 20 - 1 \times 100 \times 0.98 = 122(\text{ns})$$

So it is only $(122 - 100) / 100 = 22\%$ slowdown in memory access time.

■ Implementing Page Table

■ Paging Hardware with TLB

■ Advanced TLB Aspects

- Some TLBs store address-space identifiers (ASIDs) in each TLB entry – uniquely identifies each process to provide address-space protection for that process:
 - Otherwise need to flush at every context switch.
- TLBs typically small (64 to 1,024 entries).
- On a TLB miss, value is loaded into the TLB for faster access next time:
 - Replacement policies must be considered.
 - Some entries can be *wired down* (not removed from the TLB) for permanent fast access.

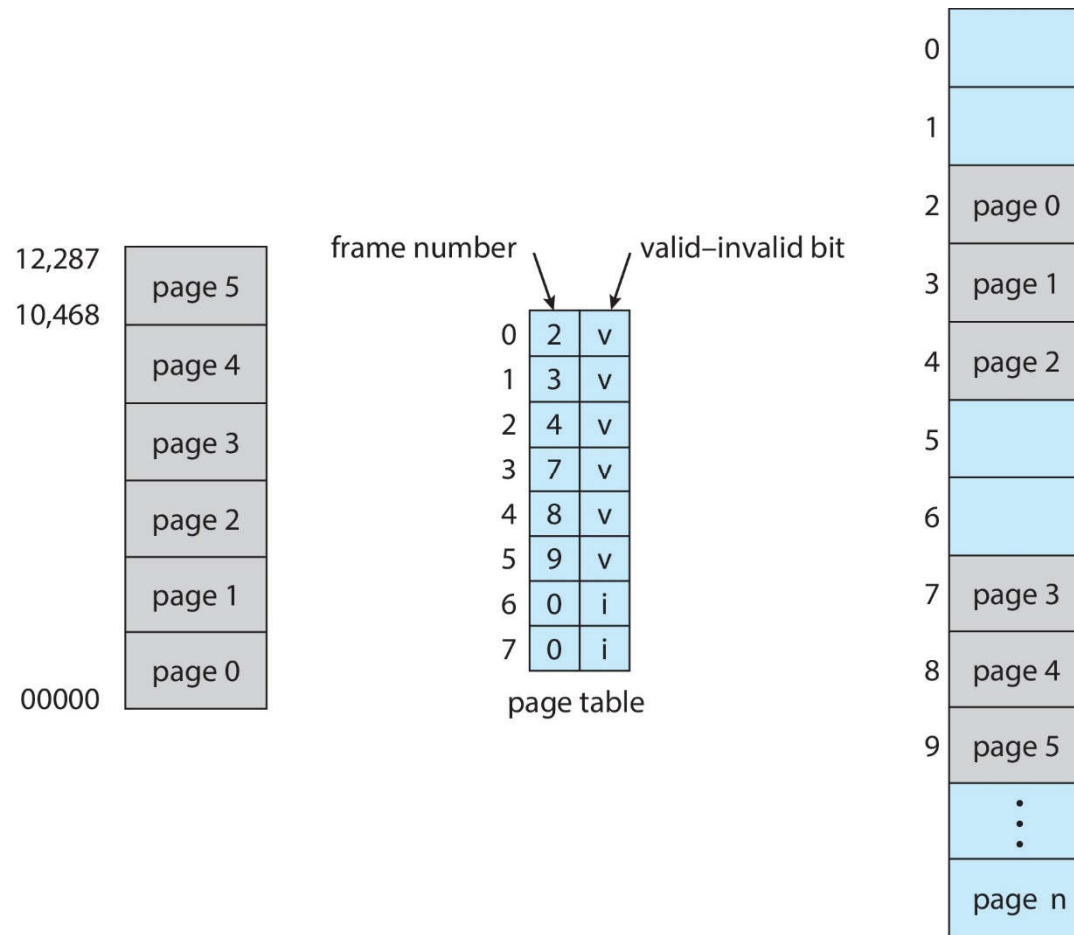


■ Memory Protection


- Memory protection is implemented by associating protection bit with *each frame* to indicate if read-only or read-write access is allowed.
 - can also add more bits to indicate page execute-only, and so on.
- Valid-invalid bit is attached to each entry in the page table.
 - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page.
 - “invalid” indicates that the page is not in the process’ logical address space.
 - Or use *page-table length register (PTLR)*
- Any violations result in a trap to the kernel.

Memory Protection

- Valid (v) or Invalid (i) Bit in Page Table.



■ Transfer of a Paged Memory to Contiguous Disk Space



■ Shared Pages

■ Shared code


- One copy of read-only (*reentrant*) code shared among processes (e.g., text editors, compilers, window systems).
 - Similar to multiple threads sharing the same process space.
- Shared code must appear in the same physical location in the logical address spaces of all processes.

■ Private code and data

- Each process keeps separate copy of code and data.
- The pages for the private code and data can appear anywhere in the logical address space.

Shared Pages

Shared Pages Example.




■ Structures of Page Tables

- Memory structures for paging can get huge using straight-forward methods.
 - Consider a 32-bit logical address space as on modern computers with page size of 4KB (2^{12}).
 - Page table would have 1 million entries ($2^{32 - 12} = 2^{20}$).
 - If each entry is 4 bytes, then 4MB of physical address space is needed for page table alone.
 - That amount of memory used to cost a lot.
 - It is hard to allocate that contiguously in main memory.
- There are three structures for page tables.
 - Hierarchical Page Tables
 - Hashed Page Tables
 - Inverted Page Tables.

Structures of Page Tables

Hierarchical Page Tables 分级页表

- Break up the logical address space into multiple page tables.
 - A simple technique is a two-level page table.




A two-level page-table scheme

■ Structures of Page Tables

■ Hierarchical Page Tables

■ Consider a 32-bit logical address space.

- A logical address (on 32-bit machine with $2^{12} = 4K$ page size) is divided into:
 - a page number consisting of 20 bits
 - a page offset consisting of 12 bits.
- Since the page table is paged, the page number is further divided into:
 - a 10-bit outer page number
 - a 10-bit inner page offset.
- Thus, a logical address is as follows:

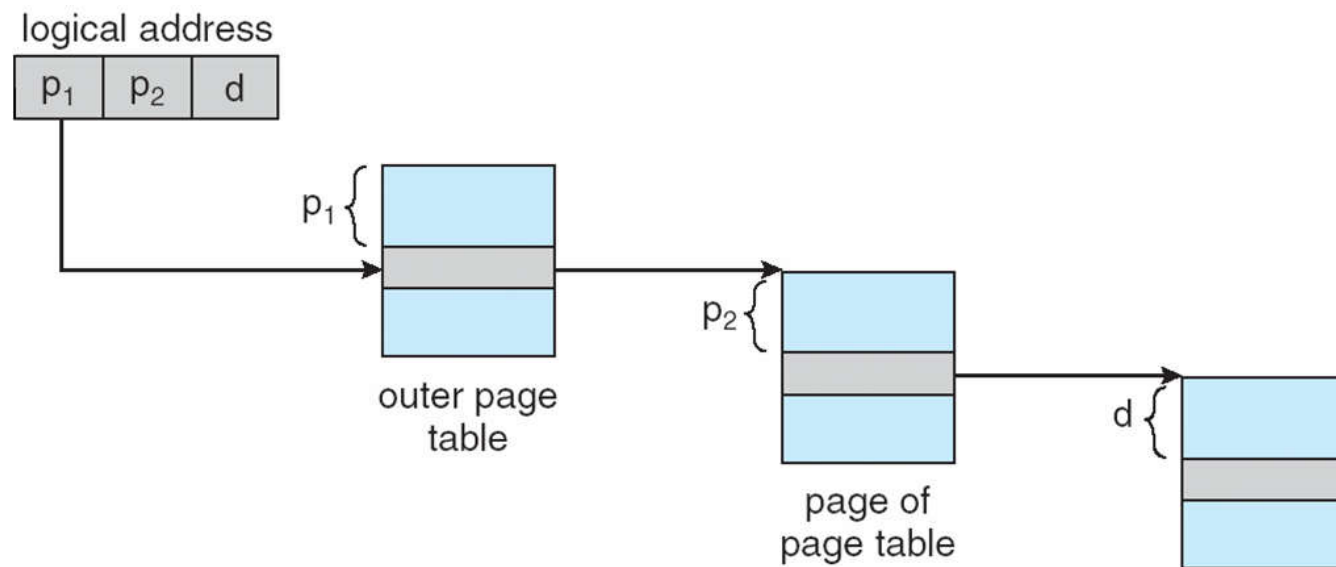


where p_1 is an index into the **outer** page table, and p_2 is the **displacement** within the page of the **inner** page table.

- Known as ***forward-mapped page table***

Structures of Page Tables

Hierarchical Page Tables.




Two-level Paging Scheme.

Structures of Page Tables

Hierarchical Page Tables

Consider a 64-bit Logical Address Space.


- Even two-level paging scheme is not sufficient.
- If page size is 4KB ($= 2^{12}$), then page table has 2^{52} entries.
 - For two level scheme, inner page tables could have 2^{10} entries of 4 bytes length. Address would look like



- Outer page table has 2^{42} entries (2^{44} bytes, 4 bytes for each entry).
- One solution is to add a 2nd outer page table.
- But in the following example the 2nd outer page table is still 2^{34} bytes in size.
 - And possibly 4 memory access to get to one physical memory location.

Structures of Page Tables

- Hierarchical Page Tables
 - Three-level Paging Scheme.



■ Structures of Page Tables


■ Hashed Page Tables

哈希页表

- Common in address spaces > 32 bits.
- The virtual page number is hashed into a page table.
 - This page table contains a chain of elements hashing to the same location.
- Each element of the chain table contains:
 - (1) The virtual page number
 - (2) The value of the mapped page frame
 - (3) A pointer to the next element.
- Virtual page numbers are compared in this chain searching for a match.
 - If a match is found, the corresponding physical frame is extracted.
- Variation for 64-bit addresses is *clustered page tables*.
 - Similar to hashed but each entry refers to several pages (such as 16) rather than 1.
 - Especially useful for *sparse* address spaces (where memory references are non-contiguous and scattered).

Structures of Page Tables

Hashed Page Tables



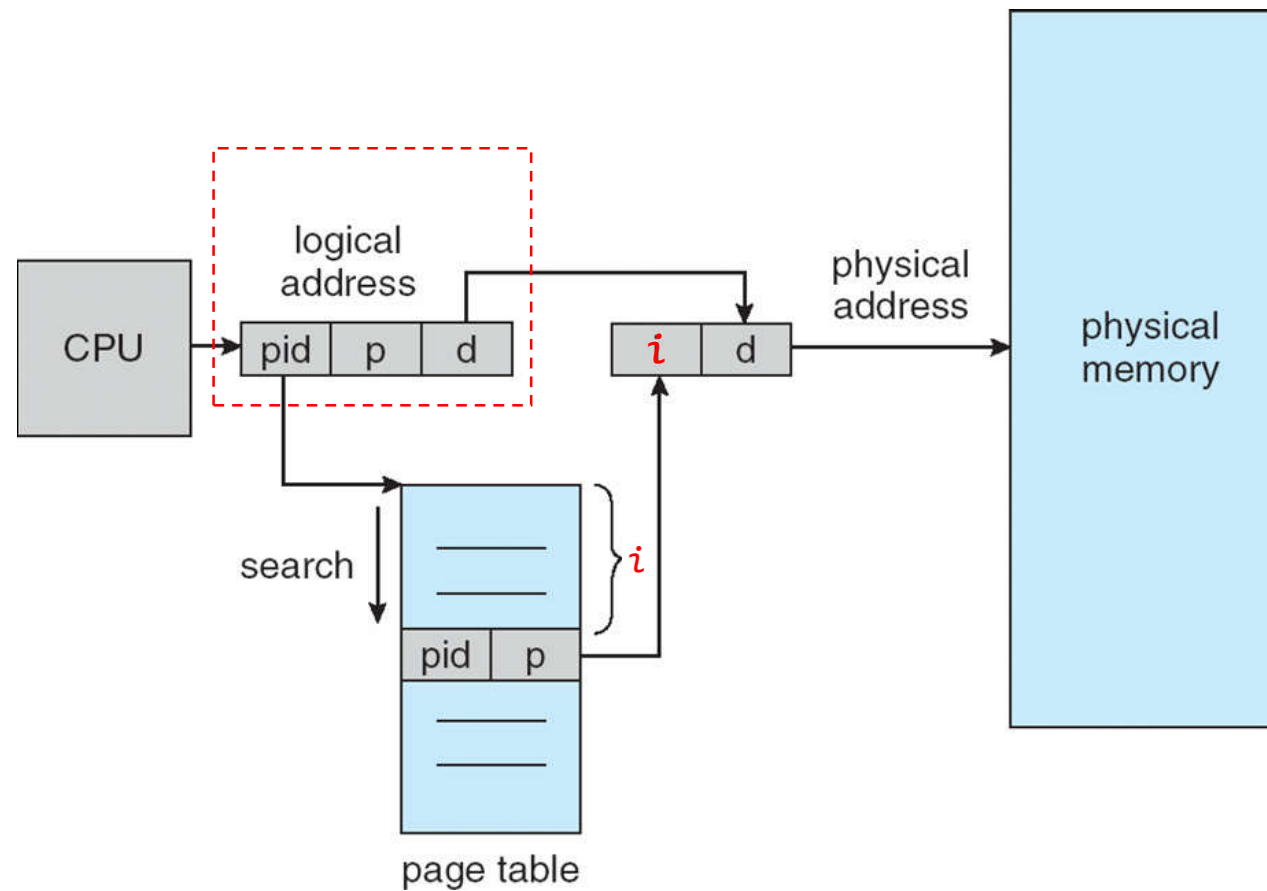
■ Structures of Page Tables

■ Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages.
- Entries of an inverted page table consists of the virtual address of the page stored in that real (physical) memory location, with information about the process that owns that page.
 - one entry for each real page of memory
 - *e.g.*, <process-id, page-number>
 - used by 64-bit Ultra SPARC, PowerPC, ...
- This scheme decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs.
- Use hash table to limit the search to one — or at most a few — page-table entries.
 - TLB can accelerate access.
- But how to implement shared memory?
 - One mapping of a virtual address to the shared physical address.

Structures of Page Tables

Inverted Page Table



Inverted page table architecture

■ Simple Segmentation/Paging Comparison

- Segmentation is visible to the programmer whereas Paging is transparent.
- Naturally supports protection/sharing.
- Segmentation can be viewed as commodity offered to the programmer to logically organize a program into segments while using different kinds of protection (example: execute-only for code but read-write for data).
- Segments are variable-size; Pages are fixed-size.
- Segmentation requires more complicated hardware for address translation than Paging.
- Segmentation suffers from external fragmentation. Paging only yields a small internal fragmentation.
- Maybe combine Segmentation and Paging?





■ The Intel 32 Architecture

- Dominant industry chips
- Pentium CPUs are 32-bit and called IA-32 architecture
 - Current Intel CPUs are 64-bit and called IA-64 architecture
- Many variations in the chips, cover the main ideas here
- Supports both segmentation and segmentation with paging
 - Each segment can be 4GB
 - Up to 16K segments per process
 - Divided into two partitions
 - First partition of up to 8K segments are private to process (kept in *local descriptor table* (LDT))
 - Second partition of up to 8K segments shared among all processes (kept in *global descriptor table* (GDT))
- CPU generates logical address
 - Given to segmentation unit which produces linear addresses.
 - Linear address given to paging unit:
 - Which generates physical address in main memory.
 - Paging units form equivalent of MMU.
 - Pages sizes can be 4KB or 4MB.

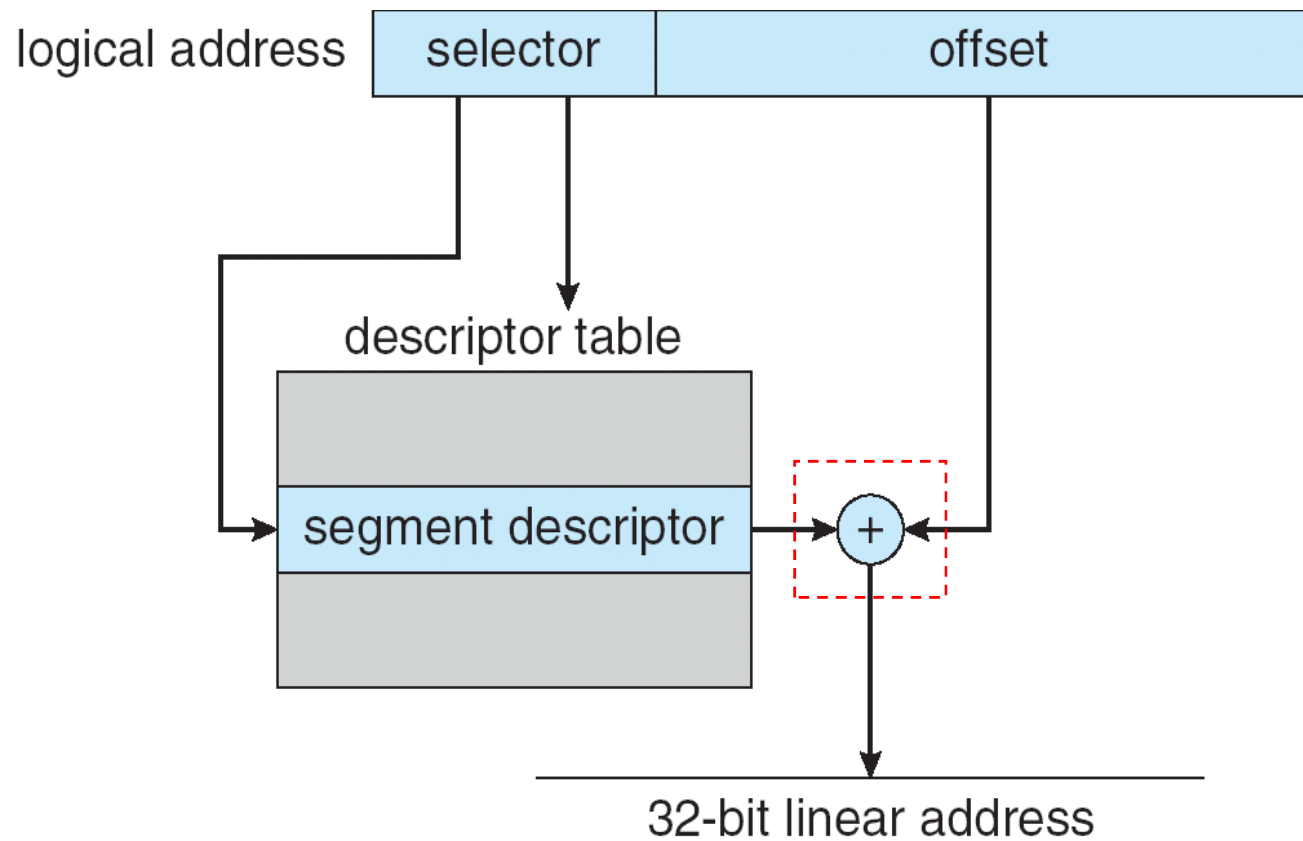
■ The Intel 32 Architecture

- Logical to Physical Address Translation.




■ The Intel 32 Architecture

- Intel Pentium Segmentation.



■ The Intel 32 Architecture

■ Intel Pentium Paging Architecture.




Two-level paging scheme
for 4 KB page.

Direct paging scheme
for 4 MB page.

■ The Intel 32 Architecture

■ Intel IA-32 Page Address Extensions

- 32-bit address limits led Intel to create *page address extension (PAE)*, allowing 32-bit apps access to more than 4GB of memory space.
 - Paging went to a three-level scheme
 - Top two bits refer to a *page directory pointer table*
 - Page-directory and page-table entries moved to 64-bits in size
 - Net effect is increasing address space to 36 bits – 64GB of physical memory






■ The Intel 32 Architecture

■ Intel X86-64

- current generation Intel x86 architecture
- 64 bits is ginormous (> 16 exabytes)
- in practice only implement 48 bit addressing
 - Page sizes of 4KB / 2MB / 1GB
 - Four levels of paging hierarchy
- can also use PAE so virtual addresses are 48 bits and physical addresses are 52 bits



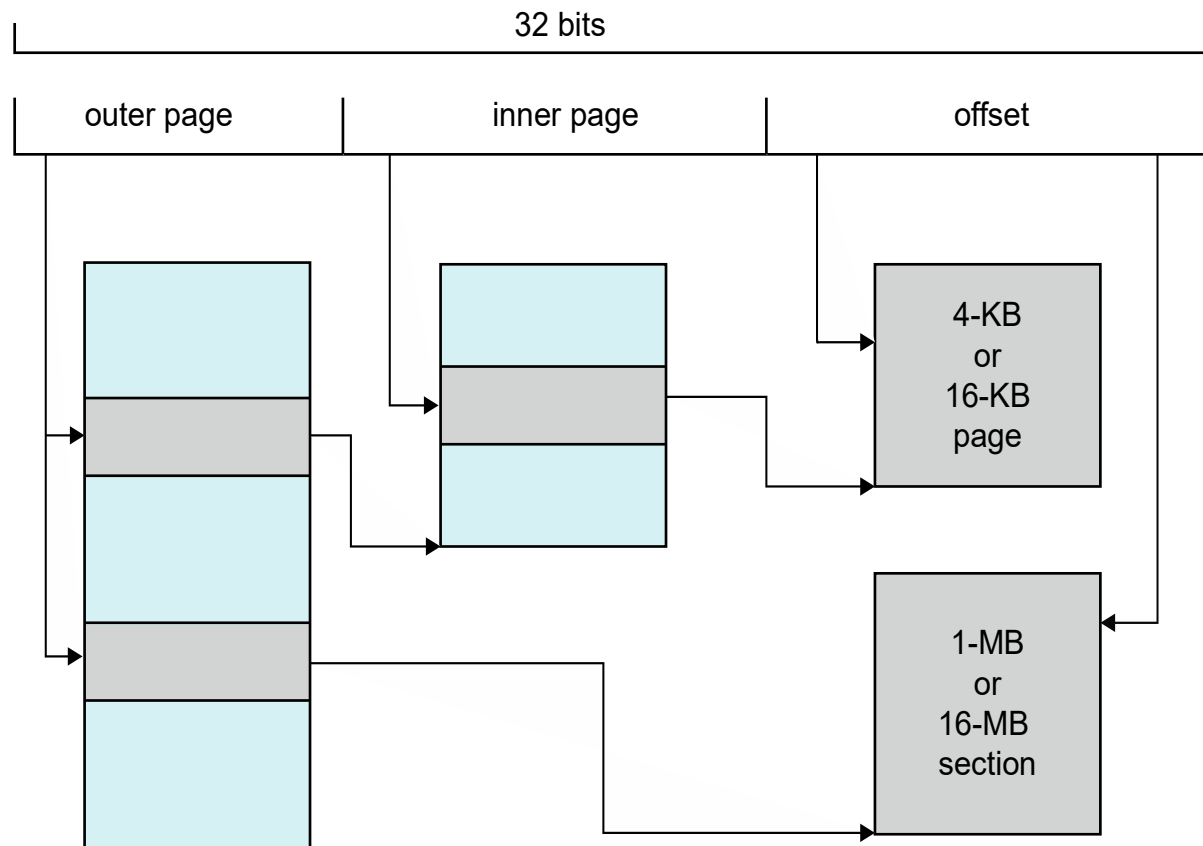


■ ARM Architecture

- Dominant mobile platform chip (Apple iOS and Google Android devices for example)
- Modern, energy efficient, 32-bit CPU
- 4KB / 16KB pages
- 1MB / 16MB pages (termed *sections*)
- One-level paging for sections, two-level for smaller pages
- Two levels of TLBs
 - Outer level has two micro TLBs (one for data, one for instruction)
 - Inner is single main TLB
 - First inner is checked, on miss outers are checked, and on miss page table walk performed by CPU



■ ARM Architecture



Linux

Three-level Paging in Linux.

