

Overview of Operating System Concepts

1. Introduction

- Modern Computer Hardware Organization
 - ✧ Processor(s)
 - ✧ Main memory
 - ✧ Mass storage
 - ✧ Various input/output devices
- What is Operating Systems
 - ✧ An Operating System is a program that acts as an intermediary or interface between a user of a computer and the computer hardware.
 - Convenient abstraction of complex hardware devices
 - Protection of access to shared resources
 - Security and Authentication
 - Communication amongst logical entities
 - ✧ An Operating System goals
 - Control/execute user or application programs
 - Make the computer system convenient to use
 - Ease the solving of user problems
 - Use the computer hardware in an efficient manner
- Computer System Components
 - ✧ Hardware
 - ✧ Operating system
 - ✧ Application programs (Application systems)
 - ✧ Users
- Classical views of Operating Systems
 - ✧ Resource manager
 - ✧ Control program
 - ✧ Command executer

2. Computer System Organization and Architecture

- Computer Hardware Architecture
 - ✧ *von Neumann* architecture
 - ✧ Harvard Architecture
 - ✧ Data representation
 - ✧ Addressing mode
 - ✧ Registers
 - ✧ Instruction system
 - ✧ Memory system
 - ✧ Interrupt controller
 - ✧ Input/output controller

- ✧ Information protecting mechanism
- Computer System Organizations
 - ✧ CPU
 - ✧ Main memory
 - ✧ Bus
 - ✧ Device controller
 - ✧ Device driver
 - ✧ Device
- Interrupt
 - ✧ The basic interrupt mechanism
 - ✧ Asynchronous processing
 - ✧ Interrupt request line
 - ✧ Interrupt handler routine
 - ✧ Interrupt vector
 - ✧ Interrupt controller
 - ✧ Maskable and non-maskable interrupts
 - ✧ Interrupt chaining
 - ✧ Interrupt priority
 - ✧ Interrupt types
 - Traps (Exceptions)
 - External interrupts
 - System calls
 - ✧ Various interrupt attributes.
 - Asynchronous vs. Synchronous.
 - External/Hardware vs. Internal/Software.
 - Implicit vs. Explicit.
- Storage Structures
 - ✧ Registers
 - ✧ Cache
 - ✧ Main memory
 - ✧ Secondary storage (NVM)
 - ✧ Hard disks
- Main Memory Management
 - ✧ Memory management dynamics
 - ✧ Dual-mode operation
 - ✧ Memory Protection
- I/O Structure
- Multiprogrammed batch systems
- Time-Sharing Systems
- Real-Time Systems
- Personal/Desktop Computers
- Multiprocessor Systems
 - ✧ SMP
 - ✧ Multicore

- Clustered Systems
- Networked Systems
- Distributed Systems
- Web-based Systems
- Handheld & Mobile Systems

3. Functional Views of Operating Systems

- Computer Dynamics
 - ✧ Instruction cycle with interrupts
 - ✧ External interrupt
 - ✧ Interrupt handler
 - ✧ Interrupt-driven I/O
 - ✧ Interrupt timeline of CPU and I/O device
 - ✧ Synchronous and asynchronous I/O methods
 - ✧ Direct memory access (DMA)
 - ✧ System calls
- Hardware Protection
 - ✧ Dual-mode operation
 - Monitor mode/Kernel mode
 - User mode
 - Privileged instructions
 - User mode to kernel mode transfer
 - ⊙ System call
 - ⊙ Interrupt
 - ⊙ Trap
 - ✧ I/O protection
 - System calls for I/O protection
 - ✧ Memory protection
 - Base register (relocation)
 - Limit register (bounded)
 - ✧ CPU protection
 - Timer
- Fundamental OS Concepts
 - ✧ Thread
 - ✧ Address space
 - ✧ Process
 - ✧ Dual-mode operation
- Resource Management
 - ✧ Process management
 - ✧ Memory management
 - ✧ File management
 - ✧ Mass-storage management
 - ✧ Cache management
 - ✧ I/O management

- Free and Open Source Operating Systems
 - ✧ GNU project
 - ✧ GNU/Linux
 - ✧ BSD UNIX

4. Structures of Operating Systems (1)

- Operating System Services
 - ✧ A user oriented view
 - ✧ User interface
 - ✧ Program execution
 - ✧ I/O operations
 - ✧ File systems
 - ✧ Communication
 - ✧ Error detection
 - ✧ Resource allocation
 - ✧ Accounting
 - ✧ Protection and security
- Common Operating System Components
 - ✧ A system oriented view
 - ✧ Process management
 - ✧ Main memory management
 - ✧ File management
 - ✧ Mass-storage management
 - ✧ I/O management
 - ✧ Networking
 - ✧ Command interpreter
 - ✧ GUI
 - ✧ Protection and security
 - ✧ Error detection and response
 - ✧ Accounting
- System Calls and APIs
 - ✧ Relationship of API, system call and OS
 - ✧ Standard C library and POSIX
- System Programs
- Operating System Design and Implementation

5. Structures of Operating Systems (2)

- Structure/Organization/Layout of OS
 - ✧ Monolithic (one unstructured program)
 - Traditional UNIX system structure
 - Linux system structure
 - Google's Android
 - ✧ Layered
 - Win3.1

- OS/2
- ✧ Microkernel
 - Architecture of a typical microkernel
 - Mach 3.0
 - MAC OS
 - MINIX
 - Windows NT, XP, 7.0
 - QNX Neutrino RTOS
- ✧ LKMS
- ✧ Virtual Machines

6. Introduction to Process (1)

- Basic Concepts
 - ✧ Segments in a process
 - Text, Data, Heap and Stack
 - ✧ Context of a process
 - ✧ Process virtual memory (PVM) in IA-32
 - ✧ Attributes of process
 - Process ID
 - Parent process ID
 - User ID
 - Process state
 - Process priority
 - Program counter
 - CPU registers
 - Memory management information
 - I/O status information
 - Access control
 - Accounting information
- Process Table and Process Control Block
 - ✧ Process table and its contents
 - ✧ PCB and its contents
- Process States and Transitions
 - ✧ Running, ready and waiting/blocked states
 - ✧ Process states transitions
 - ✧ Five-states process model
- Operations on Process
 - ✧ Process creation
 - Details in creating a process
 - [Programming](#): forking a separate process
 - [Programming](#): vforking a sharing space process
 - ✧ Process termination
 - Why terminating a process
 - Details in terminating a process

- Zombie process
- Orphan process
- [Programming](#): forking without waiting child termination
- [Programming](#): forking with waiting child termination
- [Programming](#): vfork, execv and wait
- [Programming](#): vfork, execv without wait

7. Introduction to Process (2)

- Unix and Linux Examples
 - ✧ UNIX SVR4 process states model
 - ✧ Linux process representation
 - ✧ Linux process state
- Process Scheduling
 - ✧ What is process scheduling
 - ✧ Scheduling queues
 - Job queue
 - Ready queue
 - Device/Waiting queues
 - ✧ Types of process schedulers
 - Long-term scheduler
 - Medium-term scheduler
 - Short-term scheduler
 - ✧ Process swapping
 - ✧ Process scheduling activities
 - ✧ Process scheduling queues
- Process Switching
 - ✧ Context switching (steps)
 - ✧ Mode switching

8. Inter-process Communication (1)

- Overview
 - ✧ Independent process and cooperating process
 - ✧ Why cooperating
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
 - ✧ Two fundamental models of IPC
 - Shared memory and message passing
- Shared-memory Systems
 - ✧ Producer-consumer problem with shared-memory
 - A producer and a consumer
 - Concurrency
 - Unbounded and bounded buffer

- ✧ Linux IPCs Limits
- ✧ Linux shared memory
 - System calls
 - [Programming](#): Single-writer-single-reader problem
- ✧ POSIX shared memory
 - POSIX APIs for memory sharing
 - [Programming](#): Producer-consumer problem

9. Inter-process Communication (2)

- Message-passing Systems
 - ✧ Primitives for message-passing
 - ✧ Variable and fixed message size
 - ✧ Message format
 - Linux message structure
 - ✧ Logical communication link
 - Direct or indirect communication
 - Synchronous or asynchronous communication
 - Automatic or explicit buffering
 - ✧ Synchronization in Message-passing Systems
 - Blocking and unblocking
 - ✧ Buffering
 - ✧ Linux message passing
 - System calls
 - [Programming](#): Sending & receiving process demo
 - ✧ POSIX message passing
 - POSIX APIs for memory sharing
 - [Lab work](#): Sending & receiving process demo

10. Inter-process Communication (3)

- Pipes
 - ✧ Ordinary pipes
 - ✧ Naked pipes
 - ✧ Linux ordinary pipes
 - System calls
 - ✧ Linux named pipes
 - System calls
 - [Programming](#): single pipe buffer
 - [Programming](#): total pipes capacity
 - [Programming](#): ordinary pipe
 - [Programming](#): named pipe between parent and children
 - [Programming](#): named pipe between two arbitrary Processes

11. Inter-process Communication (4)

- Communications in Client-Server Systems

- ✧ Sockets
- ✧ Linux socket programming
 - Socket APIs
 - [Programming](#): socket-server-BBS
- ✧ Remote Procedure Calls

12. Threads (1)

- Overview
 - ✧ Tasks/Processes vs. Threads
 - ✧ Benefits of multithreaded programming
 - Responsiveness
 - Resource sharing
 - Economy
 - Scalability
 - ✧ Thread states
 - Running
 - Ready
 - Blocked
 - No suspend state
 - ✧ Terminating a thread
- Multicore Programming
 - ✧ Parallelism and concurrency
 - ✧ Multithreading
 - ✧ Programming challenges
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging
 - Identifying tasks
 - ✧ Data parallelism and task parallelism
- User and Kernel Level Threads
 - ✧ Thread libraries
 - POSIX Pthreads
 - ✧ User level threads (ULT)
 - Kernel activity for ULTs
 - Advantages and inconveniences of ULTs
 - ✧ Kernel level threads (KLT)
 - Linux threads
 - Advantages and inconveniences of KLTs
 - ✧ Hybrid ULT/KLT Approaches
- Multithreading Models
 - ✧ Many-to-one model,
 - ✧ One-to-one model, and
 - ✧ Many-to-many model.

- ✧ Two-level model
- ✧ Light-weight process (LWP)

13. Threads (2)

- Threads Libraries
 - ✧ POSIX Pthreads
 - POSIX Pthreads APIs
 - [Programming](#): pthread creating
 - [Programming](#): pthread memory sharing
 - [Programming](#): pthread setting stack
- Implicit Threading
 - ✧ Threads pools
 - [Lab work](#): Implementation of a threads pool
 - ✧ OpenMP
 - [Programming](#): OpenMp demo
 - [Programming](#): OpenMP matrix adding
- Threading Issues
 - ✧ fork() and execv()
 - [Programming](#): pthread forking demo
 - ✧ Signal handling
 - Linux system calls
 - [Programming](#): sigaction-demo
 - ✧ Thread cancelation
 - Asynchronous cancellation
 - Deferred cancellation

14. Threads (3)

- ✧ Thread local storage (TLS)
 - [Programming](#): TLS implemented by __thread in language level
 - [Programming](#): TLS implemented by pthread_key_create
 - [Programming](#): tls_key and bounding data structure
- Linux clone()
 - ✧ Constants with flags in clone()
 - [Programming](#): Linux clone() demo
 - [Lab work](#): Programming with Linux clone()
 - ✧ Windows Threads

15. Cooperating Processes

- Introduction
 - ✧ Process cooperation
 - ✧ Data consistency
- Race Condition
 - ✧ Race condition

- ✧ Critical section
- ✧ Atomic operations
- The Critical-Section Problem
 - ✧ The critical-section problem
 - ✧ Three essential criteria to solve the critical-section problem
 - Mutual exclusion
 - Progress
 - Bounded waiting
 - ✧ Preemptive and Non-preemptive Kernels
 - ✧ Types of solutions to critical-section problem
 - Software-based solutions
 - Hardware-based solutions
 - Operating system solutions
 - Programming language solutions
 - ✧ Software solutions to critical-section problem
 - [Algorithm 1-5](#)
 - [Peterson's algorithm](#)
 - [Peterson's solution for n processes](#)
 - [Lamport's Bakery algorithm for n processes](#)
 - [Eisenberg-McGuire's algorithm for n processes](#)
 - **Lab work:** [Peterson's algorithm](#)

16. Process Synchronization (1)

- Synchronization Hardware
 - ✧ Atomic (non-interruptible) hardware instructions
 - test_and_set()
 - compare_and_swap()
 - ✧ Mutual exclusion with hardware instructions
- Mutex Lock
 - ✧ Concept of mutex lock
- Semaphores
 - ✧ Concepts
 - Semaphore
 - Semaphore operations
 - ⊙ wait(S) and signal(S)
 - ⊙ P and V operations
 - Solution to critical-section problem with semaphores
 - ✧ Implementation of semaphore
 - Binary Semaphore Implementation
 - Counting Semaphore Implementing
 - ✧ Deadlock and starvation with semaphore
 - ✧ Incorrect use of semaphores
- Solution to Classical Problems with Semaphores
 - ✧ Classical Problems

- Bounded-buffer problem
- Readers-writers problem
- Dining-philosophers problem

17. Process Synchronization (2)

- Monitors
 - ✧ The monitor type
 - ✧ Monitor usage
 - ✧ Monitor condition type
 - ✧ Monitor condition variables
 - ✧ Monitor condition operations
 - cwait(x) and csignal(x)
 - ✧ Solution to the dining-philosophers problem with monitor
 - ✧ Monitor implementation using semaphores
- Deadlock
 - ✧ What is deadlock (system model)
 - ✧ Four necessary conditions of deadlock
 - ✧ Resource allocation graph
 - ✧ Methods for handling deadlock
 - Preventing and Avoiding
 - Detecting and Recovering
 - Ignoring
 - ✧ The Banker's algorithm
 - [Safety algorithm](#)
 - [Resource-request algorithm](#)
 - Limitations of Deadlock Avoidance
 - [Programming](#): the Banker's algorithm

18. Process Synchronization (3)

- Synchronization Examples
 - ✧ Linux gcc __sync atomic family
 - [Programming](#): demo for Linux __sync__ type
 - ✧ POSIX mutex locks
 - [Programming](#): demo for POSIX mutex lock
 - ✧ POSIX unnamed semaphore
 - [Programming](#): demo for POSIX unnamed semaphores
 - ✧ POSIX named semaphore
 - [Programming](#): demo for POSIX named semaphores
 - ✧ POSIX Synchronization
 - [Programming](#): Solution to multi-producer-multi-consumer problem with POSIX semaphores and shared memory

19. CPU Scheduling (1)

- Basic Concepts

- ✧ CPU-I/O burst cycle
- ✧ CPU burst time
- ✧ When to make CPU scheduling
- ✧ preemptive scheduling
 - during system call
 - during interruption
- ✧ Dispatcher and dispatcher latency
- CPU Scheduling Criteria
 - ✧ Scheduling goals of different systems
 - All systems
 - Batch systems
 - Interactive systems
 - Real-time systems
 - ✧ Scheduling criteria
 - CPU utilization
 - Throughput
 - Turnaround time
 - Waiting time
 - Response time
- Simple Scheduling Algorithms
 - ✧ [First-come, first-served scheduling](#)
 - ✧ [Shortest-job-first scheduling](#)
 - ✧ [Priority scheduling](#)
- Advanced Scheduling Algorithms
 - ✧ [Round-robin scheduling](#)
 - ✧ [Multiple-priority queues scheduling](#)
 - ✧ [Multilevel feedback queue scheduling](#)
 - ✧ Thread scheduling

20. CPU Scheduling (2)

- Multiple-Processor Scheduling
 - ✧ Processor affinity
 - Soft affinity
 - Hard affinity
 - ✧ Load balancing
 - Push migration
 - Pull migration
 - Load balancing vs. processor affinity
 - ✧ Multicore processors
 - Memory stall
- Real-Time CPU Scheduling
 - ✧ Soft real-time and hard real-time systems
 - ✧ Minimizing Latency
 - Event Latency

- Interrupt latency
- Dispatch latency
- ✧ Priority-based scheduling
- ✧ Rate-monotonic scheduling
- ✧ Earliest-deadline-first scheduling
- ✧ Least Laxity First algorithm
- ✧ Proportional share scheduling
- ✧ POSIX real-time scheduling
 - [Programming](#): demo for POSIX real-time scheduling
- Linux Scheduling
 - ✧ Completely Fair Scheduler
 - ✧ RTLinux
- Algorithms Evaluation
 - ✧ Deterministic modeling
 - ✧ Queuing models
 - ✧ Simulations
 - ✧ **Lab work**: CPU scheduling simulations

21. Introduction to Memory Management

- Basic Concepts
 - ✧ Hardware address protection
 - ✧ Address binding
 - ✧ Logical address space and physical address space
 - ✧ Memory-mapped unit (MMU)
 - ✧ Relocation register
 - ✧ Hardware support for relocation and limit registers
 - ✧ Hardware translation of address
 - ✧ Dynamic linking
 - ✧ Swapping
 - ✧ Program size vs. memory size
- Memory Management Requirements
 - ✧ Relocation
 - ✧ Protection
 - ✧ Sharing
 - ✧ Logical organization
 - ✧ Physical organization
- Memory Partitioning
 - ✧ Contiguous allocation
 - ✧ Real memory management techniques
 - ✧ Fixed partitioning
 - Dynamics of fixed partitioning
 - Fragmentation
 - ✧ Variable partitioning
 - Internal fragmentation

- External fragmentation
- Compaction
- Basic placement algorithms: First-fit, Next-fit, Best-fit and Worst-fit
- ✧ Buddy system

22. Memory Segmentation and Paging

- Simple Segmentation
 - ✧ User view of a program
 - ✧ Dynamics of simple segmentation
 - Segment table
 - Base & limit
 - Segment-table base register (STBR)
 - Segment-table length register (STLR)
 - ✧ Address translation
 - ✧ Protection
 - ✧ Sharing
- Simple Paging
 - ✧ Non-contiguous to avoid external fragmentation
 - ✧ Page and page frame
 - ✧ Free-frame list
 - ✧ Logical address structure
 - ✧ Page table
 - ✧ Page-table base register (PTBR)
 - ✧ Page-table length register (PTLR)
 - ✧ Translation Look-aside Buffer (TLB)
 - ✧ Address translation
 - ✧ Paging hardware with TLB
 - ✧ Protection
 - ✧ Sharing
 - ✧ Structure of page tables
 - Hierarchical page tables
 - Hashed page tables
 - Inverted page tables
- Examples

23. Virtual Memory and Demand Paging (1)

- Virtual Memory
 - ✧ Concept of virtual memory
 - ✧ Virtual memory addressing
 - ✧ Virtual address space
 - ✧ Shared libraries with virtual memory
 - ✧ Process execution with virtual memory
- Demand Paging

- ✧ Concepts of demand paging
- ✧ Page table
- ✧ Present-bit and modified-bit
- ✧ Page fault
 - Steps in handling a page fault
- ✧ Page replacement
 - Steps in handling a page replacement
- ✧ Demand paging with TLB
- ✧ Stages in demand paging (worse case)
- ✧ Performance of demand paging
- Demand Paging Considerations
 - ✧ Locality
 - ✧ Thrashing
 - ✧ Memory-mapped files
 - ✧ Buddy system
 - ✧ Slab system
 - ✧ Copy-on-write
 - ✧ Other issues

24. Virtual Memory and Demand Paging (2)

- Page Replacement Algorithms
 - ✧ Basic page replacement
 - ✧ FIFO algorithm
 - ✧ Optimal page replacement
 - ✧ Least recently used (LRU) algorithm
 - Counter implementation
 - Stack implementation
 - Hardware matrix LRU implementation
 - ✧ LRU approximation algorithms
 - Reference bit and reference byte
 - Second-chance algorithm (CLOCK algorithm)
 - Enhanced second-chance algorithm
 - ✧ Cleaning policy
 - ✧ Lab work: page replacement algorithms
- Combined Segmentation and Paging
- Virtual Memory Policies

25. Mass Storage Systems

- Introduction
 - ✧ Hard disk Structure
 - Disk Head and head crash
 - Track, section and cylinder
 - Host controller and Hard disk controller
 - Transfer rate and positioning time

- Hard disk performance
 - Logical block
- ✧ Hard disk attachment
- Disk Scheduling
 - ✧ Data access time
 - ✧ I/O request queues
 - ✧ First-come first-serve (FCFS)
 - Zigzag effects
 - ✧ Shortest seek time first (SSTF)
 - ✧ Elevator algorithms
 - SCAN scheduling and C-SCAN scheduling
 - LOOK scheduling and C-LOOK scheduling
 - ✧ Lab work: disk scheduling algorithms
- Disk Management
 - ✧ Swap-space Management
 - ✧ RAID structure
- Stable-Storage Implementation

26. File System Interface

- Basic Concepts
 - ✧ Files
 - ✧ File attributes
 - ✧ File type extensions
 - ✧ File operations
 - ✧ Access methods
 - Sequential access
 - Direct access
 - ✧ Types of file systems
- File Directories
 - ✧ Elements
 - ✧ Operations
 - ✧ Single-level directory
 - ✧ Two-level directory
 - ✧ Tree-structure directories
 - ✧ DAG directories
 - ✧ General graph directory
 - ✧ File system mounting
- File Sharing and Protection

27. File System Implementation

- File-System Structure
 - ✧ File control block (FCB)
 - ✧ File system layers
 - I/O control level

- Basic file system
- File organization module
- Logical file system
- ✧ On-storage structures
 - Boot control block, volume control block, directory structure, FCB
- ✧ In-memory structures
 - Mount table, directory-structure cache, system-wide open-file table, open-file table, buffers
- ✧ Directory Implementation
 - Linear list
 - Hash table
- Allocation Methods
 - ✧ Contiguous allocation
 - ✧ Linked allocation/Chained allocation
 - File allocation table (FAT)
 - ✧ Indexed allocation
 - One block scheme
 - Linked scheme
 - Two-level index scheme
 - ✧ Combined schemes
 - ✧ Performance analysis
- Free-Space Management
 - ✧ Bitmap with n blocks
 - ✧ Free-block list (Linked list)
 - ✧ Grouping and counting
 - ✧ Space maps
- Efficiency and Performance
- Recovery
- NFS
- Example: WAFL File System

28. I/O Systems

- Overview
- I/O Hardware
 - ✧ Port, Bus, PCI, PCIe, Expansion bus, Controller (host adapter)
 - ✧ Typical PC bus structure
 - ✧ Memory mapped I/O (MMIO)
 - ✧ I/O device-control registers
 - ✧ Polling (steps)
 - ✧ Interrupt
 - Interrupt-request line
 - Interrupt handler routine
 - Maskable and non-maskable interrupts

- Interrupt vector
- Difference among interrupt, fault, trap and exception
- ✧ Direct memory access (DMA)
 - Cycle stealing
 - DMA steps
- Application I/O Interface
 - ✧ Kernel I/O structure
 - ✧ I/O devices vary in many dimensions
 - Data transfer mode
 - Access mode
 - Transfer schedule
 - Sharing
 - Speed of operation
 - I/O direction
- Kernel I/O Subsystem
 - ✧ Services provided by kernel's I/O subsystem
 - Scheduling
 - Buffering
 - Caching
 - Spooling
 - Device reservation
 - Error handling.
 - ✧ I/O protection
 - ✧ I/O kernel structures
- Power management
- Transforming I/O Requests to Hardware Operations
 - ✧ 10 Steps
- **STREAMS
- Performance Analysis

== The End